

# Convolutional Neural Networks (CNNs) and Support Vector Machines (SVMs)

**CSCI-P556 Applied Machine Learning**  
**Lecture 18**

**D.S. Williamson**

# Agenda and Learning Outcomes

## Today's Topics

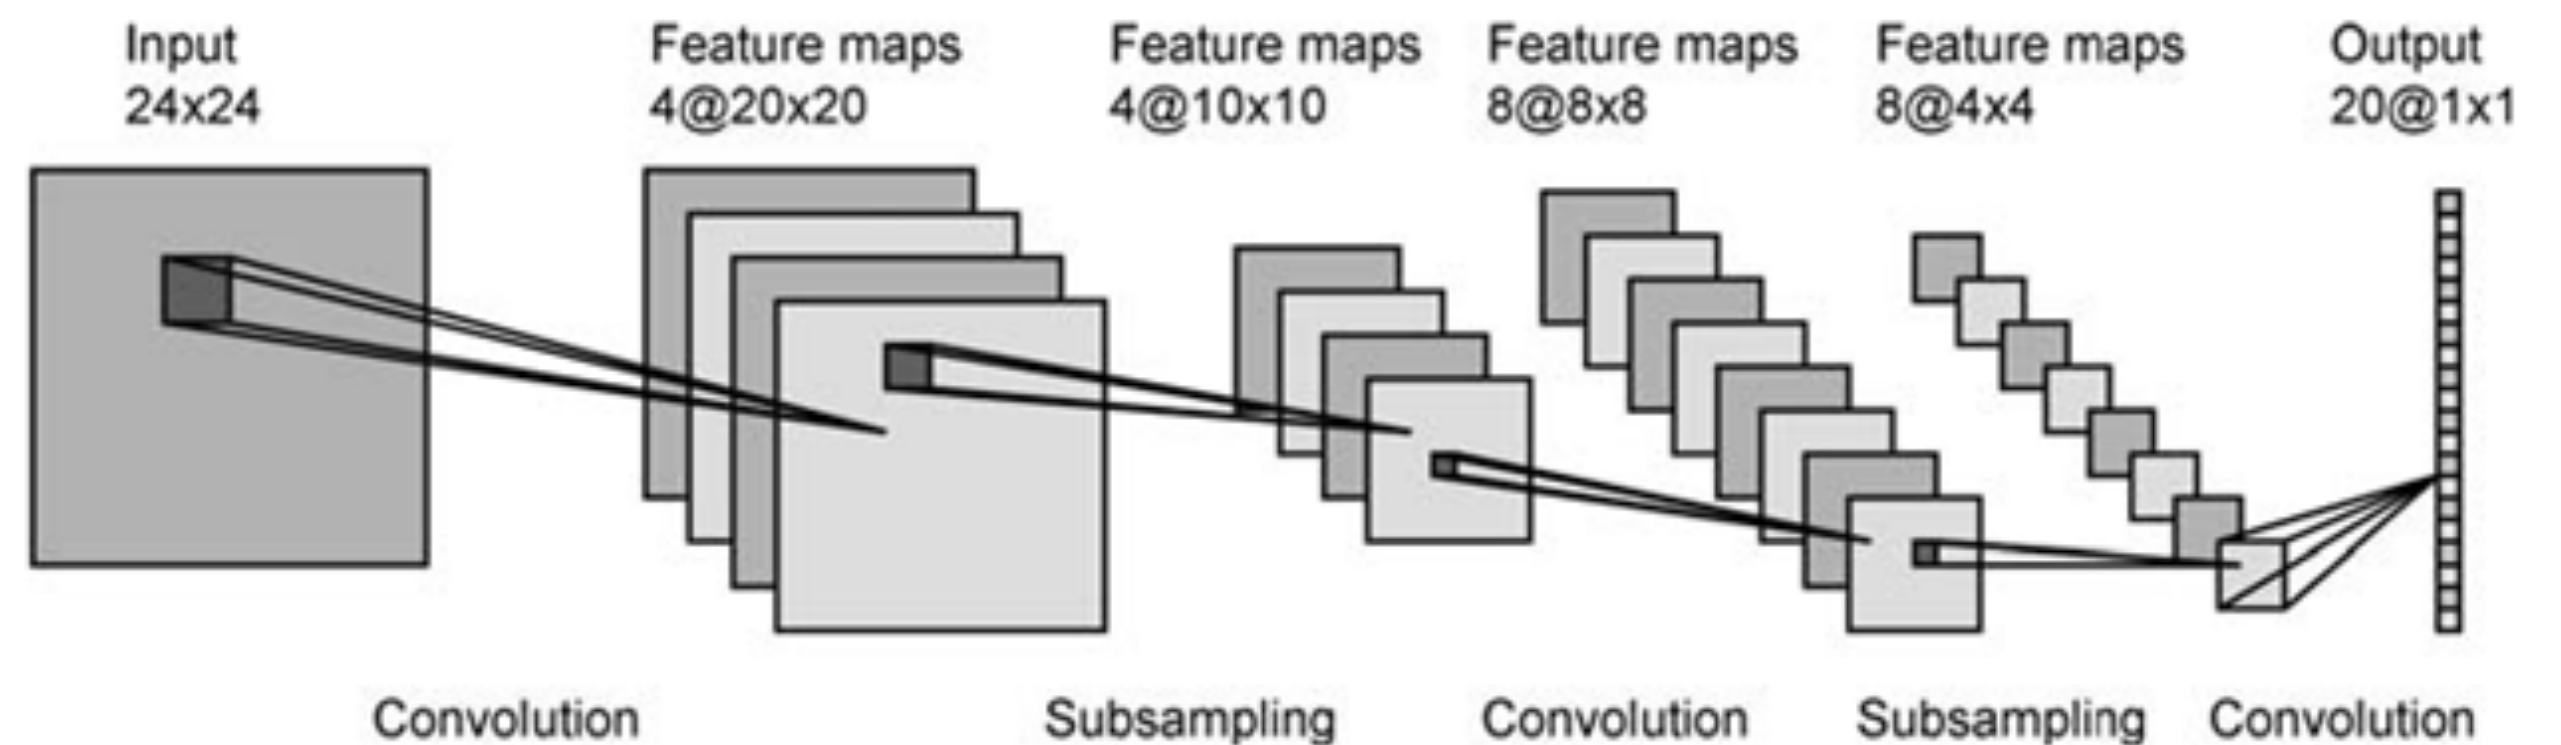
- **Topics:**
  - Convolutional Neural Networks (CNNs)
  - Support Vector Machines, Part I
- **Announcements**
  - Quiz #2 on Thursday
  - Project proposal comments (tomorrow)
  - hw#3 coming soon (today or tomorrow)

# Convolutional Neural Networks

# Convolutional Neural Networks (CNNs)

- Used for processing data with grid-like topology (i.e. images). Networks use convolution in place of general matrix multiplication
- Good at capturing local (short-term) dependencies and correlations (e.g. correlations amongst adjacent pixels in an image, or dependencies across nearby frequencies of an audio signal)
- There are four main operations in CNNs

- Convolution
- Nonlinear Activation Function (i.e. ReLU)
- Pooling (or Subsampling)
- Classification



# Convolution

- Convolution is a linear mathematical operation on two functions
- Given functions  $x(t)$  and  $w(t)$ , the convolution of  $x(t)$  and  $w(t)$  is as follows

$$s(t) = x(t) * w(t) = \sum_{a=-\infty}^{\infty} x(a)w(t-a)$$

- This step is used to extract “features” from an input, so it is also referred to as the **feature map** stage

# Example: Convolution on an Image

- Suppose you are given the following binary image, X

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

- You want to convolve this image with matrix, W as shown below

1	0	1
0	1	0
1	0	1

1 <sub>x1</sub>	1 <sub>x0</sub>	1 <sub>x1</sub>	0	0
0 <sub>x0</sub>	1 <sub>x1</sub>	1 <sub>x0</sub>	1	0
0 <sub>x1</sub>	0 <sub>x0</sub>	1 <sub>x1</sub>	1	1
0	0	1	1	0
0	1	1	0	0

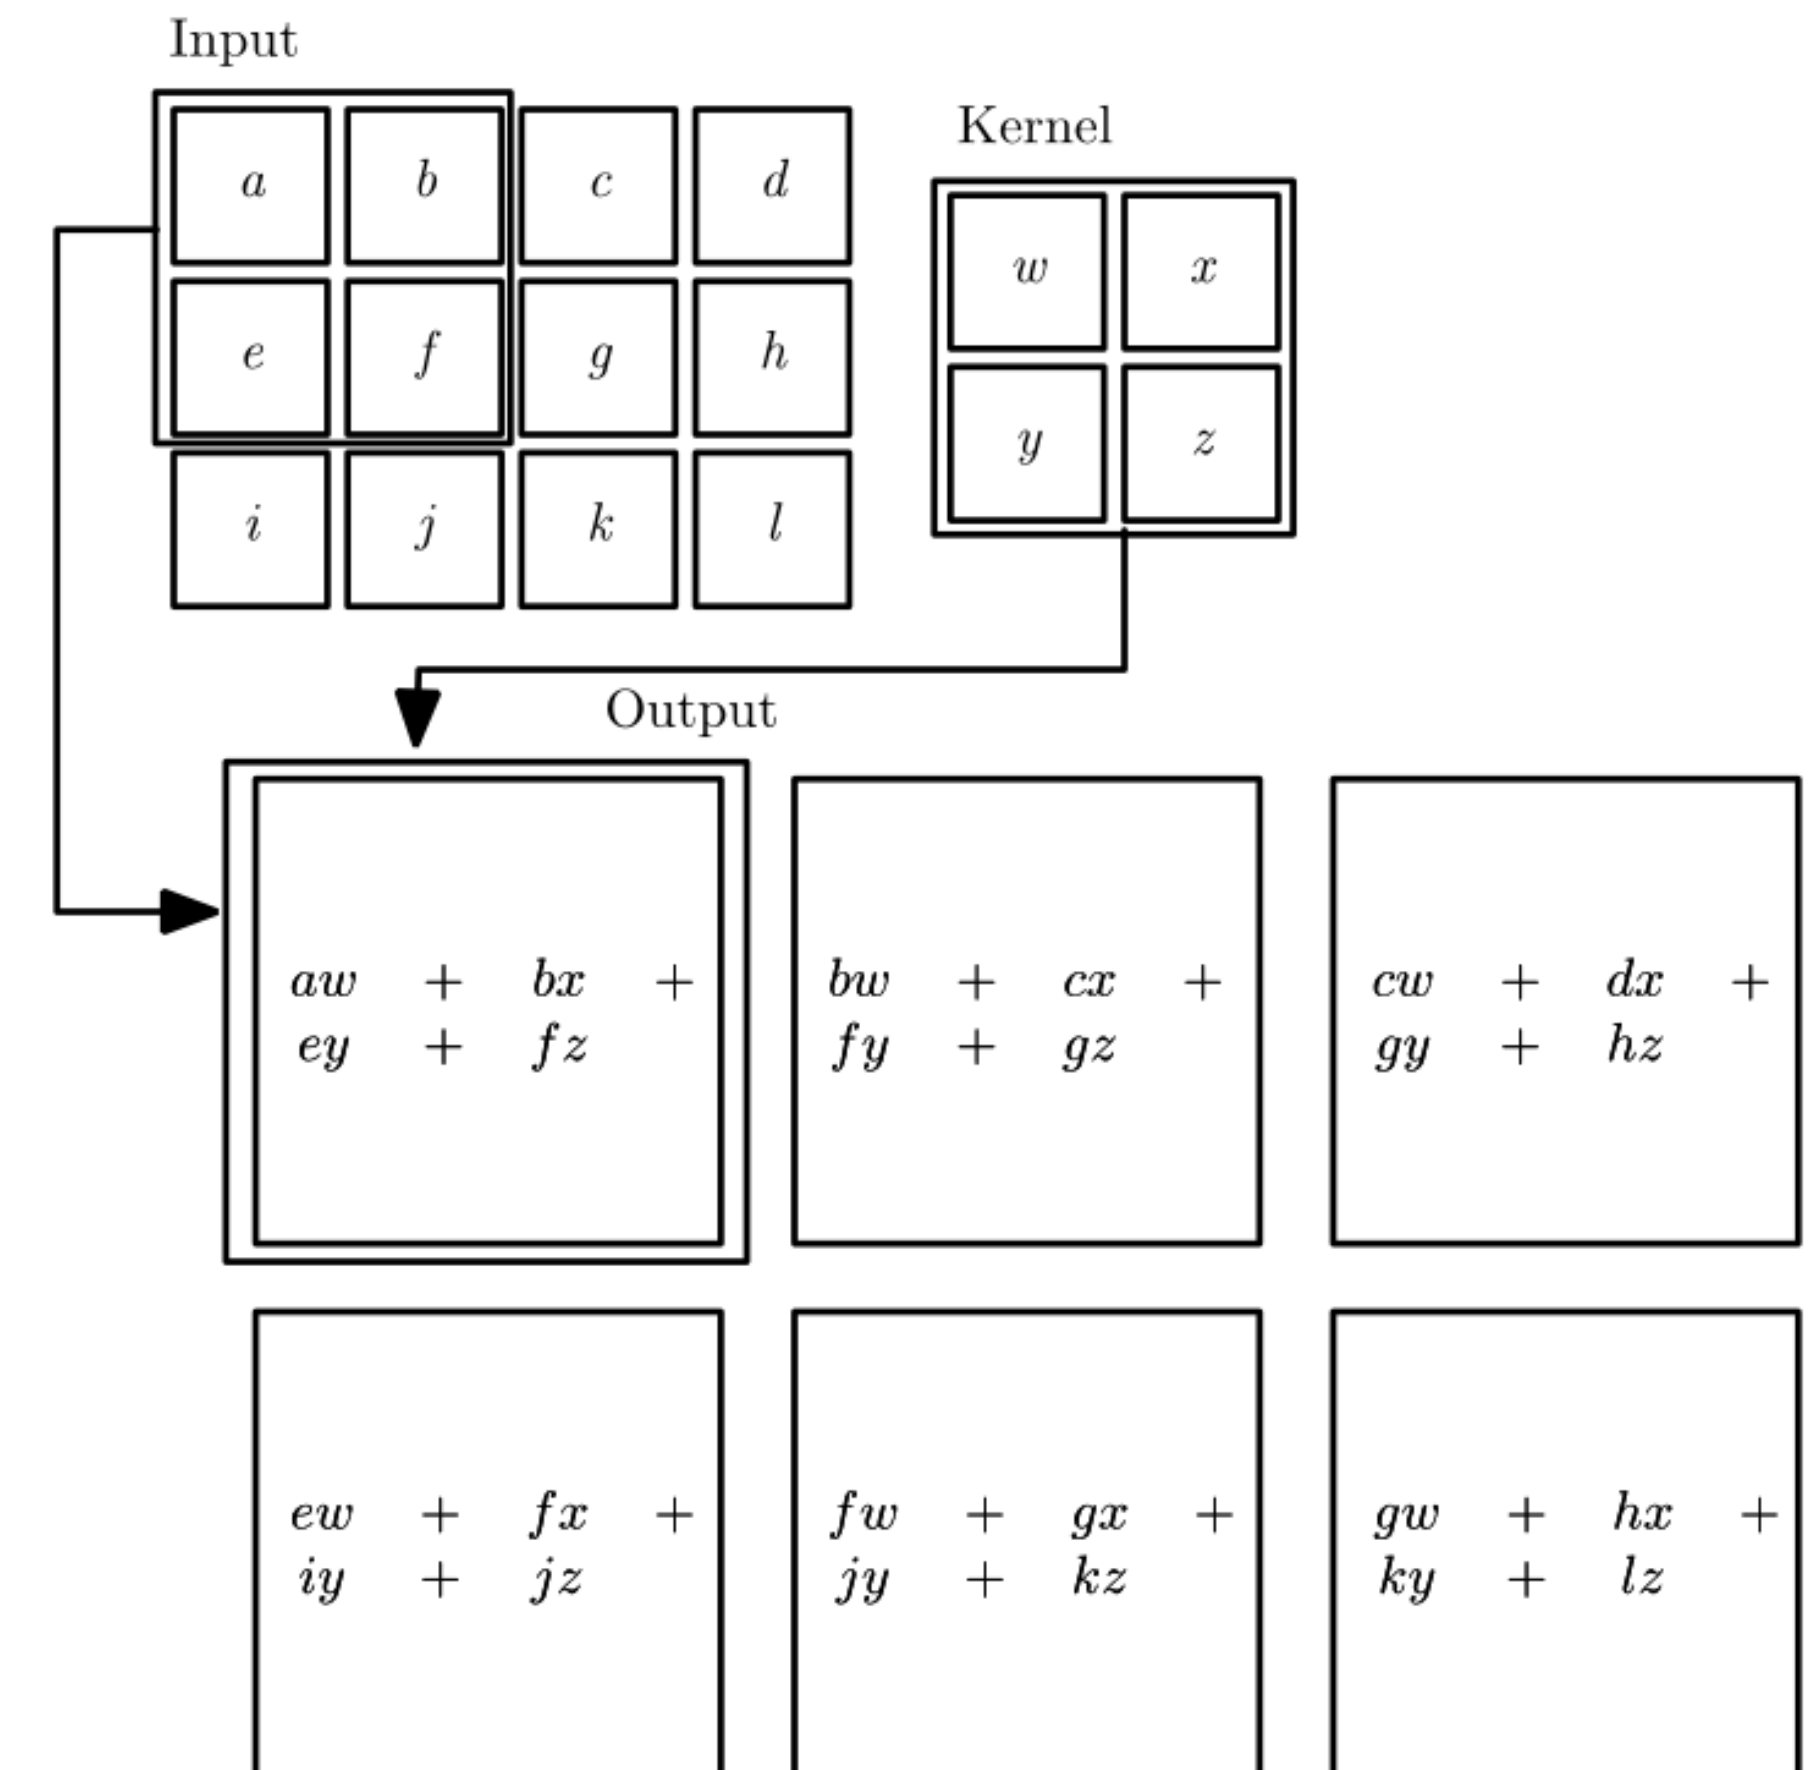
Image

4		

Convolved  
Feature


# Convolution Operation

- A mathematical depiction of the convolution operation on an input image
- A CNN learns the values of the filter (or kernel) on its own during the training process



# Feature Map


- The size of the Feature Map (resulting image after convolution) is controlled by three parameters
  - **Depth:** Number of different filters to use for the convolution operation
  - **Stride:** Number of pixels used to slide the filter across the input
  - **Zero-padding:** May pad the input with zeros around the border





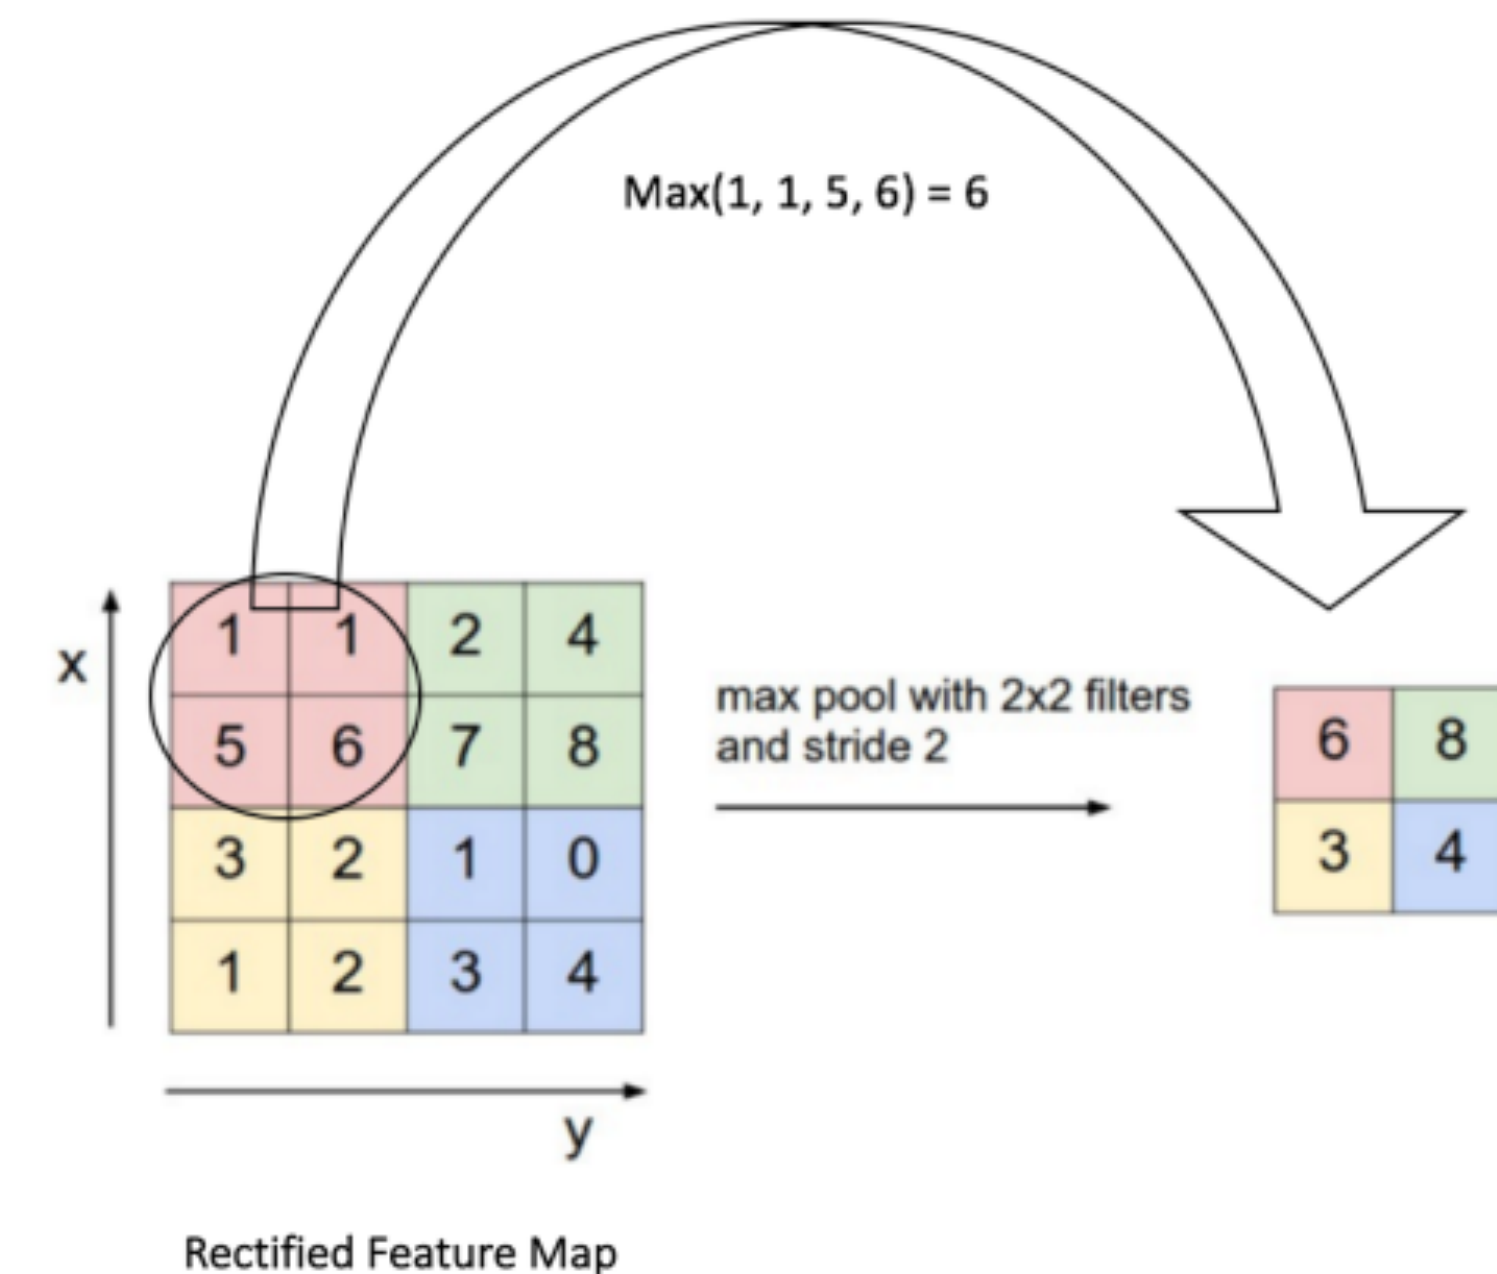
# Rectified Linear (ReLU) Activation

- A rectified linear (ReLU) activation operation may be applied after the convolution operation
  - Introduces nonlinearity to the network
  - $\text{ReLU}(x) = \max(0, x)$
  - Applied to every element (pixel)
  - Negative values are replaced by 0
- Other nonlinear activation functions may be used instead




# Pooling

- Pooling (aka spatial pooling, subsampling, or downsampling) is used to reduce the dimensionality of the feature map
- Different types of pooling include: Max, Average, Sum, etc.
- A window is defined, and the pooling operation is performed over the elements within that window
- The pooling window slides over the feature map by the stride amount
- It is applied to each feature map



# CNN

- Multiple layers of Convolution, Activation, and Pooling may be used in a CNN
- These layers act as feature extraction, to find useful features from the input
- Generally, a final Fully Connected layer is added via a DNN for classification or regression purposes



# CNN Training


- The Backpropagation algorithm is used to train the parameters of a CNN
- Basic steps:
  - Randomly initialize all filters (or kernels) and weights
  - Propagate the input forward through the layers of the CNN (convolution, activation, pooling, DNN) to get an output(s)
  - Calculate the error between the actual output(s) and the desired output(s)
  - Use Backpropagation to calculate the gradients and deltas, and then update the filters and weights accordingly



# An Early CNN Application

## Handwritten Zip code Recognition (1989)

- **Input:** binary pixels for each digit (16 x 16 dimensional)
- Padded with -1.
- **Output:** 10 digits, using sigmoid activations
- Scaled hyperbolic tangent activation function
- **Architecture:** 4 layers (12x8x8 – 12x4x4 -30 -10)
  - Layer 1: kernel of size 5x5, stride of 2, depth of 12
  - Layer 2: kernel of size 5x5, stride of 2, depth of 12
  - Layer 3: Fully-connected layer with 30 units
  - Layer 4: 10 unit output layer
- **Performance:** Trained on 7300 digits and tested on 2000 new ones
  - Achieved 1% error on the training set and 5% error on the test set
  - If allowing rejection (no decision), 1% error on the test set
  - This task is not easy



# CNN in PyTorch

- The following classes can be used within a defined class that implements the desired CNN architecture. (e.g. in the `__init__()` and `forward()` functions)
- **Convolution:** Conv2d (or Conv1d) to perform convolution operation
- **Pooling:** Maxpool2D, Avgpool2D,...

## CONV2D

Conv2d

**CLASS** `torch.nn.Conv2d(in_channels, out_channels, kernel_size, stride=1, padding=0, dilation=1, groups=1, bias=True, padding_mode='zeros')` [SOURCE]

Applies a 2D convolution over an input signal composed of several input planes.

In the simplest case, the output value of the layer with input size  $(N, C_{in}, H, W)$  and output  $(N, C_{out}, H_{out}, W_{out})$  can be precisely described as:

$$out(N_i, C_{out_j}) = bias(C_{out_j}) + \sum_{k=0}^{C_{in}-1} weight(C_{out_j}, k) \star input(N_i, k)$$

where  $\star$  is the valid 2D **cross-correlation** operator,  $N$  is a batch size,  $C$  denotes a number of channels,  $H$  is a height of input planes in pixels, and  $W$  is width in pixels.

This module supports **TensorFloat32**.

- `stride` controls the stride for the cross-correlation, a single number or a tuple.
- `padding` controls the amount of implicit padding on both sides for `padding` number of points for each dimension.
- `dilation` controls the spacing between the kernel points; also known as the à trous algorithm. It is harder to describe, but this [link](#) has a nice visualization of what `dilation` does.
- `groups` controls the connections between inputs and outputs. `in_channels` and `out_channels` must both be divisible by `groups`. For example,

Applies a 2D max pooling over an input signal composed or several input planes.

In the simplest case, the output value of the layer with input size  $(N, C, H, W)$ , output  $(N, C, H_{out}, W_{out})$  and `kernel_size`  $(kH, kW)$  can be precisely described as:

$$out(N_i, C_j, h, w) = \max_{m=0, \dots, kH-1} \max_{n=0, \dots, kW-1} input(N_i, C_j, stride[0] \times h + m, stride[1] \times w + n)$$

If `padding` is non-zero, then the input is implicitly zero-padded or controls the spacing between the kernel points. It is harder to describe, but this [link](#) has a nice visualization of what `dilation` does.

## AVGPOOL2D

AvgPool2d

**CLASS** `torch.nn.AvgPool2d(kernel_size, stride=None, padding=0, ceil_mode=False, count_include_pad=True, divisor_override=None)` [SOURCE]

Applies a 2D average pooling over an input signal composed of several input planes.

In the simplest case, the output value of the layer with input size  $(N, C, H, W)$ , output  $(N, C, H_{out}, W_{out})$  and `kernel_size`  $(kH, kW)$  can be precisely described as:

$$out(N_i, C_j, h, w) = \frac{1}{kH * kW} \sum_{m=0}^{kH-1} \sum_{n=0}^{kW-1} input(N_i, C_j, stride[0] \times h + m, stride[1] \times w + n)$$

If `padding` is non-zero, then the input is implicitly zero-padded on both sides for `padding` number of points.

### • NOTE

When `ceil_mode=True`, sliding windows are allowed to go off-bounds if they start within the left padding or the input. Sliding windows that would start in the right padded region are ignored.

# Support Vector Machines


# Warning!

- This is just a first introduction to SVMs – very basic version of SVM introduced
- We will discuss more on the advanced techniques later
  - Hence, Strong assumptions in this case




# The Classification Problem

- **Binary classification** can be viewed as the task of **separating classes in the feature space**
- This is accomplished by formulating a decision boundary
  - $\mathbf{w}$  is the slope of the line
  - $\mathbf{b}$  is the intercept




# Recall: Perceptrons

- **Perceptron Convergence Theorem:** If a classification problem is linearly separable, a perceptron will reach a solution in a finite number of iterations
- The solution weight vector is **not** unique. There are ***infinite possible solutions*** and decision boundaries.
  - Perceptrons find any separating hyperplane
  - The hyperplane depends on initialization and ordering of training points




# Recall: Perceptrons

- **Perceptron Convergence Theorem:** If a classification problem is linearly separable, a perceptron will reach a solution in a finite number of iterations
- The solution weight vector is **not** unique. There are **infinite possible solutions** and decision boundaries.
  - Perceptrons find any separating hyperplane
  - The hyperplane depends on initialization and ordering of training points
- **If done differently**




# Recall: Perceptrons

- **Perceptron Convergence Theorem:** If a classification problem is linearly separable, a perceptron will reach a solution in a finite number of iterations
- The solution weight vector is **not** unique. There are **infinite possible solutions** and decision boundaries.
  - Perceptrons find any separating hyperplane
  - The hyperplane depends on initialization and ordering of training points
- **If done differently....Again**




# Recall: Perceptrons

- **Perceptron Convergence Theorem:** If a classification problem is linearly separable, a perceptron will reach a solution in a finite number of iterations
- The solution weight vector is **not** unique. There are **infinite possible solutions** and decision boundaries.
  - Perceptrons find any separating hyperplane
  - The hyperplane depends on initialization and ordering of training points
- **If done differently....Again...And Again**



# Linear Separators


- Which is the best linear separator?
  - Depends on the goal
  - Goal is to classify accurately and generalize to new examples.






# Notion of Margins

- Many different hyperplanes can classify the data. Which one will work best?
- The **hyperplane** that **maximizes** the **separation between the two classes** (the margin)



# Intuition of a Margin



- Consider points A, B, and C
- We are quite confident in our prediction for A because it is far from the decision boundary.
- In contrast, we are not so confident in our prediction for C because a slight change in the decision boundary may flip the decision
- Given a training set, we would like to make all predictions correct and confident! This leads to the concept of margin





# Why Max Margin?

- Minimizes generalization error. Works well on **Future data**
- Minimizes Complexity. **Fewer support vectors**
- Minimizes the capacity of the classifier. **Eliminates overfitting**




# Decision Boundary

**Need to know distance from point to the decision boundary**

- Given a decision boundary (e.g. linear discriminant function)

$$g(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + b = 0$$

- To find its distance to a given pattern  $\mathbf{x}$ , project  $\mathbf{x}$  onto the decision boundary




# Decision Boundary (cont.)

- $\mathbf{x}$  can be re-written as a function of the projection and the weights

$$\mathbf{x} = \mathbf{x}_p + r \frac{\mathbf{w}}{||\mathbf{w}||}$$

- $\mathbf{x}_p$  is  $\mathbf{x}$ 's projection
- The second term arises from the fact that the weight vector is perpendicular to the decision boundary
- The algebraic distance  $r$  is positive if  $\mathbf{x}$  is on the positive side of the boundary and negative if  $\mathbf{x}$  is on the negative side

$$g(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + b = 0$$



# Decision Boundary (cont.)


- Since  $\mathbf{x}$  can be written in terms of its projection and  $r$ , then the decision boundary can as well:

$$\begin{aligned} g(\mathbf{x}) &= g\left(\mathbf{x}_p + r \frac{\mathbf{w}}{\|\mathbf{w}\|}\right) \\ &= \mathbf{w}^T \left(\mathbf{x}_p + r \frac{\mathbf{w}}{\|\mathbf{w}\|}\right) + b \\ &= \mathbf{w}^T \mathbf{x}_p + b + r \|\mathbf{w}\| \\ &= r \|\mathbf{w}\| \end{aligned}$$

- Thus,  $r = \frac{g(\mathbf{x})}{\|\mathbf{w}\|}$


As a special case, for the origin,  $r = b/\|\mathbf{w}\|$

$$g(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + b = 0$$



# Margins

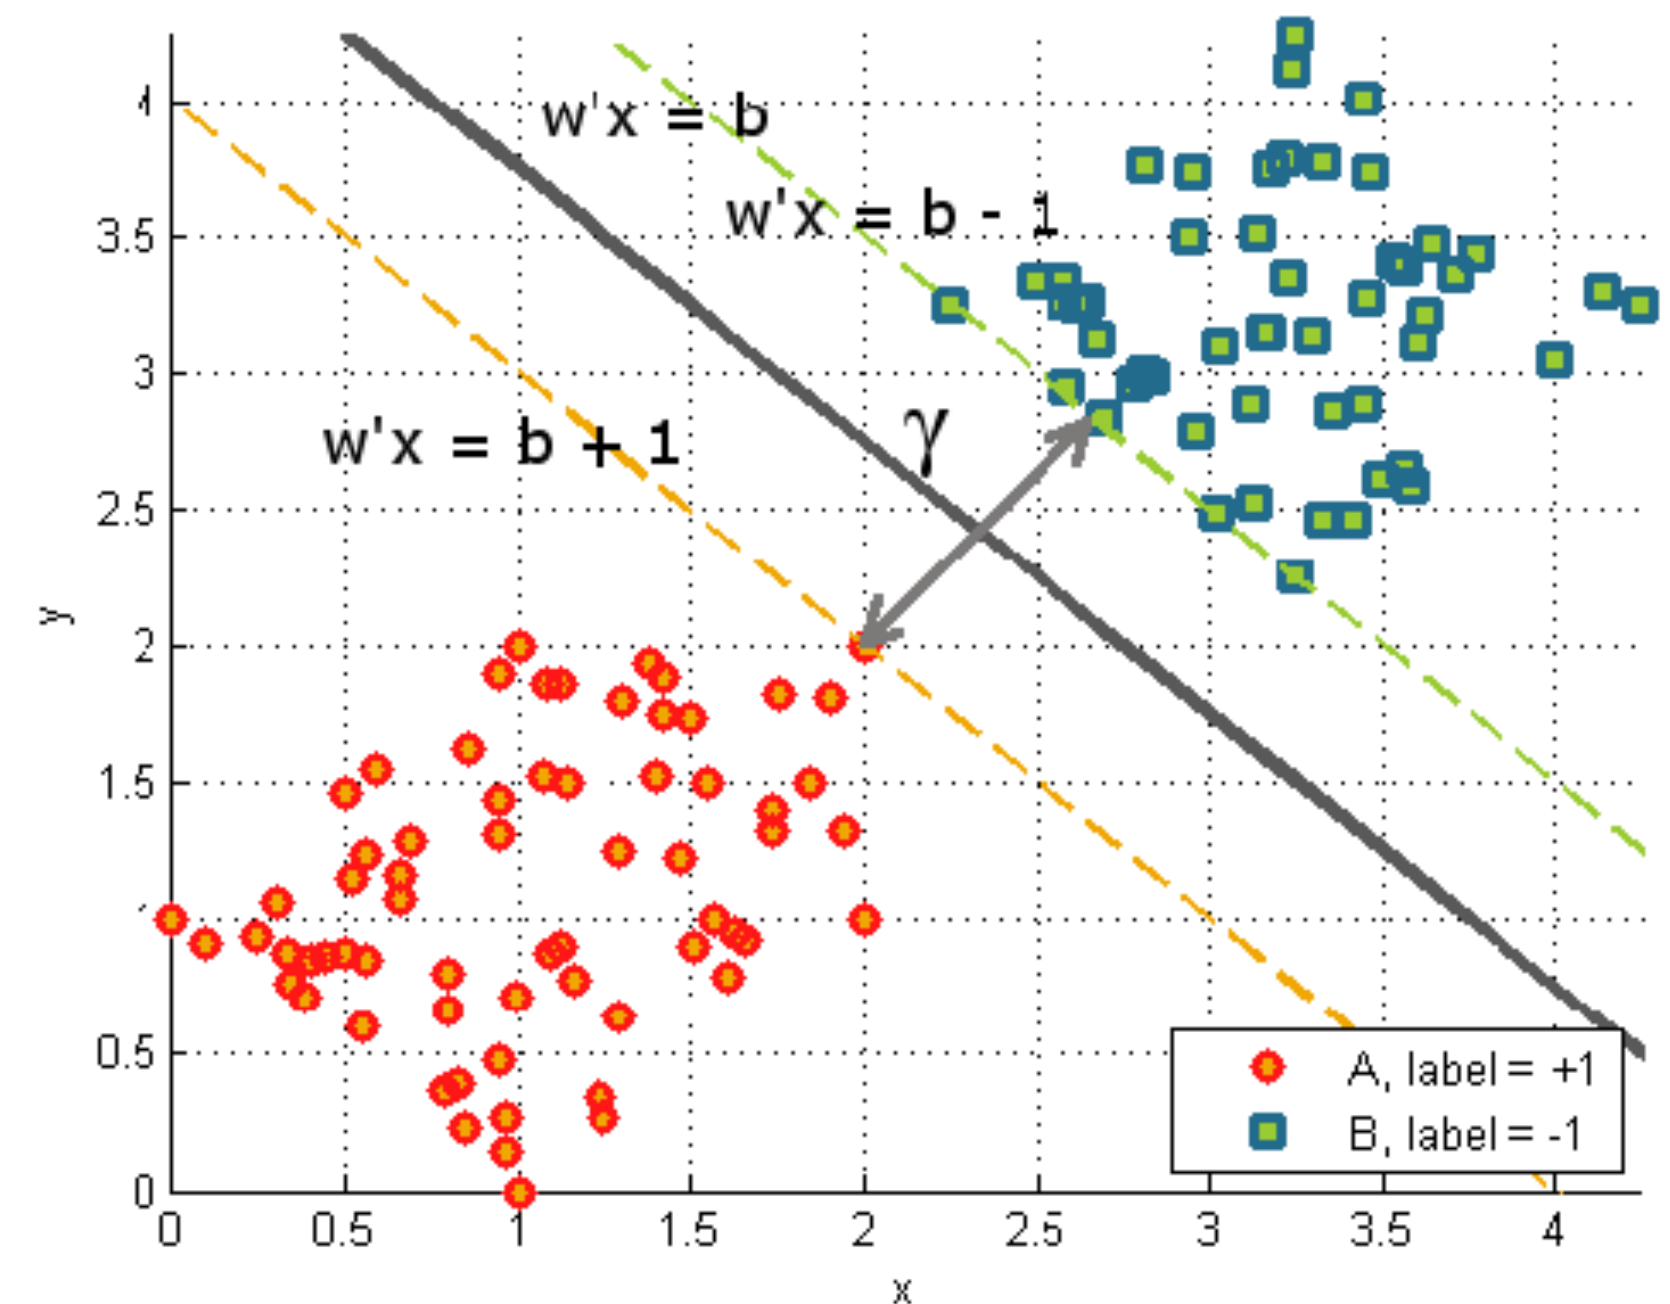
- Distance from example  $\mathbf{x}_i$  to the separator is  $r = \frac{\mathbf{w}^T \mathbf{x}_i + b}{\|\mathbf{w}\|}$
- Examples closest to the hyperplane are **support vectors**.
- **Margin**  $\gamma$  of the separator is the distance between support vectors.



# Notation

- We denote the classifier,  $f(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + b$ , for all  $\mathbf{x} \in \mathbb{R}^n$
- **Assumptions**
  - Supporting hyperplanes  $\mathbf{w}^T \mathbf{x} + b = \pm 1$
  - The distance between the two supporting hyperplanes is the margin, which is  $\gamma = 2$
- To achieve **scale invariance** we divide the classifier by  $\|\mathbf{w}\|_2$ .  
Then the supporting hyperplanes are  $\mathbf{w}^T \mathbf{x} + b = \pm \frac{1}{\|\mathbf{w}\|_2}$

margin is  $\gamma = \frac{2}{\|\mathbf{w}\|_2}$ .





# Max margin Classifier

- Given a **linearly separable** training set  $S = \{(\mathbf{x}^{(i)}, y^{(i)}) : i = 1, \dots, N\}$ , we would like to find a classifier with a maximum margin,  $\gamma$
- This can be represented as an optimization problem.

$$\max_{\mathbf{w}, b, \gamma} \quad \text{subject to: } y^{(i)} \frac{(\mathbf{w}^T \mathbf{x}^{(i)} + b)}{\|\mathbf{w}\|} \geq \gamma, \quad i = 1, \dots, N$$

Constraint ensures accurate classification

Nasty optimization problem! Let's make it look nicer!


- Let  $\gamma' = \gamma \|\mathbf{w}\|$ , this is equivalent to

$$\max_{\mathbf{w}, b, \gamma'} \frac{\gamma'}{\|\mathbf{w}\|} \quad \text{subject to: } y^{(i)}(\mathbf{w}^T \mathbf{x}^{(i)} + b) \geq \gamma', \quad i = 1, \dots, N$$

# Max margin Classifier

- Note that rescaling  $\mathbf{w}$  and  $b$  by  $(1/\gamma')$  will not change the classifier, we can thus further reformulate the optimization problem

$$\max_{\mathbf{w}, b, \gamma'} \frac{\gamma'}{||\mathbf{w}'||} \quad \text{subject to: } y^{(i)}(\mathbf{w}'^T \mathbf{x}^{(i)} + b) \geq \gamma', \quad i = 1, \dots, N$$



$$\max_{\mathbf{w}, b} \frac{1}{||\mathbf{w}'||} \quad \text{subject to: } y^{(i)}(\mathbf{w}'^T \mathbf{x}^{(i)} + b) \geq 1, \quad i = 1, \dots, N$$

- Note that maximizing the geometric margin is equivalent to minimizing the magnitude of  $\mathbf{w}$  subject to maintaining a functional margin of at least 1



# Solving the problem

$$\min_{\mathbf{w}, b} \frac{1}{2} \|\mathbf{w}\|^2$$

subject to :  $y^i (\mathbf{w} \cdot \mathbf{x}^i + b) \geq 1, \quad i = 1, \dots, N$

- This results in a **quadratic optimization problem** with *linear inequality constraints*.
- This is a well-known class of mathematical programming problems for which several (non-trivial) algorithms exist.
  - One could solve for  $\mathbf{w}$  using any of these methods
- We will see that it is useful to first formulate an equivalent dual optimization problem and solve it instead
  - This requires a bit of machinery.

# Constrained Optimization

- The general optimization problem can be written as such, for generic functions

$$\min_x f(x) \quad \text{subject to: } g_i(x) \leq 0, \quad i = 1, \dots, m$$

- To solve the above optimization problem, ***consider the following cost function known as the Lagrangian***

$$\mathcal{L}(x, \alpha) = f(x) + \sum_i \alpha_i g_i(x)$$

- Under certain conditions it can be shown that for a solution  $x'$  to the above problem, we have

$$f(x') = \min_x \max_{\alpha} \mathcal{L}(x, \alpha) = \max_{\alpha} \min_x \mathcal{L}(x, \alpha)$$

Primal Form

Dual Form subject to  $\alpha_i \geq 0$

# Dual Problem

- After simplifying the inequality, the problem becomes:

$$\min \frac{1}{2} ||\mathbf{w}||^2 \quad \text{subject to: } 1 - y^{(i)}(\mathbf{w}^T \mathbf{x}^{(i)} + b) \leq 0, \quad i = 1, \dots, N$$

- **The Lagrangian is then**

$$\mathcal{L}(\mathbf{w}, b, \alpha) = \frac{1}{2} \mathbf{w}^T \mathbf{w} + \sum_{i=1}^N \alpha_i [1 - y^{(i)}(\mathbf{w}^T \mathbf{x} + b)], \quad \text{subject to } \alpha_i \geq 0$$

- We want to solve  $\max_{\alpha} \min_x \mathcal{L}(\mathbf{w}, b, \alpha)$  s.t.  $\alpha_i \geq 0$

- **Setting the gradient of  $\mathcal{L}$  w.r.t.  $\mathbf{w}$  and  $b$  to zero, we have**

$$\mathbf{w} - \sum_{i=1}^N \alpha_i y^i \mathbf{x}^i = 0 \quad \mathbf{w} = \sum_{i=1}^N \alpha_i y^i \mathbf{x}^i$$

$$\sum_{i=1}^N \alpha_i y^i = 0$$

# Dual Problem

- If we substitute  $\mathbf{w} = \sum_{i=1}^N \alpha_i y^i \mathbf{x}^i$  in  $\mathcal{L}$ , we have

$$\begin{aligned} L(\boldsymbol{\alpha}) &= \frac{1}{2} \mathbf{w} \cdot \mathbf{w} - \sum_{i=1}^N \alpha_i \{y^i (\mathbf{w} \cdot \mathbf{x}^i + b) - 1\} \\ &= \frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N \alpha_i \alpha_j y^i y^j \langle \mathbf{x}^i \cdot \mathbf{x}^j \rangle - \sum_{i=1}^N \sum_{j=1}^N \alpha_i \alpha_j y^i y^j \langle \mathbf{x}^i \cdot \mathbf{x}^j \rangle - b \sum_{i=1}^N \alpha_i y^i + \sum_{i=1}^N \alpha_i \\ &= \sum_{i=1}^N \alpha_i - \frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N \alpha_i \alpha_j y^i y^j \langle \mathbf{x}^i \cdot \mathbf{x}^j \rangle \end{aligned}$$

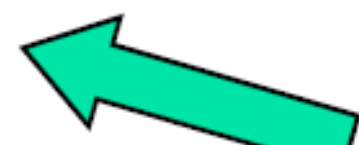

- Note that  $\sum_{i=1}^N \alpha_i y^i = 0$
- This is a function of  $\alpha_i$  only

# Dual Problem

- The new objective function is in terms of  $\alpha_i$  only. It is known as the **dual problem**: if we know all  $\alpha_i$ , then we know  $\mathbf{w}$ 
  - The original problem is known as at the primal problem
- **The objective function of the dual problem needs to be maximized!**

$$\max L(\boldsymbol{\alpha}) = \sum_{i=1}^N \alpha_i - \frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N \alpha_i \alpha_j y^i y^j < \mathbf{x}^i \cdot \mathbf{x}^j >$$

subject to  $\alpha_i \geq 0, i = 1, \dots, n,$



Properties of  $\alpha_i$  when we introduce the Lagrange multipliers

The result when we differentiate the original Lagrangian w.r.t.  $b$

# Dual Problem

$$\begin{aligned} \max L(\boldsymbol{\alpha}) &= \sum_{i=1}^N \alpha_i - \frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N \alpha_i \alpha_j y^i y^j \langle \mathbf{x}^i \cdot \mathbf{x}^j \rangle \\ \text{subject to } \quad &\alpha_i \geq 0, i = 1, \dots, n, \quad \sum_{i=1}^N \alpha_i y^i = 0 \end{aligned}$$

- This is also a **quadratic programming (QP) problem**. A global maximum of  $\alpha_i$  can always be found using a QP solver (beyond scope of this class). Luckily, SVMs have been implemented within many platforms.

- $\mathbf{w}$  can be recovered by  $\mathbf{w} = \sum_{i=1}^N \alpha_i y^i \mathbf{x}^i$
- $b$  can also be recovered as well (just a minute)



# Characteristics of the Solution

- Many of the  $\alpha_i$  are zero.  $\mathbf{w}$  is a linear combination of only a small number of data points
- In fact, optimization theory requires that the solution to satisfy the following KKT conditions:

$$\alpha_i \geq 0, i = 1, \dots, n,$$

$$y^i \left( \sum_{j=1}^N \alpha_j y^j < \mathbf{x}^j \cdot \mathbf{x}^i > + b \right) \geq 1 \quad \boxed{\text{Functional margin} \geq 1}$$

$$\alpha_i \left\{ y^i \left( \sum_{j=1}^N \alpha_j y^j < \mathbf{x}^j \cdot \mathbf{x}^i > + b \right) - 1 \right\} = 0 \quad \boxed{\alpha_i \text{ is nonzero only when functional margin} = 1}$$

- $\mathbf{x}_i$  with non-zero  $\alpha_i$  are called support vectors (SV)

- Let  $t_j, (j = 1, \dots, s)$  be the indices of the  $s$  support vectors. We can write  $\mathbf{w} = \sum_{j=1}^s \alpha_{t_j} y^{t_j} \mathbf{x}^{t_j}$
- The decision boundary is determined only by the SV

# Solving for $b$ , the bias

- Note that we know that for support vectors the functional margin = 1
- We can use this information to solve for  $b$
- We can use any support vector to achieve this

$$y^i \left( \sum_{j=1}^s \alpha_{t_j} y^{t_j} < \mathbf{x}^{t_j} \cdot \mathbf{x}^i > + b \right) = 1$$

- A numerically more stable solution is to use all support vectors (see a ML textbook)



# Classifying new examples

- For classifying with a new input  $\mathbf{z}$


Represents inner product  
between the two (e.g.  
support vector and input)

- Compute  $\mathbf{w}^T \mathbf{z} + b = \sum_{j=1}^2 \alpha_{t_j} y^{t_j} < \mathbf{x}^{t_j} \cdot \mathbf{z} > + b$
- Classify  $\mathbf{z}$  as positive if the sum is positive, and negative otherwise
- Note:  $\mathbf{w}$  need not be formed explicitly, rather we can classify  $\mathbf{z}$  by taking a weighted sum of the inner products with the support vectors
- This is useful when we generalize from inner product to kernel functions later

# Support vectors

- Only points,  $\mathbf{x}_i$ , that lie on the supporting hyperplanes have  $\alpha_i > 0$ . These are called the **support vectors**. Complexity of the solution only depends on the number of support vectors


$$\mathbf{w} = \sum_{i=1}^N y_i \alpha_i \mathbf{x}_i = \sum_{\text{support vectors}} y_i \alpha_i \mathbf{x}_i$$



Recall that  $\mathbf{w}$  is a linear combination of training data


# Support Vectors

- Learned model will not change if we delete all the data (e.g., only the support vectors matter)




# Geometric Perspective

- Maximizing margin is equivalent to maximizing the distance between the two closet points on the convex hulls of the two sets



# Geometric Perspective (2)





# Summary

- We demonstrated that we prefer to have linear classifiers with large margin.
- We formulated the problem of finding the maximum margin linear classifier as a quadratic optimization problem
- This problem can be solved by solving its dual problem, and efficient QP algorithms are available.
- Problem solved?
- How about non-linear data? – Kernels
- How about noise? – Soft Margin SVMs



# Next Class

**Support Vector Machines II: soften assumptions**