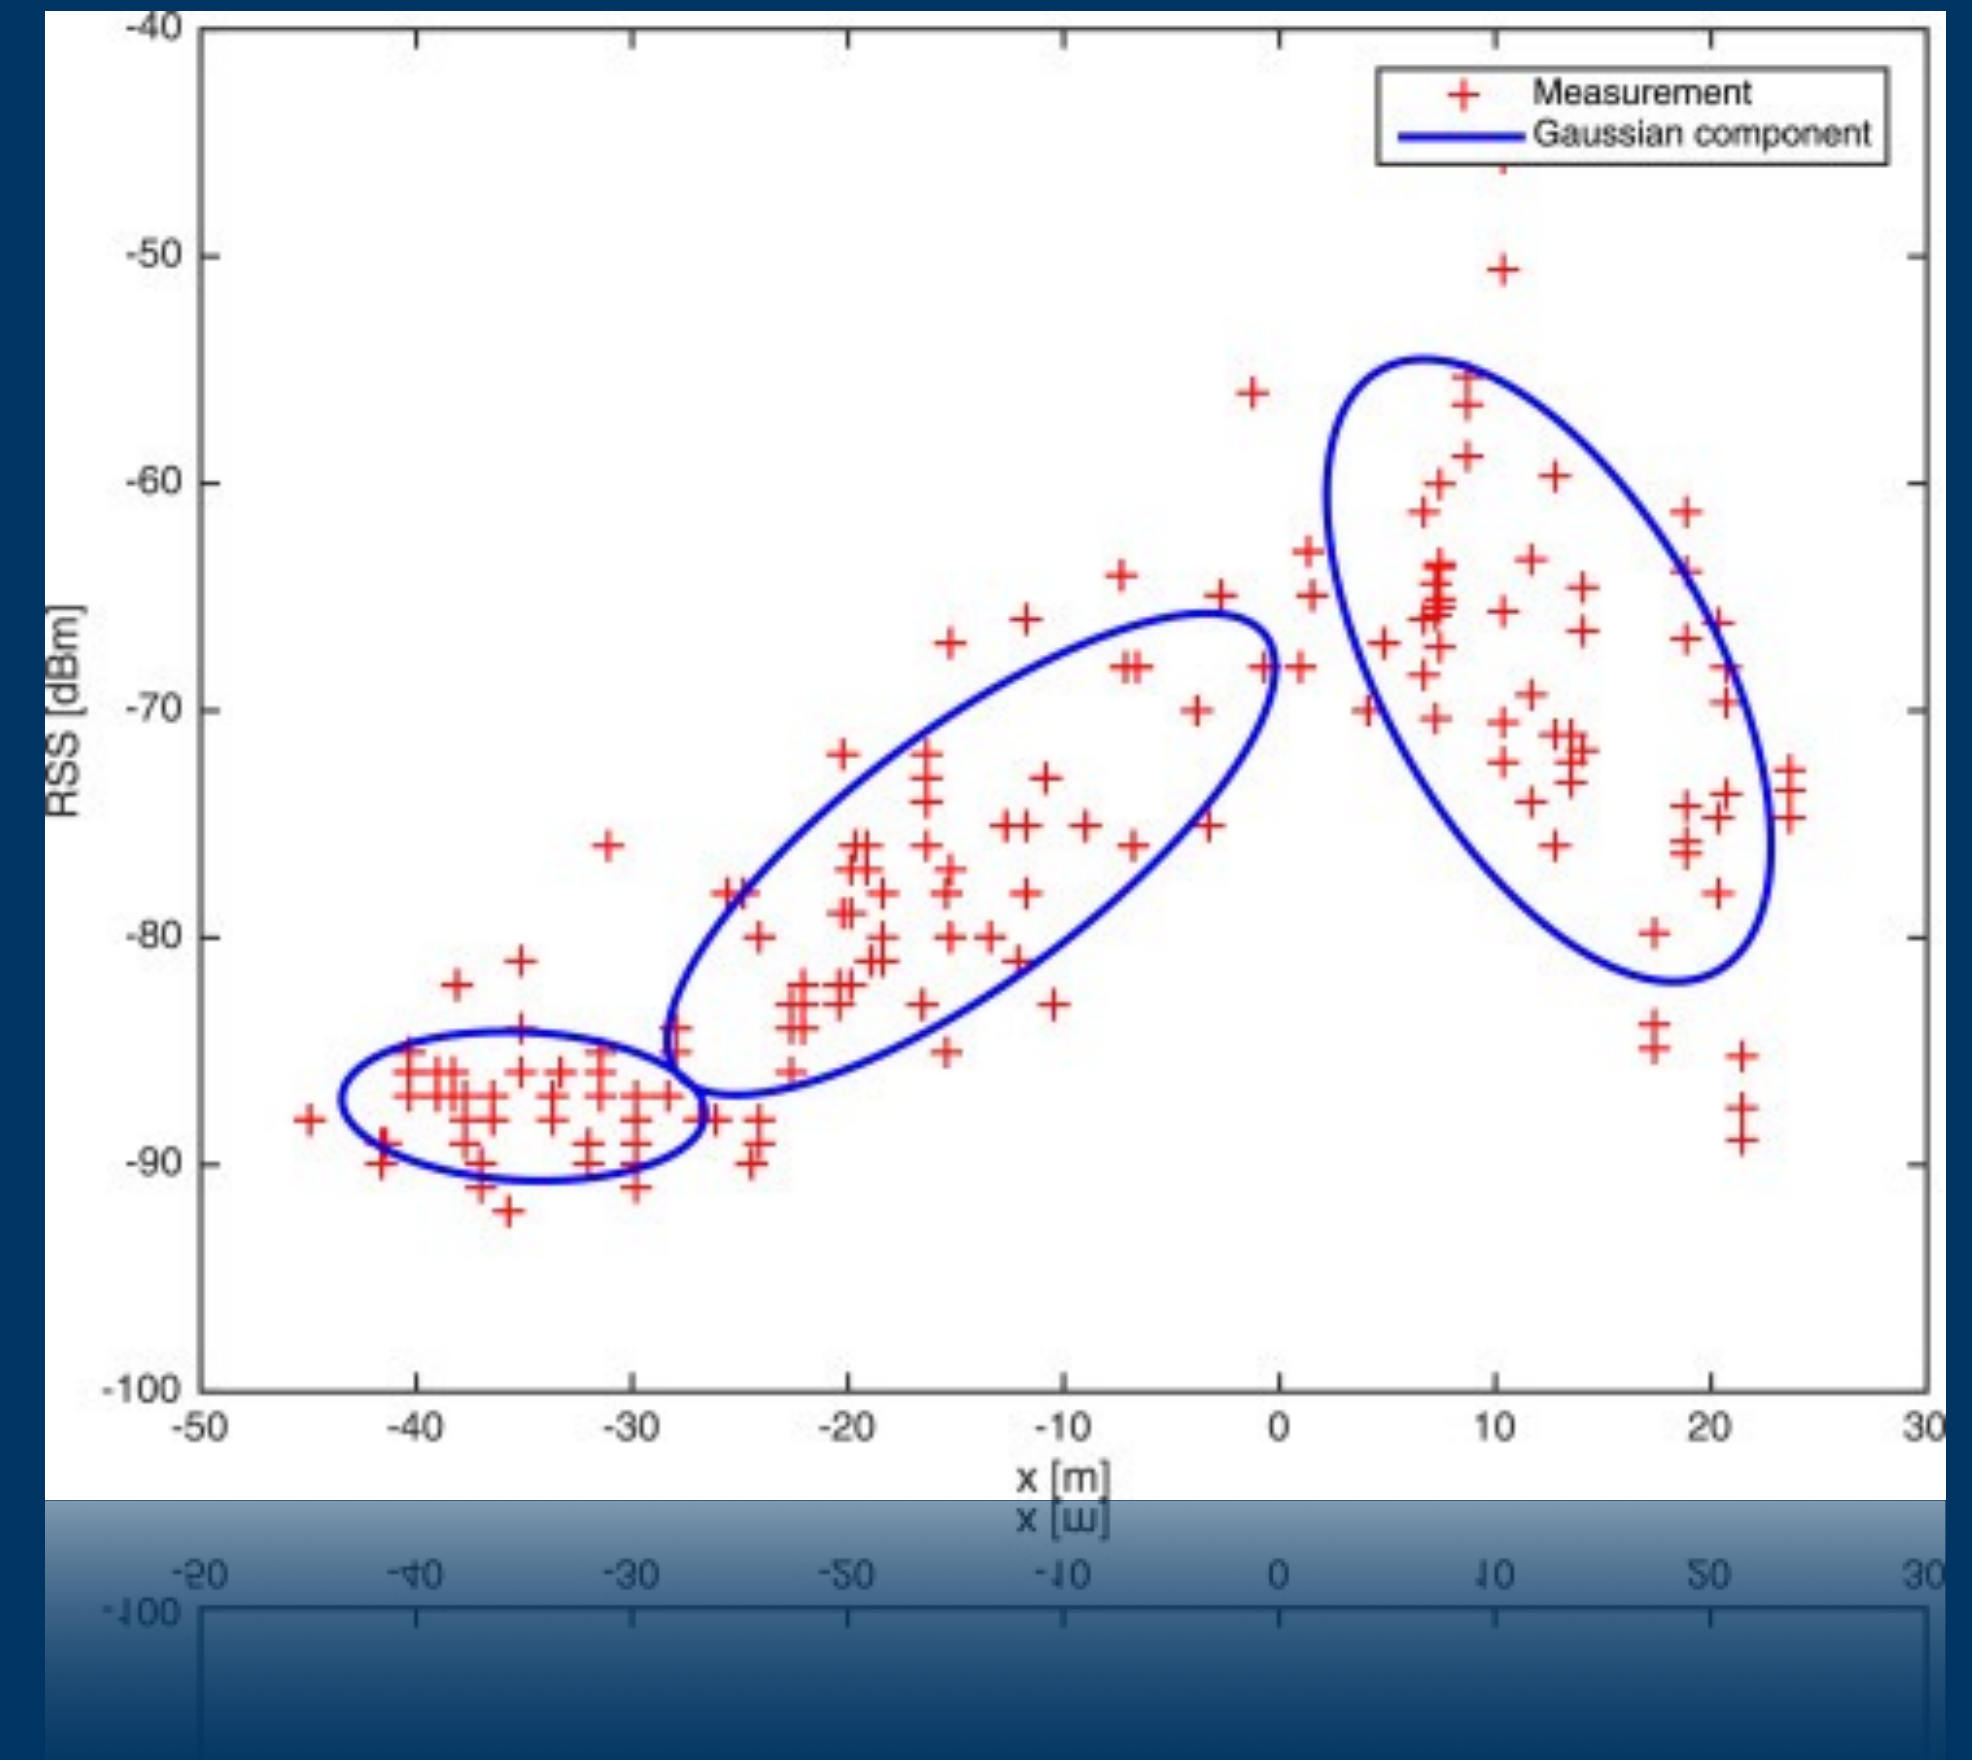


# Unsupervised Learning: Gaussian Mixture Models (GMMs)

CSCI-P556 Applied Machine Learning  
Lecture 23


D.S. Williamson



# Agenda and Learning Outcomes

## Today's Topics

- **Topics:**
  - Unsupervised Learning: Gaussian Mixture Models
- **Announcements:**
  - HW#3 Due this week
  - HW#4 Posted next week



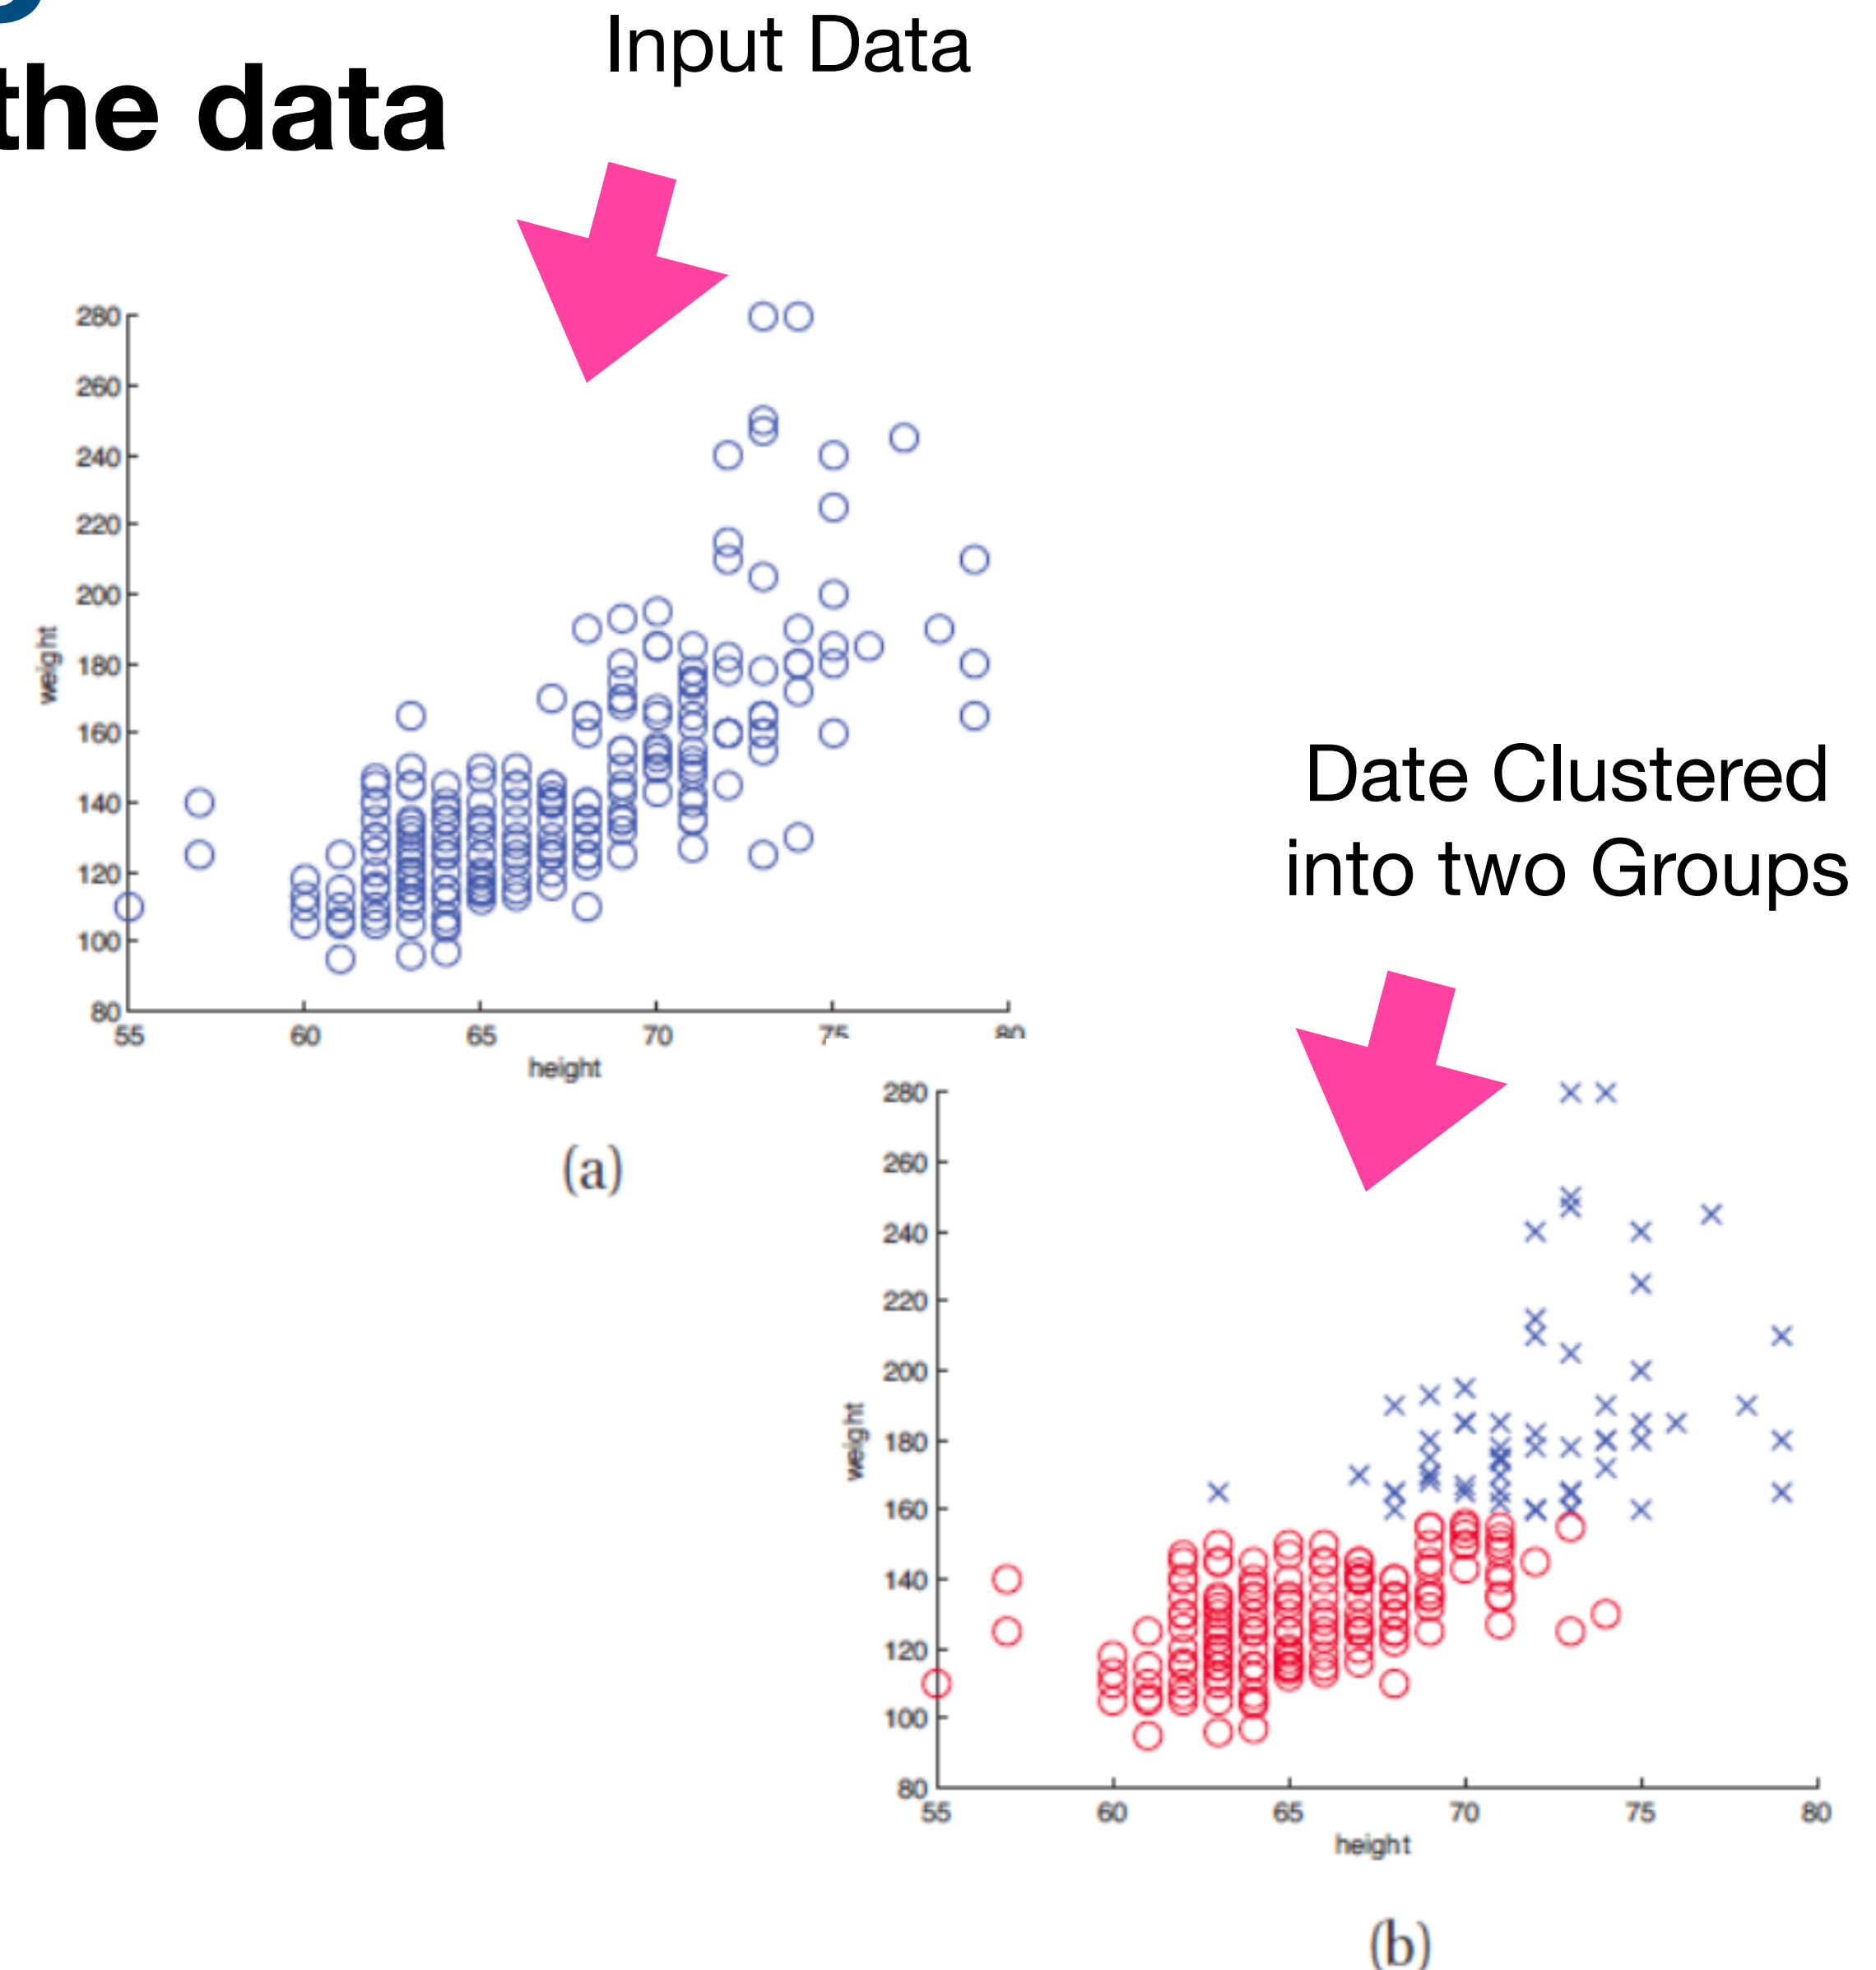
# Unsupervised Learning

Discover patterns or structure in the data

- Only have (or use) the data information (e.g. ignore labels)

$$D = \{x_1, x_2, \dots, x_i, \dots, x_N\}$$

- Examples of unsupervised learning
  - **Clustering:** K-means, vector quantization, Gaussian mixture models
  - **Dimensionality Reduction:** principal components analysis, nonnegative matrix factorization
  - **Topic Modeling:** often used in NLP



# K-Means Algorithm

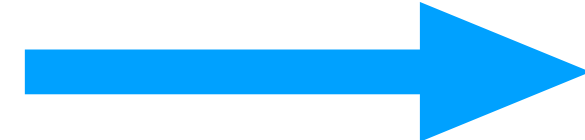
## Iterative Steps

1. Choose a set of  **$K$  cluster centers randomly** from the **input data**; or **randomly initialize**
2. Assign the  $N$  input patterns (individually) to the  $K$  clusters using the squared Euclidean distance rule.  $\mathbf{x}$  is assigned to  $C_j$  if:

$$||\mathbf{x} - \mathbf{u}_j||^2 \leq ||\mathbf{x} - \mathbf{u}_i||^2 \text{ for } i \neq j$$

3. Update cluster centers based on data that is assigned to each cluster

The number of data samples in the  $j$ -th cluster


$$\mathbf{u}_j = \frac{1}{|C_j|} \sum_{i \in C_j} \mathbf{x}_i$$

4. If any cluster center changes, go to step 2; otherwise stop. Can also specify a tolerance threshold for stopping

The K-means algorithm always converges, but the global minimum is not assured



# Hard vs. Soft Clustering

## K-Means Performs Hard Clustering

- K-Means is an example of **hard clustering**, each example in the dataset is assigned to exactly one cluster
- Hard clustering may be problematic if data points are somewhat in between the centers of multiple clusters
  - They can conceivably be in multiple clusters.
- **Soft clustering**, conversely, assigns examples to one or more clusters

# Gaussian Mixture Models

## A Soft Clustering Approach

- GMMs are a soft clustering approach that ***assign conditional probabilities for each example***.
  - Each cluster follows a Gaussian distribution with specific mean and variance
  - Each example can have a probability  $> 0$  for multiple clusters
- GMMs are also useful in better modeling probability distributions of real data (more on this later)

# Gaussian Distribution Review


## Probability Density Function

- A 1-D random variable,  $X$ , follows a Gaussian distribution if its PDF takes the following form:

$$f_X(x) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{1}{2\sigma^2}(x-\mu)^2}$$

$\mu$ : average (mean) value of  $X$

$\sigma^2$ : variance (spread) of  $X$



- Example 1-D Gaussian Random Variables:** Class grades, height, IQ, income, ...

# Gaussian Distribution Review

## Probability Density Function

- An  $N$ -D random variable,  $\mathbf{X}$ , follows a Gaussian distribution if its PDF takes the following form:

$$f_{\mathbf{X}}(\mathbf{x}) = \frac{1}{\sqrt{(2\pi)^N |\boldsymbol{\Sigma}|}} e^{-\frac{1}{2}(\mathbf{x}-\boldsymbol{\mu})^T \boldsymbol{\Sigma}^{-1}(\mathbf{x}-\boldsymbol{\mu})}$$

$\boldsymbol{\mu}$ :  $N$ -D average (mean) vector for  $\mathbf{X}$

$\boldsymbol{\Sigma}$ :  $N \times N$  co-variance matrix of  $\mathbf{X}$




# 2-D Gaussian Random Variable

## Side View

- Suppose  $\mathbf{X}$  is a 2-D Gaussian Random Variable with the mean vector and covariance matrix below. It's PDF will look like:

$$\mu = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$$

$$\Sigma = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$$




# 2-D Gaussian Random Variable

## Top-down View

- Suppose  $\mathbf{X}$  is a 2-D Gaussian Random Variable with the mean vector and covariance matrix below. It's PDF will look like:

$$\mu = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$$

$$\Sigma = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$$





# 2-D Gaussian Random Variable

Now with co-variance terms

- Suppose  $\mathbf{X}$  is a 2-D Gaussian Random Variable with the mean vector and covariance matrix below. It's PDF will look like:


$$\mu = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$$

$$\Sigma = \begin{bmatrix} 1 & 0.5 \\ 0.5 & 1 \end{bmatrix}$$



# Gaussians for Clustering


- The Gaussians can be used to model soft clustering
- Each data point has a probability of being in the cluster
- If there are multiple clusters, then multiple Gaussians will be used.





# Gaussian Mixture Models

- **Main Question:**
  - How do we learn which Gaussian each data point belongs to? **Expectation-Maximization (EM) algorithm**
- Other Questions we will answer:
  - How many Gaussians?
  - How to find means of each?
  - How to find (co-) variance of each?






# Gaussian Mixture Models

## Notation

- **Clusters are represented by  $C_i$ .** If there are  $K=4$  clusters then  $\{C_1, C_2, C_3, C_4\}$
- The conditional probability of each sample given the cluster  $P(\mathbf{X}_j | C_i)$  follows a Gaussian distribution with:
  - Mean vector:  $\mu_i$
  - Covariance matrix:  $\Sigma_i$
- Each cluster has a certain probability as well,  $P(C_i)$



# Gaussian Mixture Models


## Notation

- Our goal is to find the parameters (e.g. probabilities) that maximize the likelihood of these Gaussians accurately representing **ALL** the data

$$P(\mathbf{X}) = \sum_{j=1}^D \sum_{i=1}^K P(C_i) P(\mathbf{X}_j | C_i)$$

which Gaussian (or class)?

likelihood given Gaussian



# Expectation-maximization (EM) algorithm

- Two parts, iterated over and over again until convergence
- **Part 1: Expectation**
  - What's our best guess for every data point as to which Gaussian cluster (mixture) it comes from
  - In general, compute the conditional probability of the data point for a given cluster  $P(\mathbf{X}_j | C_i)$
- **Part 2: Maximization:**
  - Given our expectations, figure out the parameters (mean and variance) for the Gaussian distributions
  - In general, compute new parameters based on the above conditional probability

# EM: Initialization

1. Set  $P(C_i)$  to a random value, making sure the sum over all Gaussians is 1  
(e.g.  $\sum_{i=1}^K P(C_i) = 1$ )
2. Set means ( $\mu_i$ ) and variances ( $\Sigma_i$ ) for each Gaussian (e.g. cluster) to random values. This helps to get them started in the right direction.
  - With this, you can now compute the probability of  $\mathbf{X}_j$  given  $\mathbf{C}_i$ , for each class

$$P(\mathbf{X}_j | C_i) = \frac{1}{\sqrt{(2\pi)^N |\Sigma_i|}} e^{-\frac{1}{2}(\mathbf{x}-\mu_i)^T \Sigma_i^{-1}(\mathbf{x}-\mu_i)}$$

# EM: Expectation Step

## Compute Posterior Probability

- **Task #1:** for every point  $\mathbf{X}_j$ , compute the probability that cluster  $i$  (e.g.  $C_i$ ) generated that point.

From Bayes Rule

$$P(C_i | \mathbf{X}_j) = \frac{P(\mathbf{X}_j | C_i)P(C_i)}{P(\mathbf{X}_j)}$$

$$= \frac{P(\mathbf{X}_j | C_i)P(C_i)}{\sum_{i=1}^K P(\mathbf{X}_j | C_i)P(C_i)}$$

Law of Total Probability



# EM: Expectation Step

## Computer Number of Samples per Cluster

- **Task #2:** Compute the (effective) number of data samples that are assigned to each cluster,  $N_i$

$$N_i = \sum_{j=1}^D P(C_i | X_j)$$

# EM: Maximization Step

## Re-estimate the parameters

- For every cluster, compute a new class mean, variance, and prior probability

$$\mu_i = \frac{1}{N_i} \sum_{j=1}^D P(C_i | \mathbf{X}_j) \cdot \mathbf{X}_j$$

**new mean vector:** weighted average of points assigned to class  $i$

$$\Sigma_i = \frac{1}{N_i} \sum_{j=1}^D P(C_i | \mathbf{X}_j) (\mathbf{X}_j - \mu_i)(\mathbf{X}_j - \mu_i)^T$$

**new variance matrix:** calculated in same weighted manner

$$P(C_i) = \frac{N_i}{\sum_{i=1}^K N_i}$$

**new class prior:** proportion of weighted samples attributed to class

# EM: Iterations of EM algorithm


## Summary of EM

1. Initialize probabilities and parameters
2. Perform Expectation Step
  1. Compute posterior probability of cluster given each data point,  $P(C_i | \mathbf{X}_j)$
  2. Compute effective number of data points in each cluster,  $N_i$
3. Perform Maximization Step
  1. Update mean of each cluster,  $\mu_i$
  2. Update variance of each cluster,  $\Sigma_i$
  3. Update prior probability of each cluster,  $P(C_i)$
4. Return to step 2, unless convergence condition has been satisfied

# Graphical Depiction of EM

## Clustering Data using Two Gaussians


- 2-D Data points shown in green
- Initial 2-D Gaussians, one in Blue, one in Red
- Showing Gaussian contours out to one standard deviation



# Graphical Depiction of EM

## Clustering Data using Two Gaussians

- Result after the initial E-step
  - Compute posterior probability of cluster given each data point,  $P(C_i | \mathbf{X}_j)$
  - Color indicates probability that it was generated by the cluster
  - Points that could have come from either cluster are shown in Purple






# Graphical Depiction of EM

## Clustering Data using Two Gaussians


- Result after the initial M-step
  - For every cluster, compute a new class mean, variance, and prior probability
- The clusters have moved and better represent their data



# Graphical Depiction of EM

## Clustering Data using Two Gaussians


- Result after the completing the 2nd E and M steps



# Graphical Depiction of EM

## Clustering Data using Two Gaussians


- Result after the completing the 5th E and M steps



# Graphical Depiction of EM

## Clustering Data using Two Gaussians

- Result after the completing the 20th E and M steps



# EM Performance Evaluation

## How to judge performance


- The log likelihood is used to evaluate performance
- For a set of data  $\mathbf{X}$ , we define the log likelihood given as (assuming independence):

$$\ln P(\mathbf{X}) = \sum_{j=1}^D \ln \left[ \sum_{i=1}^K P(C_i) P(\mathbf{X}_j | C_i) \right]$$



# EM Performance Evaluation

- **Key points:**
  - EM algorithm monotonically increases the data log likelihood until it reaches a local optimum
  - **If you do not observe monotonic increase, then you must have an error in your math and/or code!**
  - In practice, the algorithm converges when the change in log likelihood (or the parameters) falls below some threshold

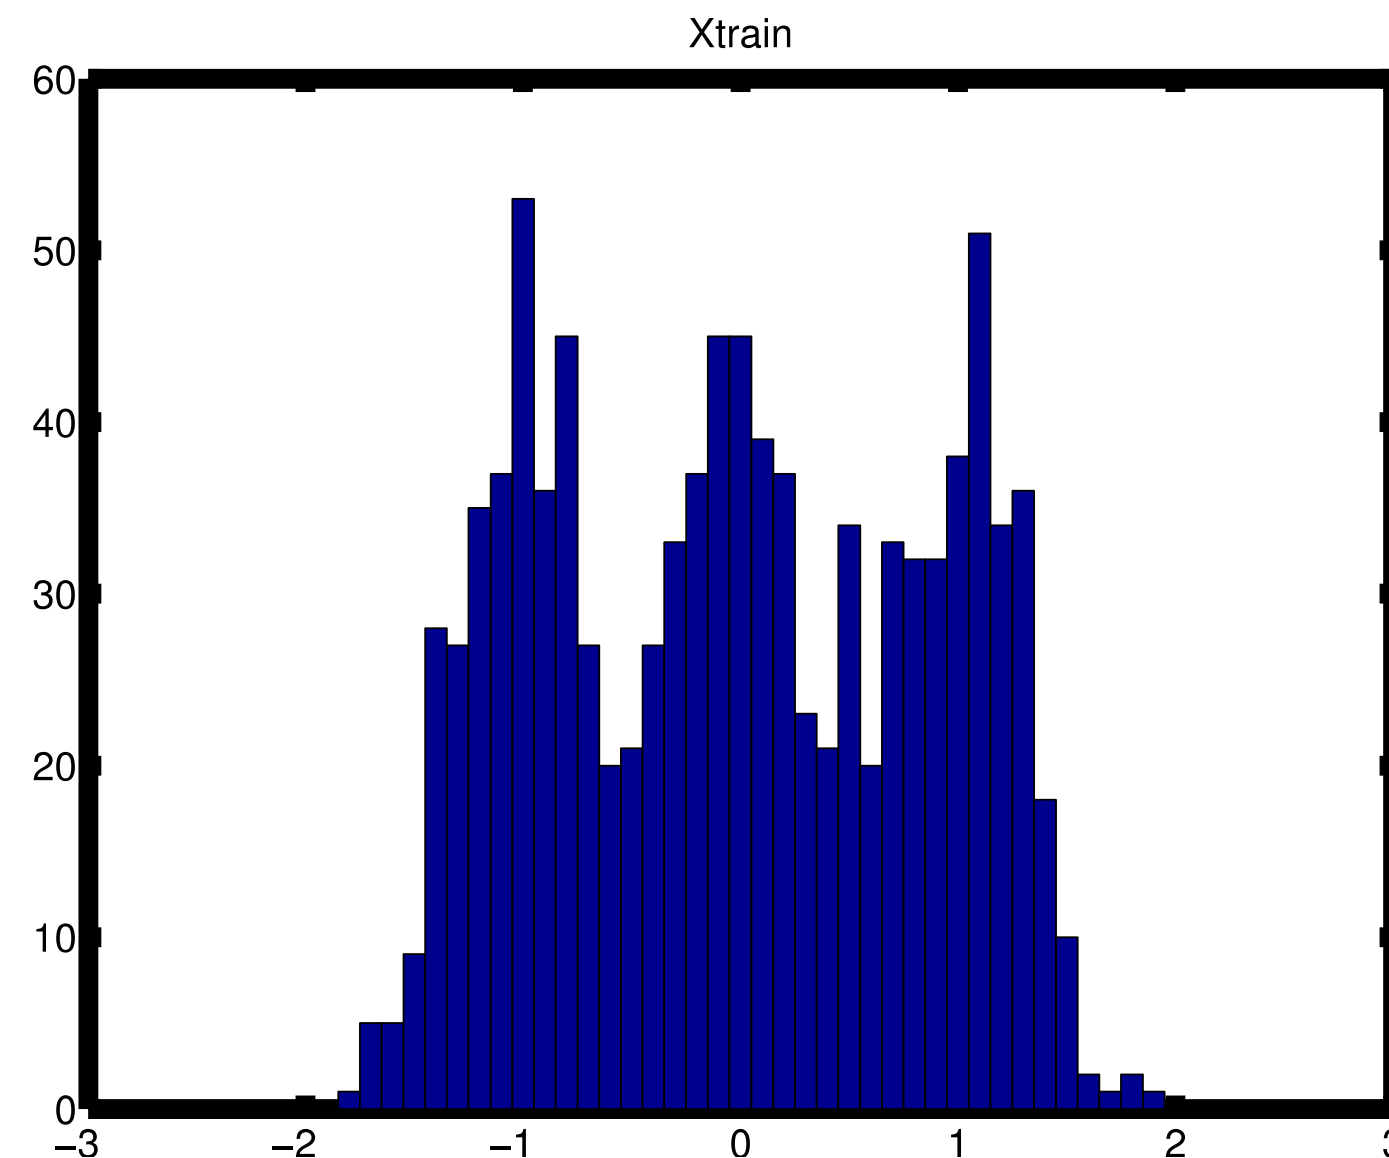


# EM Design Decision


## How do you determine how many Gaussians to use?

- Simple approach:
  - Try different number of Gaussians and pick the one with the largest (log) likelihood (or smallest negative log likelihood)
    - Pros: Easy to do
    - Cons: slow and time-consuming
  - Alternatively, use the **elbow method** as before for *K-Means*

True “Distribution”



Modeled “Distributions”



# Gaussian Mixture Model in Python

## `sklearn.mixture.GaussianMixture`

```
class sklearn.mixture.GaussianMixture(n_components=1, *, covariance_type='full', tol=0.001, reg_covar=1e-06,
max_iter=100, n_init=1, init_params='kmeans', weights_init=None, means_init=None, precisions_init=None,
random_state=None, warm_start=False, verbose=0, verbose_interval=10)
```

[\[source\]](#)

### Parameters:

**`n_components` : int, default=1**

The number of mixture components.

**`covariance_type` : {'full', 'tied', 'diag', 'spherical'}, default='full'**

String describing the type of covariance parameters to use. Must be one of:

**'full'**

each component has its own general covariance matrix

**'tied'**

all components share the same general covariance matrix

**'diag'**

each component has its own diagonal covariance matrix

**'spherical'**

each component has its own single variance

**`tol` : float, default=1e-3**

The convergence threshold. EM iterations will stop when the lower bound average gain is below this threshold.

**`reg_covar` : float, default=1e-6**

Non-negative regularization added to the diagonal of covariance. Allows to assure that the covariance matrices are all positive.

**`max_iter` : int, default=100**

```
>>> import numpy as np
>>> from sklearn.mixture import GaussianMixture
>>> X = np.array([[1, 2], [1, 4], [1, 0], [10, 2], [10, 4], [10, 0]])
>>> gm = GaussianMixture(n_components=2, random_state=0).fit(X)
>>> gm.means_
array([[10.,  2.],
       [ 1.,  2.]])
```

# **Next Class: Unsupervised Learning**

## **Dimensionality Reduction**