

Operating Systems

Xinu – Hardware and Runtime Environment

1

Processors and Systems

- The textbook describes two architectures – The ARM-based BeagleBone Black and the Intel Galileo
 - This covers both the ARM and x86 Instruction Set Architectures (ISA)
 - ARM is a Reduced Instruction Set Computing (RISC) ISA, whereas x86 is referred to as a Complex Instruction Set Computing (CISC) ISA
 - The previous edition of the book describes the Linksys E2100L home router which uses the MIPS ISA
- We will focus on the BeagleBone, which is based on a system on a chip (SoC) from TI, the AM335x Sitara
 - The Raspberry PI system is based on a Broadcom BCM2835 which includes an ARM1176JZF-S processor

2

Registers

- High speed memory used by the processor
- The compiler generates code to load (fetch) and store to memory, and operates on values in registers
- Some are general purpose, and some have a purpose (by convention only in some cases)
- Hold part of the state of the computation and thus must be saved and restored when switching processes

3

ARM Processor - Registers

- ARM registers are banked (or windowed) such that some registers are renamed when the mode changes (in e.g. interrupts)
 - r0 – r6 are visible across modes

Name	Use	Notes
r0-r3	General Purpose	Argument registers for 1 st four args
r4-r12	General Purpose	r4-r11 are "callee-save"; r12 = ip
r13	Stack Pointer	by procedure calling convention
r14	Link Register	by procedure calling convention
r15	Program Counter	
cpsr	Current Program Status Register	spsr is the Saved Program Status Register

4

ARM Program Status Register (CPSR)

- The CPSR includes
 - Processor mode (0-4)
 - Interrupt disable (7)
 - Fast interrupt disable (6)
 - Condition code flags
 - Negative result from CPU logic (31)
 - Zero result from CPU logic (30)
 - ALU operation Carried out (29)
 - ALU operation oVerflowed (28)

5

Intel Registers

Name	Use
EAX	Accumulator
EBX	Base
ECX	Count
EDX	Data
ESI	Source Index
EDI	Destination Index
EBP	Base Pointer
ESP	Stack Pointer

Also a Program Counter (PC) register called the Instruction Pointer (RIP)


6

The Bus

- The bus is for communication with peripherals and memory
 - Load (Fetch) / Store operation
- The CPU places an address on the bus, and signals ready
- The appropriate target (e.g. memory bank) retrieves the value from the requested address and puts it on the bus
- In memory-mapped I/O – devices appear as “memory” with particular addresses
 - reading and writing to a device uses fetch/store in the same way


7

System Overview



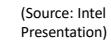
8

- Device Controllers are connected or integrated in various locations
 - PCI-e
 - PCI
- Location determines the speed of access to memory



More Modern Architecture (NUMA)

- No longer a single memory controller
- Symmetric Multiprocessor with Non-uniform Memory Access (NUMA)



BCM2835 (Raspberry PI)



ISA Implications for the OS

- The compiler generates machine code from C language source
- There are a few places that we need to consider the system details
 - Hardware Abstraction Layer (HAL) in Windows
- Mostly in startup, context switching and interrupt handling as we will see

12

Stack: Function Calling Conventions

- Function invocation is central to OS operation
- Calling conventions refers to what happens when a function is called and when it returns
 - Conventions in that it is really decided by the compiler, and it simply needs to be consistent
- The run-time stack stores state associated with a function call
 - with addresses growing downward (generally)
- There is space in this stack frame (or activation record) that includes
 - local variables
 - return address
 - arguments
 - return values

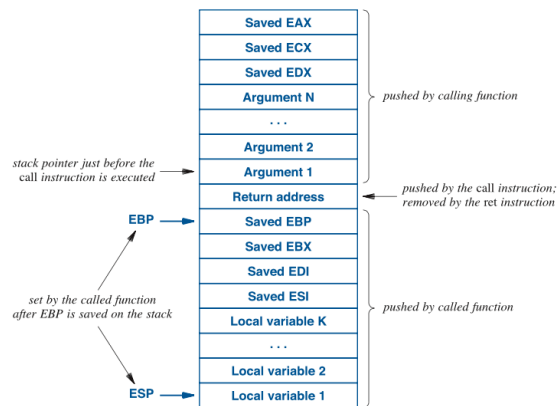
13

Stack: Function Calling

- Arguments / Parameters
 - Fixed number of arguments are passed in registers (r0-r3) and the remainder are passed via memory on the stack
- The compiler computes the size needed for the stack frame
 - This is why you need a function prototype accessible when the calling code is compiled
- Argument save area in stack frame for function X is used by function Y

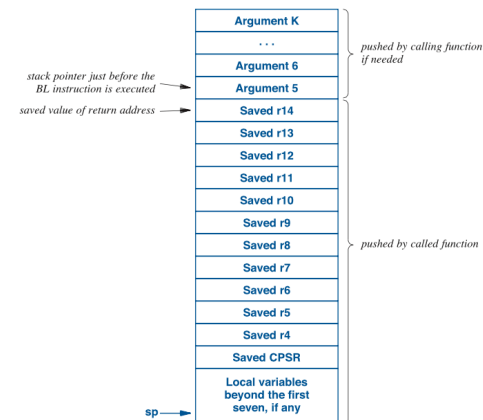
14

x86 Stack Layout



15

ARM Stack



16

Interrupts

- As discussed, I/O devices interrupt the system
 - Similar handling for other *exceptions* like divide by zero or page fault
- I/O interrupts are generated by placing a signal on the bus (control lines)
 - Monitored by the processor or a co-processor
- When an interrupt occurs
 - Record the address of the next instruction
 - Jump to memory location 0x18
 - Set a control bit in cpsr to disable further interrupts

17

Interrupt Processing

- The interrupt controller has registers for interrupt masks
 - Controlling what interrupts can and cannot occur
 - There is also a bit to disable all of them
 - Xinu uses this when manipulating internal structures
 - Functions `disable()` and `restore()` to save and restore the mask
- At startup, an instruction must be stored at 0x18 (on ARM) for servicing interrupts

18

Timer Hardware

- Timer device works similarly to I/O
 - When the timer expires, it generates an interrupt
 - Once the timer has been set, the OS must be ready to handle interrupts
- The timer is programmable with two registers
 - Counter – initial value
 - Limit – the value at which to interrupt (the length of time)
- An alternative is a real-time clock that interrupts with a fixed period (60x / second)

19

Direct Memory Access - DMA

- Hardware devices communicating directly with memory
 - Rather than the processor reading from the bus, then writing to the bus...
 - The I/O device can perform a series of operations while the processor runs code
- Network device example
 - The OS allocates a list of packet buffers and starts the network device
 - Packet arrives, device transfers the complete packet into memory, then interrupts

20

Serial Device

- Serial devices are the simplest I/O device
- Universal Asynchronous Receiver Transmitter (UART)
- Handles both input and output
 - Called transmit (TX) and receive (RX)
- When an interrupt occurs, the system must examine a device register
 - Byte(s) received
 - Byte(s) transmission complete

21

Memory Layout

- When a program is compiled the compiler generates partitions of memory segments

- Text segment
- Data segment
- BSS segment
- Stack Segment



- As the program creates new processes:



22