

Operating Systems

Xinu – Internals
List and Queue Manipulation

1

Linked Lists in the OS

- Manipulating lists of processes is an important operation in the OS
 - A process's lifecycle consists of moving between, and in, queues and lists
- Xinu implements a unified approach to list management
 - All list management uses this common infrastructure
 - Common functions to create a new list, insert an element at the end of the list, insert or remove from the middle, remove an item from the front
- We can think of this code as being sequential
 - Protected with mutual exclusion as necessary

2

Unified Lists in Xinu

- The process manager handles processes
 - A process moves among the lists frequently
 - At any time, a process is only in one list
- Rather than store all the information about a process, the process manager is free to store only the process ID (PID) or Thread ID (TID) in a list
 - So when we refer to putting a process on a list, it really means putting the PID there – the process control block need not move
- Unified implementation means that not every subsystem uses all the list features


3

List Properties

- All lists are doubly linked – each node points to predecessor and successor
- Each node stores a key as well as a process ID, even though a key is not used in a FIFO list
- Each list has head and tail nodes; the head and tail nodes have the same memory layout as other nodes
- Non-FIFO lists are ordered in descending order; the key in the head node is greater than the maximum valid key value, and the key value in the tail node is less than the minimum valid key


4

A doubly-linked List



5

An Empty List



6

A compact list structure

- Compact memory use is important
 - in general, but in embedded systems in particular
- Two ideas related to process specific representation
- Relative pointers
 - Given that there is some (small) fixed number of processes (NPROC)
 - One might use a pointer in this situation, which is 4 bytes
 - For NPROCS < 62, we only need 6 bits
 - Allocate the nodes in a contiguous array and use the array index as a "pointer"

7

A compact list structure

- Implicit data structure – based on the fact that a process can only be in one list, we can use the list position to indicate the ID
- To omit the PID, use an array and use the *i*th element of the array for process *i*
- To put process 3 in a linked list, insert node 3 into the list
- The relative address of a node is the same as the ID of a process being stored
 - List membership is implicit

8

Illustration of implicit identifiers

	KEY	PREV	NEXT	
0				
1				
2	14	4	61	
3				
4	25	60	2	
5				
...				
NPROC-1				
...				
60	MAXKEY	-	4	
61	MINKEY	2	-	
...				

Each row corresponds to a single process

Conceptual boundary

Pairs of rows form the head and tail of lists

9

Notes

- The table contains NQENT entries
- There is an implicit divider at NPROCS-1
- Between queue[tab[NPROC]] and queue[tab[NQENT-1]] are the heads and tails of lists
- The default $NPROC + 4 + NSEM + NSEM$ allocates enough space for
 - each process
 - head and tail entries for the ready and sleep lists
 - a head and tail entry for each of the NSEM semaphores in the system
- This can be changed at compile time

10

Implementation of the Queue Data Structure

```

/* queue.h - firstid, firstkey, isempty, lastkey, nonempty */

/* Queue structure declarations, constants, and inline functions */

/* Default # of queue entries: 1 per process plus 2 for ready list plus*/
/*      2 for sleep list plus 2 per semaphore */

#ifndef NQENT
#define NQENT (NPROC + 4 + NSEM + NSEM)
#endif

#define EMPTY (-1) /* null value for qnext or qprev index */
#define MAXKEY 0x7FFFFFFF /* max key that can be stored in queue */
#define MINKEY 0x80000000 /* min key that can be stored in queue */

```

11

```

struct gentry { /* one per process plus two per list */
    int32  qkey; /* key on which the queue is ordered */
    qid16  qnext; /* index of next process or tail */
    qid16  qprev; /* index of previous process or head */
};

extern struct gentry queue[tab[]];

/* Inline queue manipulation functions */

#define queuehead(q) (q)
#define queuetail(q) ((q) + 1)
#define firstid(q) (queue[tab[queuehead(q)]] .qnext)
#define lastid(q) (queue[tab[queuetail(q)]] .qprev)
#define isempty(q) (firstid(q) >= NPROC)
#define nonempty(q) (firstid(q) < NPROC)
#define firstkey(q) (queue[tab[firstid(q)]] .qkey)
#define lastkey(q) (queue[tab[lastid(q)]] .qkey)

```

12

```

/* assumes interrupts are disabled */

/* Inline to check queue id */

#define isbadqid(x) (((int32)(x) < 0) || (int32)(x) >= NQENT-1)

/* Queue function prototypes */

pid32 getfirst(qid16);
pid32 getlast(qid16);
pid32 getitem(pid32);
pid32 enqueue(pid32, qid16);
pid32 dequeue(qid16);
status insert(pid32, qid16, int);
status insertd(pid32, qid16, int);
qid16 newqueue(void);

```

13

Basic Functions to Extract A Process From A List

```

/* getitem.c - getfirst, getlast, getitem */
#include <xinu.h>
/*----- *
getfirst - Remove a process from the front of a queue
*----- */

pid32 getfirst(
    qid16 q /* ID of queue from which to */
) /* remove a process (assumed */
/* valid with no check) */

{
    pid32 head;

    if (isempty(q)) {
        return EMPTY;
    }

    head = queuehead(q);
    return getitem(queuestab[head].qnext);
}

```

14

```

/*----- *
* getlast - Remove a process from end of queue
*----- */

pid32 getlast(
    qid16 q /* ID of queue from which to */
) /* remove a process (assumed */
/* valid with no check) */

{
    pid32 tail;

    if (isempty(q)) {
        return EMPTY;
    }

    tail = queuetail(q);
    return getitem(queuestab[tail].qprev);
}

```

15

```

/*----- *
* getitem - Remove a process from an arbitrary point in a queue
*----- */

pid32 getitem(
    pid32 pid /* ID of process to remove*/
)

{
    pid32 prev, next;

    next = queuestab[pid].qnext; /* following node in list */
    prev = queuestab[pid].qprev; /* previous node in list */
    queuestab[prev].qnext = next;
    queuestab[next].qprev = prev;
    return pid;
}

```

16

FIFO Queue Manipulation

```

/* queue.c - enqueue, dequeue */

#include <xinu.h>

struct gentry queuetab[NQENT]; /* table of process queues*/

/*-----
 * enqueue - Insert a process at the tail of a queue
 *-----
 */
pid32 enqueue(
    pid32 pid, /* ID of process to insert*/
    qid16 q /* ID of queue to use */
)

```

17

```

{
    int tail, prev; /* tail & previous node indexes */

    if (isbadqid(q) || isbadpid(pid)) {
        return SYSERR;
    }

    tail = queuetail(q);
    prev = queuetab[tail].qprev;

    queuetab[pid].qnext = tail; /* insert just before tail node */
    queuetab[pid].qprev = prev;
    queuetab[prev].qnext = pid;
    queuetab[tail].qprev = pid;
    return pid;
}

```

18

```

/*-----
 * dequeue - Remove and return the first process on a list
 *-----
 */
pid32 dequeue(
    qid16 q /* ID queue to use */
)
{
    pid32 pid; /* ID of process removed */

    if (isbadqid(q)) {
        return SYSERR;
    } else if (isempty(q)) {
        return EMPTY;
    }

    pid = getfirst(q);
    queuetab[pid].qprev = EMPTY;
    queuetab[pid].qnext = EMPTY;
    return pid;
}

```

19

Manipulation of Priority Queues

```

/* insert.c - insert */

#include <xinu.h>

/*-----
 * insert - Insert a process into a queue in descending key order
 *-----
 */
status insert(
    pid32 pid, /* ID of process to insert*/
    qid16 q, /* ID of queue to use */
    int32 key /* key for the inserted process */
)

```

20

```

{
    int16    curr;          /* runs through items in a queue*/
    int16    prev;          /* holds previous node index */

    if (isbadqid(q) || isbadpid(pid)) {
        return SYSERR;
    }

    curr = firstid(q);
    while (queuetab[curr].qkey >= key) {
        curr = queuetab[curr].qnext;
    }
    /* insert process between curr node and previous node */

    prev = queuetab[curr].qprev; /* get index of previous node */
    queuetab[pid].qnext = curr;
    queuetab[pid].qprev = prev;
    queuetab[pid].qkey = key;
    queuetab[prev].qnext = pid;
    queuetab[curr].qprev = pid;
    return OK;
}

```

21

List Initialization

```

/* newqueue.c - newqueue */

#include <xinu.h>

/*-----
 * newqueue - Allocate and initialize a queue in the global queue
 *            table
 *-----
 */
qid16 newqueue(void)
{
    static qid16    nextqid=NPROC; /* next list in queuetab to use */
    qid16          q;             /* ID of allocated queue */

    q = nextqid;
    if (q > NQENT) { /* check for table overflow */
        return SYSERR;
    }
}

```

22

```

nextqid += 2; /* increment index for next call*/

/* initialize head and tail nodes to form an empty queue */

queuetab[queuehead(q)].qnext = queuetail(q);
queuetab[queuehead(q)].qprev = EMPTY;
queuetab[queuehead(q)].qkey = MAXKEY;
queuetab[queuetail(q)].qnext = EMPTY;
queuetab[queuetail(q)].qprev = queuehead(q);
queuetab[queuetail(q)].qkey = MINKEY;
return q;
}

```

23