

Operating Systems

Context Switching (and Scheduling)

1

Context Switching Implementation

- Registers can't be manipulated directly in C, so *resched* calls an assembly language function called *ctxsw*
 - This function is machine dependent
- The last step involves the program counter as Xinu must restore all other elements of program state before jumping to the new process
- Some architectures provide single instructions for (atomically) saving and restoring registers
 - The RISC model often requires a series of instructions – saving each register explicitly
 - ARM has push and pop macros


2

Saving State in Memory

- We have said that elements of state need to be stored on a per-process basis
- A process table entry is per-process but each process also has a stack
 - In use by the currently executing process
 - When a function is called, the executing process makes space on the stack for local variables and arguments
 - When it returns, they are popped off
 - Per-process even if processes are otherwise identical
- Xinu processes execute in this way already, and the call to *ctxsw* can use the same mechanism

3

Picture of Memory with 3 processes



4

ctxsw

- Xinu saves the process's register state on the process's stack
- Written in assembly
 - Rewritten for each architecture
- Thus *ctxsw*:
 - Executes instructions that push the register contents onto the stack of the running process
 - Saves the stack pointer in the process table entry for the current process, and loads the SP of the next process
 - Executes instructions to reload the processor registers from the values previously saved for the next process
 - Jumps to the location in the new process at which execution should resume

5

Context switch on ARM

- ARM has `push` and `pop` macros to copy multiple registers to the stack
 - On ARM (and MIPS) each instruction can store only one register; N instructions for N registers
- A “coprocessor” stores the status register, instruction `mrs` copies it to a general-purpose register
- `r0 – r3` are caller-save, and `ctxsw` only uses 2 arguments so `r2` and `r3` are available
 - No need to save an extra register as on x86

6

```
/* ctxsw.S - ctxsw (for ARM) */

.text
.globl ctxsw

/*-----
 * ctxsw - ARM context switch; the call is ctxsw(&old_sp, &new_sp)
 *-----
 */

ctxsw:
    push    {r0-r11, lr}      /* Push regs 0 - 11 and lr */
    push    {lr}              /* Push return address */
    mrs     r2, cpsr           /* Obtain status from coprocess. */
    push    {r2}              /* and push onto stack */
    str     sp, [r0]           /* Save old process's SP */
    ldr     sp, [r1]           /* Pick up new process's SP */
    pop     {r0}              /* Use status as argument and */
    bl      restore           /* call restore to restore it */
    pop     {lr}              /* Pick up the return address */
    pop     {r0-r12}          /* Restore other registers */
    mov     pc, r12           /* Return to the new process */
```

7


Switching and the Program Counter

- What we can observe is that a process calling `resched` and `ctxsw` will be suspended at that point – in `ctxsw`
- Saving and restoring the PC is a special case
 - It is changing as the instructions to save the registers are executed
 - It can specifically calculated
- When a process restarts it should begin executing in *resched* immediately following the call to `ctxsw`
- All processes call *resched* and it calls `ctxsw` so at context switch time, they are all frozen in the same place
 - But the stack above it is distinct

8

Concurrent execution

- The concurrent execution model means that a process must always be executing
 - and eventually execute the scheduler
- The scheduler's only function is to switch the process from one function to another



9

The Null process

- There are times when no regular process is ready to execute
 - blocked on a semaphore, waiting for I/O, etc.
- The scheduler needs a process to execute!
- So at least one process must remain to be switched to
- This is the Null process
 - PID 0
 - Priority 0
 - Infinite loop

10

Making a Process Ready

- When resched moves the current process onto the ready list, it does it directly
- Making a process ready is a frequent activity so there is a function
- Args: PID and a boolean to indicate whether resched should be called
 - RESCHED_{YES/NO}
- The scheduling policy says that the highest priority process is executing at any time
 - Scheduling invariant: a function assumes that the highest priority ready process was executing when it was called and must insure that is also true when it returns
 - If the function alters process state, then it must insure the invariant
- *ready* is an exception in that it may move multiple processes to the ready queue at once
 - Thus, RESCHED_NO

11

```

/* ready.c - ready */

#include <xinu.h>

qid16  readylist;      /* index of ready list */

/*-----
 * ready - Make a process eligible for CPU service
 *-----
 */
status ready(
    pid32  pid,        /* ID of process to make ready */
    bool8   resch       /* reschedule afterward? */
)
{
    register struct procent *prptr;

    if (isbadpid(pid)) {
        return(SYSERR);
    }

    /* Set process state to indicate ready and add to ready list */

    prptr = &proctab[pid];
    prptr->prstate = PR_READY;
    insert(pid, readylist, prptr->prprio);
}
    
```

12

```

if (resched == RESCHED_YES) {
    resched();
}
return OK;
}

```

13

Deferred Rescheduling

- Resched uses Defer.ndefers to determine if rescheduling is deferred
 - sched_cntl(DEFER_START); /* defers, and */
 - sched_cntl(DEFER_STOP); /* ends deferral */
- It is a counter as multiple processes may request deferral
- When it is 0, then no process has requested deferral
- The Defer.attempt flag indicates whether it was attempted so that resched can be called

14

```

/* sched.h */

/* Constants and variables related to deferred rescheduling */

#define DEFER_START1 /* start deferred rescheduling */
#define DEFER_STOP 2 /* stop deferred rescheduling */

/* Structure that collects items related to deferred rescheduling */

struct defer {
    int32 ndefers; /* number of outstanding defers */
    bool8 attempt; /* was resched called during the */
                /* deferral period? */
};

extern struct defer Defer;

```

15

```

/* sched_cntl.c - sched_cntl */

#include <xinu.h>

struct defer Defer;

/*-----
 * sched_cntl - control whether rescheduling is deferred or allowed
 *-----
 */

status sched_cntl( /* assumes interrupts are disabled */
    int32 def /* either DEFER_START or DEFER_STOP */
)
{
    switch (def) {

        /* Process request to defer:
         * 1) Increment count of outstanding deferral requests */
        /* 2) If this is the start of deferral, initialize Boolean */
        /* to indicate whether resched has been called */
    }
}

```

16

```

case DEFER_START:
    if (Defer.ndefers++ == 0) { /* increment deferrals */
        Defer.attempt = FALSE; /* no attempts so far */
    }
    return OK;

    /* Process request to stop deferring: */
    /* 1) Decrement count of outstanding deferral requests */
    /* 2) If last deferral ends, make up for any calls to */
    /*    resched that were missed during the deferral */

case DEFER_STOP:
    if (Defer.ndefers <= 0) { /* none outstanding */
        return SYSERR;
    }
    if (--Defer.ndefers == 0) { /* end deferral period */
        if (Defer.attempt) { /* resched was called */
            resched(); /* during deferral */
        }
    }
    return OK;

default:
    return SYSERR;
}
}

```

17

Context Switch on x86

- Must save registers for the old process on its stack
- x86 has a dedicated instruction to do this – `pushal`
– and `popal` does the inverse
- First save EBX as it is needed to access the arguments
- Next, processor flags and general purpose registers
- Save the old process's stack pointer in the location pointed to by argument 1
- Load the new process's stack pointer in the location pointed to by argument 2

18

```

/* ctxsw.S - ctxsw (for x86) */

        .text
        .globl ctxsw

/*-----
 * ctxsw - X86 context switch; the call is ctxsw(&old_sp, &new_sp)
 *-----
 */
ctxsw:
        pushl    %ebp          /* Push ebp onto stack */
        movl     %esp,%ebp     /* Record current SP in ebp */
        pushfl   /* Push flags onto the stack */
        pushal   /* Push general regs. on stack */

        /* Save old segment registers here, if multiple allowed */

```

19

```

        movl     8(%ebp),%eax   /* Get mem location in which to */
                                /* save the old process's SP */
        movl     %esp,(%eax)    /* Save old process's SP */
        movl     12(%ebp),%eax  /* Get location from which to */
                                /* restore new process's SP */

        /* The next instruction switches from the old process's */
        /* stack to the new process's stack. */

        movl     (%eax),%esp    /* Pop up new process's SP */

        /* Restore new seg. registers here, if multiple allowed */

        popal    /* Restore general registers */
        movl     4(%esp),%ebp   /* Pick up ebp before restoring */
                                /* interrupts */
        popfl    /* Restore interrupt mask */
        add      $4,%esp        /* Skip saved value of ebp */
        ret      /* Return to new process */

```

20