

Operating Systems

Xinu - Process Management


1

Suspending and Resuming

- Sometimes OS functions temporarily stop a process (suspend)
 - Waiting for one of a set of conditions to become true before resuming
 - Cannot be in *ready* state – new state: *suspended*
- Operations:
 - suspend (make ineligible to use the CPU)
 - resume (make eligible to run again)
- Augmented state transition diagram

2

Process State Transition Diagram



3

Suspending and Implementation

- One process can suspend another process, or a process may suspend itself
 - The function takes a PID as an argument: *suspend(PID)*;
- The function *getpid()* returns the PID of the calling process
 - Despite the fact that this value is stored in the global variable *currp*
- The principle of information hiding states that implementation details should be hidden unless necessary
 - The implementation can change without affecting programs
 - Outweighs efficiency concerns (in general)

4

System Calls

- Suspension and resumption are straightforward
 - change state, manipulate queues
 - the *ready()* call discussed previously does these things
- System calls in Xinu are distinguished from internal functions (like *ready*)
 - They provide the external interface to the system
 - Like system calls in Linux even though the entire system is in one address space
- Thus, system calls must protect the system from illegal / invalid use and hide information about the underlying implementation

5

System Calls

- To provide protection, system calls do things that internal functions need not:
 - Check all arguments
 - Ensure that actions and changes leave global data structures in a consistent state
 - Report success or failure to the caller
- Make no assumptions about the calling process
- Take steps to prevent other processes from executing concurrently
 - Avoid functions that give up the CPU – avoid direct or indirect calls to *resched*
 - Disable interrupts to prevent involuntary preemption
 - Return with the same interrupt status as when called

6

System Call Example: *resume()*

```
/* resume.c - resume */

#include <xinu.h>

/*-----
 * resume - Unsuspend a process, making it ready
 *-----*/

pril6 resume(
    pid32 pid /* ID of process to unsuspend */
)
{
    intmask mask; /* saved interrupt mask */
    struct procent *prptr; /* ptr to process' table entry */
    pril6 prio; /* priority to return */

    mask = disable();
    if (isbadpid(pid)) {
        restore(mask);
        return (pril6)SYSERR;
    }
    prptr = &proctab[pid];
```

7

```
    if (prptr->prstate != PR_SUSP) {
        restore(mask);
        return (pril6)SYSERR;
    }

    prio = prptr->prprio; /* record priority to return */
    ready(pid, RESCHED_YES);
    restore(mask);
    return prio;
}
```

8

Interrupt state

- Disable interrupts to prevent involuntary preemption
- Return with the same interrupt status as when called
- Rather than just enabling interrupts before returning, save and restore interrupt status
- `disable()` turns off interrupts and records the prior state (mask)
 - Which may have been disabled already
- `restore()` takes a prior state and configures the processor with it
- Calling *disable* works called with interrupts enabled or disabled
- *An operating system function always returns to its caller with the same interrupt status as when it was called.*

9

System Call Skeleton

```
syscall function_name ( args ) {  
    intmask mask;          /* interrupt mask */  
    mask = disable ( ) ; /* disable interrupts at start of function */  
    if ( args are incorrect ) {  
        restore (mask) ; /* restore interrupts before error return */  
        return (SYSERR) ;  
    }  
    ... other processing ...  
    if ( an error occurs ) {  
        restore (mask) ; /* restore interrupts before error return */  
        return (SYSERR) ;  
    }  
    ... more processing ...  
    restore (mask) ;      /* restore interrupts before normal return */  
    return ( appropriate value ) ;  
}
```

10

Implementation and Suspend

- Some system calls simply return the constant OK, while others return a relevant value
 - SYSERR is returned otherwise
- Suspend can be called on a process that is ready
 - easy: remove from ready list, change state to suspended (PR_SUSP)
- Applying to the current process is slightly more involved
 - it must result in calling *resched*

11

Suspend

```
/* suspend.c - suspend */  
  
#include <xinu.h>  
  
/*-----  
 * suspend - Suspend a process, placing it in hibernation  
 *-----  
 */  
syscall suspend(  
    pid32 pid /* ID of process to suspend */  
)  
{  
    intmask mask;          /* saved interrupt mask */  
    struct procent *prptr; /* ptr to process' table entry */  
    pri16 prio;            /* priority to return */  
  
    mask = disable();  
    if (isbadpid(pid) || (pid == NULLPROC)) {  
        restore(mask);  
        return SYSERR;  
    }  
}
```

12

```

/* Only suspend a process that is current or ready */

prptr = &proctab[pid];
if ((prptr->prstate != PR_CURR) && (prptr->prstate != PR_READY)) {
    restore(mask);
    return SYSERR;
}
if (prptr->prstate == PR_READY) {
    getitem(pid);          /* remove a ready process */
                          /* from the ready list */
    prptr->prstate = PR_SUSP;
} else {
    prptr->prstate = PR_SUSP; /* mark the current process */
    resched();              /* suspended and reschedule */
}
prio = prptr->prprio;
restore(mask);
return prio;
}

```

13

Priority used to communicate information

```

newprio = suspend ( getpaid ( ) );
if ( newprio == 25 ) {
    ... event 1 occurred ...
} else {
    ... event 2 occurred ...
}

```

- Suspend returns the priority that the process had when suspend returns

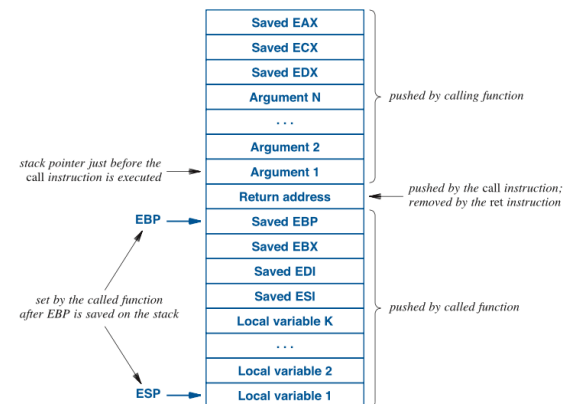
14

Process Creation

- The create() call configures the stack as if the function had been called from elsewhere in the code
- Configures the stack according to the compiler and architecture calling conventions
- Then it creates entries on the stack as if *ctxsw* had been called

15

x86 Stack Layout



16

```

/* create.c - create, newpid */
#include <xinu.h>

local int newpid();

#define roundmb(x) ((x+3)& ~0x3)

/*-----
 * create - create a process to start running a procedure
 *-----
 */
pid32 create(
    void *funcaddr, /* procedure address */
    uint32 ssize, /* stack size in words */
    pri16 priority, /* process priority > 0 */
    char *name, /* name (for debugging) */
    uint32 nargs, /* number of args that follow */
    ...
)
{
    uint32 savsp, *pushsp;
    intmask mask; /* interrupt mask */
    pid32 pid; /* stores new process id */
    struct proctab *prptr; /* pointer to proc. table entry */
    int32 i;

```

17

```

uint32 *a; /* points to list of args */
uint32 *saddr; /* stack address */

mask = disable();
if (ssize < MINSTK)
    ssize = MINSTK;
ssize = (uint32) roundmb(ssize);
if (((saddr = (uint32 *)getstk(ssize)) ==
    (uint32 *)SYSERR) ||
    (pid=newpid()) == SYSERR || priority < 1) {
    restore(mask);
    return SYSERR;
}

prcount++;
prptr = &proctab[pid];

/* initialize process table entry for new process */
prptr->prstate = PR_SUSP; /* initial state is suspended */
prptr->prprio = priority;
prptr->prstkbase = (char *)saddr;
prptr->prstklen = ssize;
prptr->prname[PNMLEN-1] = NULLCH;
for (i=0; i<PNMLEN-1 && (prptr->prname[i]!=NULLCH;i++);
prptr->prsem = -1;
prptr->prparent = (pid32) getpid();
prptr->prhasmsg = FALSE;

```

18

```

/* set up initial device descriptors for the shell */
prptr->prdesc[0] = CONSOLE; /* stdin is CONSOLE device */
prptr->prdesc[1] = CONSOLE; /* stdout is CONSOLE device */
prptr->prdesc[2] = CONSOLE; /* stderr is CONSOLE device */

/* Initialize stack as if the process was called */

*saddr = STACKMAGIC;
savsp = (uint32)saddr;

/* push arguments */
a = (uint32 *)(&nargs + 1); /* start of args */
a += nargs - 1; /* last argument */
for (; nargs > 0; nargs--) /* machine dependent; copy args */
    *--saddr = *a--; /* onto created process' stack */
*--saddr = (long)INITRET; /* push on return address */

/* The following entries on the stack must match what ctxsw */
/* expects a saved process state to contain: ret address, */
/* ebp, interrupt mask, flags, registers, and an old SP */

*--saddr = (long)funcaddr; /* Make the stack look like it's */
/* half-way through a call to */
/* ctxsw that "returns" to the */
/* new process */

```

19

```

*--saddr = savsp; /* This will be register ebp */
/* for process exit */
savsp = (uint32)saddr; /* start of frame for ctxsw */
*--saddr = 0x00000200; /* New process runs with */
/* interrupts enabled */

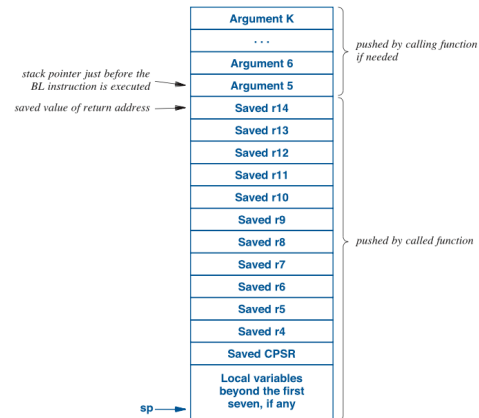
/* Basically, the following emulates a x86 "pushal" instruction */

*--saddr = 0; /* %eax */
*--saddr = 0; /* %ecx */
*--saddr = 0; /* %edx */
*--saddr = 0; /* %ebx */
*--saddr = 0; /* %esp; value filled in below */
pushsp = saddr; /* remember this location */
*--saddr = savsp; /* %ebp (while finishing ctxsw) */
*--saddr = 0; /* %esi */
*--saddr = 0; /* %edi */
*pushsp = (unsigned long) (prptr->prstkptr = (char *)saddr);
restore(mask);
return pid;
}

```

20

ARM Stack



21

ARM version

```
/* push arguments */
a = (uint32 *)(&nargs + 1); /* start of args */
a += nargs - 1; /* last argument */
for ( ; nargs > 4 ; nargs--) /* machine dependent; copy args */
    *--saddr = *a--; /* onto created process' stack */

*--saddr = (long)procaddr;
for(i = 11; i >= 4; i--)
    *--saddr = 0;
for(i = 4; i > 0; i--) {
    if(i <= nargs)
        *--saddr = *a--;
    else
        *--saddr = 0;
}
*--saddr = (long)INITRET; /* push on return address */
*--saddr = (long)0x00000053; /* CPSR, A, F bits set, */
/* Supervisor mode */

prptr->prstkptr = (char *)saddr;
restore(mask);
return pid;
```

22

newpid()

```
/*-----
 * newpid - Obtain a new (free) process ID
 *-----
 */
local pid32 newpid(void)
{
    uint32 i; /* iterate through all processes*/
    static pid32 nextpid = 1; /* position in table to try or */
    /* one beyond end of table */

    /* check all NPROC slots */

    for (i = 0; i < NPROC; i++) {
        nextpid %= NPROC; /* wrap around to beginning */
        if (proctab[nextpid].prstate == PR_FREE) {
            return nextpid++;
        } else {
            nextpid++;
        }
    }
    return (pid32) SYSERR;
}
```

23


userret()

```
/* userret.c - userret */

#include <xinu.h>

/*-----
 * userret - Called when a process returns from the top-level function
 *-----
 */
void userret(void)
{
    kill(getpid()); /* force process exit */
}
```

24



25

kill()

```

/* kill.c - kill */

#include <xinu.h>

/*-----
 * kill - Kill a process and remove it from the system
 *-----
 */
syscall kill(
    pid32 pid /* ID of process to kill */
)
{
    intmask mask; /* saved interrupt mask */
    struct procent *prptr; /* ptr to process' table entry */
    int32 i; /* index into descriptors */

    mask = disable();
    if (isbadpid(pid) || (pid == NULLPROC)
        || ((prptr = &proctab[pid])->prstate) == PR_FREE) {
        restore(mask);
        return SYSERR;
    }
}

```

26

```

if (--prcount <= 1) { /* last user process completes */
    xdone();
}

send(prptr->prparent, pid);
for (i=0; i<3; i++) {
    close(prptr->prdesc[i]);
}
freestk(prptr->prstkbase, prptr->prstklen);

switch (prptr->prstate) {
case PR_CURR:
    prptr->prstate = PR_FREE; /* suicide */
    resched();

case PR_SLEEP:
case PR_RECTIM:
    unsleep(pid);
    prptr->prstate = PR_FREE;
    break;

case PR_WAIT:
    semtab[prptr->prsem].scount++;
    /* fall through */
}

```

27

```

case PR_READY:
    getitem(pid); /* remove from queue */
    /* fall through */

default:
    prptr->prstate = PR_FREE;
}

restore(mask);
return OK;
}

```

28

xdone()

```
/* xdone.c - xdone */

#include <xinu.h>

/*-----
 * xdone - Print system completion message as last thread exits
 *-----
 */
void xdone(void)
{
    kprintf("\r\n\r\nAll user processes have completed.\r\n\r\n");
    /* gpioLEDOff(GPIO_LED_CISCOWHT); turn off LED "run" light */
    halt(); /* halt the processor */
}
```

29

getprio()

```
/* getprio.c - getprio */

#include <xinu.h>

/*-----
 * getprio - Return the scheduling priority of a process
 *-----
 */
syscall getprio(
    pid32 pid /* process ID */
)
{
    intmask mask; /* saved interrupt mask */
    uint32 prio; /* priority to return */

    mask = disable();
    if (isbadpid(pid)) {
        restore(mask);
        return SYSERR;
    }
    prio = proctab[pid].prprio;
    restore(mask);
    return prio;
}
```

30

chprio()

```
/* chprio.c - chprio */

#include <xinu.h>

/*-----
 * chprio - Change the scheduling priority of a process
 *-----
 */
pril6 chprio(
    pid32 pid, /* ID of process to change */
    pril6 newprio /* new priority */
)
{
    intmask mask; /* saved interrupt mask */
    struct procent *prptr; /* ptr to process' table entry */
    pril6 oldprio; /* priority to return */

    mask = disable();
    if (isbadpid(pid)) {
        restore(mask);
        return (pril6) SYSERR;
    }
    prptr = &proctab[pid];
    oldprio = prptr->prprio;
    prptr->prprio = newprio;
    restore(mask);
    return oldprio;
}
```

31