

# Operating Systems

Xinu - Coordination of Concurrent Processes

1

## Coordination and Synchronization

- Concurrent processes need to cooperate when sharing global data
  - Mutual exclusion accessing variables
- Allow a subset of processes to contend for access to a resource
  - Do not disable all interrupts
  - Block only appropriate processes
- Fair access
  - Each contending process eventually receives access
  - No starvation

2

## Semaphores

- *wait* and *signal* operations
  - also known as down/up or P/V
- Provide a solution for
  - Mutual exclusion
    - Binary semaphore
  - Producer-consumer interaction
    - Counting semaphore
- Busy waiting
  - *Generally* to be avoided as it deprives other processes of the CPU
  - Sometimes necessary for correctness
  - Sometimes desirable for performance (must be used judiciously)


3

## Scheduler Interaction

- Semaphores without busy waiting involve process lists and management by the scheduler
- When a semaphore is signaled, the scheduler should activate a waiting process (if the list is not empty)
- Scheduler decision
  - Priority? Wait time? Random?
- Random works but requires random number generation
- Xinu uses FCFS

4

## Process State: Waiting (PR\_WAIT)



5

## Semaphore Structures

```

/* semaphore.h - isbadsem */

#ifndef NSEM
#define NSEM 120 /* number of semaphores, if not defined */
#endif

/* Semaphore state definitions */

#define S_FREE 0 /* semaphore table entry is available */
#define S_USED 1 /* semaphore table entry is in use */

/* Semaphore table entry */
struct sentry {
    byte sstate; /* whether entry is S_FREE or S_USED */
    int32 scount; /* count for the semaphore */
    qid16 squeue; /* queue of processes that are waiting */
                /* on the semaphore */
};

extern struct sentry semtab[]; /* index identifies semaphore */

#define isbadsem(s) ((int32)(s) < 0 || (s) >= NSEM);
  
```

6

## Wait system call - wait.c

```

/* wait.c - wait */

#include <xinu.h>

/*-----
 * wait - Cause current process to wait on a semaphore
 *-----*/

syscall wait( sid32 sem ) /* semaphore on which to wait */
{
    intmask mask; /* saved interrupt mask */
    struct procent *prptr; /* ptr to process' table entry */
    struct sentry *semptr; /* ptr to semaphore table entry */

    mask = disable();

    if (isbadsem(sem)) {
        restore(mask);
        return SYSERR;
    }
  
```

7

```

    semptr = &semtab[sem];
    if (semptr->sstate == S_FREE) {
        restore(mask);
        return SYSERR;
    }

    if (--(semptr->scount) < 0) { /* if caller must block */
        prptr = &proctab[currpid];
        prptr->prstate = PR_WAIT; /* set state to waiting */
        prptr->prsem = sem; /* record semaphore ID */
        enqueue(currpid, semptr->squeue); /* enqueue on semaphore */
        resched(); /* and reschedule */
    }

    restore(mask);
    return OK;
}
  
```

- Once enqueued, a process remains in PR\_WAIT until it reaches the head of the queue and another process signals the semaphore
- Then it is moved to the ready queue
- ctxsw -> resched -> wait -> process code

8

## Signal system call - signal.c

```

/* signal.c - signal */

#include <xinu.h>

/*-----
 * signal - Signal a semaphore, releasing a process if one is waiting
 *-----*/

/*
syscall signal( sid32 sem )      /* id of semaphore to signal */
{
    intmask mask;                /* saved interrupt mask */
    struct sentry *semptr;       /* ptr to semaphore table entry */

    mask = disable();
    if (isbadsem(sem)) {
        restore(mask);
        return SYSERR;
    }
}

```

9

```

semptr = &semtab[sem];

if (semptr->sstate == S_FREE) {
    restore(mask);
    return SYSERR;
}

if ((semptr->scount++) < 0) {      /* release a waiting process */
    ready(dequeue(semptr->squeue), RESCHED_YES);
}

restore(mask);
return OK;
}

```

- A non-negative semaphore count means that queue is empty
  - Count of N means that wait can be called N times before a process blocks
- A semaphore count of negative N means that the queue contains N waiting processes

10

## Static and Dynamic Allocation

- Static allocation – fixed set of semaphores at system compile time
- Dynamic allocation – the system includes functions for creating and destroying semaphores
- Static allocation
  - Reduction in overhead
  - Reduction in flexibility
- Most systems provide dynamic allocation as we have seen

11

## Dynamic semaphore allocation

```

/* semcreate.c - semcreate, newsem */
#include <xinu.h>

local sid32 newsem(void);

/*-----
 * semcreate - create a new semaphore and return the ID to the caller
 *-----*/

sid32 semcreate( int32 count      /* initial semaphore count */
{
    intmask mask;                /* saved interrupt mask */
    sid32 sem;                   /* semaphore ID to return */

    mask = disable();

    if (count < 0 || ((sem=newsem())==SYSERR)) {
        restore(mask);
        return SYSERR;
    }

    semtab[sem].scount = count;   /* initialize table entry */

    restore(mask);
    return sem;
}

```

12

```

/*-----
 * newsem - allocate an unused semaphore and return its index
 *-----
 */
local sid32 newsem(void)
{
    static sid32 nextsem = 0; /* next semaphore index to try */
    sid32 sem; /* semaphore ID to return */
    int32 i; /* iterate through # entries */

    for (i=0 ; i<NSEM ; i++) {
        sem = nextsem++;
        if (nextsem >= NSEM)
            nextsem = 0;
        if (semtab[sem].sstate == S_FREE) {
            semtab[sem].sstate = S_USED;
            return sem;
        }
    }
    return SYSERR;
}

```

13

```

/* semdelete.c - semdelete */

#include <xinu.h>

/*-----
 * semdelete -- Delete a semaphore by releasing its table entry
 *-----
 */
syscall semdelete( sid32 sem ) /* ID of semaphore to delete */
{
    intmask mask; /* saved interrupt mask */
    struct sentry *semptr; /* ptr to semaphore table entry */

    mask = disable();
    if (isbadsem(sem)) {
        restore(mask);
        return SYSERR;
    }

    semptr = &semtab[sem];
    if (semptr->sstate == S_FREE) {
        restore(mask);
        return SYSERR;
    }
}

```

14

```

semptr->sstate = S_FREE;

while (semptr->scount++ < 0) { /* free all waiting processes */
    ready(getfirst(semptr->squeue), RESCHED_NO);
}
resched();
restore(mask);
return OK;
}

```

- Makes all waiting processes ready
- Deleting a semaphore with waiting processes could be considered an error

15

```

/* semreset.c - semreset */

#include <xinu.h>

/*-----
 * semreset -- reset a semaphore's count and release waiting processes
 *-----
 */
syscall semreset(
    sid32 sem, /* ID of semaphore to reset */
    int32 count /* new count (must be >= 0) */
)
{
    intmask mask; /* saved interrupt mask */
    struct sentry *semptr; /* ptr to semaphore table entry */
    qid16 semqueue; /* semaphore's process queue ID */
    pid32 pid; /* ID of a waiting process */

    mask = disable();

    if (count < 0 || isbadsem(sem) || semtab[sem].sstate == S_FREE) {
        restore(mask);
        return SYSERR;
    }
}

```

16

```

sempr = &semtab[sem];

semqueue = sempr->squeue; /* free any waiting processes */
while ((pid=getfirst(semqueue)) != EMPTY)
    ready(pid, RESCHED_NO);

sempr->scount = count; /* reset count as specified */
resched();
restore(mask);
return(OK);
}

```

17

## Coordination across Multiple Processors / Cores

- Xinu semaphores as described work well on single CPUs
- Disabling interrupts on a system with multiple cores is inefficient
  - No progress on I/O on any core, and no timer interrupts!
- Test and set instruction can be used for atomic access without disabling interrupts
  - Or compare and swap
- The book describes spin locks in this context, but atomic instructions can be used with waiting and queues as well
  - Adaptive or hybrid solutions are possible

18