

Operating Systems

Inter-process Communication (IPC)

Xinu Message Passing

1

Xinu Messages

- Xinu supports completely synchronous and partially asynchronous message functionality
 - Also illustrates direct vs indirect operation with point to point exchange or rendezvous
 - Asynchronous, indirect case discussed later
- For the synchronous case, the system is designed to ensure that processes do not block and that waiting messages don't grow to consume undue memory
 - Primitive operation for resource-constrained embedded systems

2

Message Passing Design

- Limited message size
 - The system limits each message to a small, fixed size
 - In the basic implementation, it is one word (int)
- No message queues
 - The system is permitted to store only one unreceived message – per process – at any time
- First message semantics
 - If several messages are sent to a process, only the first is stored and the subsequent senders do not block
 - This is useful for determining which of a set of events completes first

3


Xinu MP Functions

- Three system calls: *send*, *receive* and *recvclr*
- *send* takes a message and a PID
- *receive* takes no arguments, waits until a message arrives, and returns it
 - Or returns with a message immediately
- *recvclr* is a non-blocking version of *receive*
 - If a message has arrived, return the message immediately
 - If no message has arrived, return value *OK* immediately
 - Also used to clear an old message if one exists

4

Xinu MP Functions

- Process state for messaging: PR_RECV



5

Implementation of Send

- Requires agreement / coordination between senders and receivers
- Sender must store the message somewhere
 - Can't be in the sender's memory, since it might exit
 - Can't be in the receiver's memory since writing into another process's memory is problematic for security, coordination
- Message size and pending limitation addresses this for Xinu
- Space is allocated in the process table entry for a message destined to that process

6

Implementation of Send

```

/* send.c - send */

#include <xinu.h>

/*-----
 * send - pass a message to a process and start recipient if waiting
 *-----*/
syscall send(
    pid32    pid,          /* ID of recipient process */
    umsg32   msg,          /* contents of message */
    intmask  mask;         /* saved interrupt mask */
    struct   procent *prptr; /* ptr to process' table entry */
)

```

7

```

mask = disable();
if (isbadpid(pid)) {
    restore(mask);
    return SYSERR;
}
prptr = &proctab[pid];
if ((prptr->prstate == PR_FREE) || prptr->prhasmsg) {
    restore(mask);
    return SYSERR;
}
prptr->prmsg = msg;          /* deliver message */
prptr->prhasmsg = TRUE;      /* indicate message is waiting */

/* If recipient waiting or in timed-wait make it ready */
if (prptr->prstate == PR_RECV) {
    ready(pid, RESCHED_YES);
} else if (prptr->prstate == PR_RECTIM) {
    unsleep(pid);
    ready(pid, RESCHED_YES);
}
restore(mask);              /* restore interrupts */
return OK;
}

```

8

Implementation of Receive

```
/* receive.c - receive */

#include <xinu.h>

/*-----
 * receive - wait for a message and return the message to the caller
 *-----
 */
umsg32 receive(void)
{
    intmask    mask;           /* saved interrupt mask */
    struct      procent *prptr; /* ptr to process' table entry */
    umsg32      msg;           /* message to return */

    mask = disable();
    prptr = &proctab[currpid];
```

9

```
    if (prptr->prhasmsg == FALSE) {
        prptr->prstate = PR_RECV;
        resched();           /* block until message arrives */
    }
    msg = prptr->prmsg;       /* retrieve message */
    prptr->prhasmsg = FALSE;  /* reset message flag */
    restore(mask);
    return msg;
}
```

10

Implementation of Non-Blocking Message Reception

```
/* recvclr.c - recvclr */

#include <xinu.h>

/*-----
 * recvclr - clear incoming message, and return message if one waiting
 *-----
 */
umsg32 recvclr(void)
{
    intmask    mask;           /* saved interrupt mask */
    struct      procent *prptr; /* ptr to process' table entry */
    umsg32      msg;           /* message to return */

    mask = disable();
    prptr = &proctab[currpid];
```

11

```
    if (prptr->prhasmsg == TRUE) {
        msg = prptr->prmsg;       /* retrieve message */
        prptr->prhasmsg = FALSE;  /* reset message flag */
    } else {
        msg = OK;
    }

    restore(mask);
    return msg;
}
```

12

Summary of Simple Messages in Xinu

- Compact and efficient code for basic message passing
 - General purpose message passing is potentially much more complex
- Synchronous, blocking the receiver if desired
- Limits message size and queue length
- First message semantics with only one outstanding message

13

High-level Message Passing in Xinu

- The low-level messaging passing interface permits a process to send a message directly to another process
 - Direct message approach – cannot coordinate multiple receivers
- High-level interface provides
 - Buffering of a specified number of messages
 - Indirect approach
- Defines an IPC port
 - Rendezvous point

14

Port Interface

- Messages are (still) a 32-bit word
- *ptsend* – deposits a message in a port
- *ptrecv* – receives a message from a port
- Sending and receiving are synchronous
- If space exists, the sender can deposit a message immediately
 - If the port is full, then the sender is blocked
- If a port is empty, receive will block
- Messages are handed FIFO, as are blocked processes
 - If multiple processes are blocked, the one waiting the longest will get the first new message, or send first into a new message slot

15

Port Implementation

- Each port consists of a queue to hold messages
 - Two semaphores – reader and writer
- Fixed number of message slots, or nodes
 - Shared among all port functions
- Initially linked into a single free list *ptfree*
- Send removes a node from the free list and adds it to the associated port queue
- Receive removes the node from the queue and restores it in the free list

16

The Implementation of Ports

```
/* ports.h - isbadport */

#define NPORTS      30          /* maximum number of ports */
#define PT_MSGS     100        /* total messages in system */
#define PT_FREE     1          /* port is free */
#define PT_LIMBO    2          /* port is being deleted/reset */
#define PT_ALLOC    3          /* port is allocated */

struct ptnode {
    uint32  ptmsg;              /* a one-word message */
    struct  ptnode *ptnext;     /* ptr to next node on list */
};

struct ptenry {
    sid32   ptssem;             /* sender semaphore */
    sid32   ptrsem;             /* receiver semaphore */
    uint16  ptstate;            /* port state (FREE/LIMBO/ALLOC) */
    uint16  ptmaxcnt;           /* max messages to be queued */
};
```

17

```
int32  ptseq;                  /* sequence changed at creation */
struct ptnode *pthead;         /* list of message pointers */
struct ptnode *pttail;        /* tail of message list */
};

extern struct ptnode *ptfree;   /* list of free nodes */
extern struct ptenry porttab[]; /* port table */
extern int32  ptnextid;        /* next port ID to try when
                                looking for a free slot */

#define isbadport(portid)      ( (portid)<0 || (portid)>=NPORTS )
```

18

Port Table Initialization

```
/* ptinit.c - ptinit */

#include <xinu.h>

struct ptnode *ptfree;         /* list of free message nodes */
struct ptenry porttab[NPORTS]; /* port table */
int32  ptnextid;               /* next table entry to try */

/*-----
 * ptinit -- initialize all ports
 *-----
 */
syscall ptinit(
    int32      maxmsgs          /* total messages in all ports */
)
```

19

```
{
    int32      i;                /* runs through port table */
    struct     ptnode *next, *prev; /* used to build free list */

    ptfree = (struct ptnode *)getmem(maxmsgs*sizeof(struct ptnode));
    if (ptfree == (struct ptnode *)SYSERR) {
        panic("pinit - insufficient memory");
    }
    /* Initialize all port table entries to free */

    for (i=0 ; i<NPORTS ; i++) {
        porttab[i].ptstate = PT_FREE;
        porttab[i].ptseq = 0;
    }
    ptnextid = 0;

    /* Create free list of message pointer nodes */

    for ( prev=next=ptfree ; --maxmsgs > 0 ; prev=next )
        prev->ptnext = ++next;
    prev->ptnext = NULL;
    return(OK);
}
```

20

Port Creation

```
/* pcreate.c - pcreate */

#include <xinu.h>

/*-----
 * pcreate -- create a port that allows "count" outstanding messages
 *-----
 */
syscall pcreate(
    int32      count
)
{
    intmask    mask;          /* saved interrupt mask          */
    int32      i;             /* counts all possible ports */
    int32      ptnum;         /* candidate port number to try */
    struct      pentry *ptptr; /* pointer to port table entry */
}
```

21

```
mask = disable();
if (count < 0) {
    restore(mask);
    return(SYSERR);
}

for (i=0 ; i<NPORTS ; i++) { /* count all table entries */
    ptnum = ptnextid;         /* get an entry to check */
    if (++ptnextid >= NPORTS) {
        ptnextid = 0;        /* reset for next iteration */
    }

    /* Check table entry that corresponds to ID ptnum */

    ptptr = &porttab[ptnum];
    if (ptptr->ptstate == PT_FREE) {
        ptptr->ptstate = PT_ALLOC;
        ptptr->ptssem = semcreate(count);
        ptptr->ptrsem = semcreate(0);
        ptptr->pthead = ptptr->pttail = NULL;
        ptptr->ptseq++;
        ptptr->ptmaxcnt = count;
        restore(mask);
        return(ptnum);
    }
}
restore(mask);
return(SYSERR);
}
```

22

Sending A Message To A Port

```
/* ptsend.c - ptsend */

#include <xinu.h>

/*-----
 * ptsend -- send a message to a port by adding it to the queue
 *-----
 */
syscall ptsend(
    int32      portid,        /* ID of port to use          */
    umsg32     msg,           /* message to send            */
)
{
    intmask    mask;          /* saved interrupt mask          */
    struct      pentry *ptptr; /* pointer to table entry        */
    int32      seq;           /* local copy of sequence num.  */
}
```

23

```
struct      ptnode *msgnode; /* allocated message node */
struct      ptnode *tailnode; /* last node in port or NULL */

mask = disable();
if ( isbadport(portid) ||
    (ptptr = &porttab[portid])->ptstate != PT_ALLOC ) {
    restore(mask);
    return SYSERR;
}

/* Wait for space and verify port has not been reset */

seq = ptptr->ptseq;          /* record original sequence */
if (wait(ptptr->ptssem) == SYSERR
    || ptptr->ptstate != PT_ALLOC
    || ptptr->ptseq != seq) {
    restore(mask);
    return SYSERR;
}

if (ptfree == NULL) {
    panic("Port system ran out of message nodes");
}
```

24

```

/* Obtain node from free list by unlinking */
msgnode = ptfree;          /* point to first free node */
ptfree = msgnode->ptnext;  /* unlink from the free list*/
msgnode->ptnext = NULL;    /* set fields in the node */
msgnode->ptmsg = msg;

/* Link into queue for the specified port */

tailnode = ptptr->pttail;
if (tailnode == NULL) {    /* queue for port was empty */
    ptptr->pttail = ptptr->pthead = msgnode;
} else {                  /* insert new node at tail */
    tailnode->ptnext = msgnode;
    ptptr->pttail = msgnode;
}
signal(ptptr->ptrsem);
restore(mask);
return OK;
}

```

25

Receiving A Message from A Port

```

/* ptrecev.c - ptrecev */

#include <xinu.h>

/*-----
 * ptrecev -- receive a message from a port, blocking if port empty
 *-----
 */
uint32 ptrecev(
    int32      portid      /* ID of port to use */
)
{
    intmask    mask;        /* saved interrupt mask */
    struct      pentry *ptptr; /* pointer to table entry */
    int32      seq;         /* local copy of sequence num.*/
    umsg32     msg;         /* message to return */

```

26

```

struct ptnode *msgnode;    /* first node on message list */

mask = disable();
if ( isbadport(portid) ||
    (ptptr = &porttab[portid])->ptstate != PT_ALLOC ) {
    restore(mask);
    return (uint32)SYSERR;
}

/* Wait for message and verify that the port is still allocated */

seq = ptptr->ptseq;        /* record original sequence num */
if (wait(ptptr->ptrsem) == SYSERR || ptptr->ptstate != PT_ALLOC
    || ptptr->ptseq != seq) {
    restore(mask);
    return (uint32)SYSERR;
}

```

27

```

/* Dequeue first message that is waiting in the port */

msgnode = ptptr->pthead;
msg = msgnode->ptmsg;
if (ptptr->pthead == ptptr->pttail) /* delete last item */
    ptptr->pthead = ptptr->pttail = NULL;
else
    ptptr->pthead = msgnode->ptnext;
msgnode->ptnext = ptfree; /* return to free list */
ptfree = msgnode;
signal(ptptr->ptssem);
restore(mask);
return msg;
}

```

28

Port Deletion and Reset

```

/* ptdelete.c - ptdelete */

#include <xinu.h>

/*-----
 * ptdelete -- delete a port, freeing waiting processes and messages
 *-----
 */
syscall ptdelete(
    int32    portid,          /* ID of port to delete */
    int32    (*dispose)(int32) /* function to call to dispose */
    )          /* of waiting messages */
{
    intmask    mask;          /* saved interrupt mask */
    struct      pentry *ptptr; /* pointer to port table entry */

```

29

```

    mask = disable();
    if ( isbadport(portid) ||
        (ptptr = &porttab[portid])->ptstate != PT_ALLOC ) {
        restore(mask);
        return(SYSERR);
    }
    _ptclear(ptptr, PT_FREE, dispose);
    ptnextid = portid;
    restore(mask);
    return(OK);
}

```

30

```

/* ptreset.c - ptreset */

#include <xinu.h>

/*-----
 * ptreset -- reset a port, freeing waiting processes and messages and
 * leaving the port ready for further use
 *-----
 */
syscall ptreset(
    int32    portid,          /* ID of port to reset */
    int32    (*dispose)(int32) /* function to call to dispose */
    )          /* of waiting messages */

```

31

```

{
    intmask    mask;          /* saved interrupt mask */
    struct      pentry *ptptr; /* pointer to port table entry */

    mask = disable();
    if ( isbadport(portid) ||
        (ptptr = &porttab[portid])->ptstate != PT_ALLOC ) {
        restore(mask);
        return SYSERR;
    }
    _ptclear(ptptr, PT_ALLOC, dispose);
    restore(mask);
    return OK;
}

```

32


```

/* ptclean.c - _ptclean */

#include <xinu.h>

/*-----
 * _ptclean -- used by ptdelete and ptreset to clear or reset a port
 *             (internal function assumes interrupts disabled
 *             and arguments have been checked for validity)
 *-----
 */
void _ptclean(
    struct ptentry *ptptr, /* table entry to clear */
    uint16 newstate, /* new state for port */
    int32 (*dispose)(int32) /* disposal function to call */
)
{
    struct ptnode *walk; /* pointer to walk message list */

```

33

```

/* Place port in limbo state while waiting processes are freed */

ptptr->ptstate = PT_LIMBO;

ptptr->ptseq++; /* reset accession number */
walk = ptptr->pthead; /* first item on msg list */

if ( walk != NULL ) { /* if message list nonempty */

    /* Walk message list and dispose of each message */

    for( ; walk!=NULL ; walk=walk->ptnext) {
        (*dispose)( walk->ptmsg );
    }

    /* Link entire message list into the free list */

    (ptptr->pttail)->ptnext = ptfree;
    ptfree = ptptr->pthead;
}

```

34

```

if (newstate == PT_ALLOC) {
    ptptr->pttail = ptptr->pthead = NULL;
    semreset(ptptr->ptssem, ptptr->ptmaxcnt);
    semreset(ptptr->ptrsem, 0);
} else {
    semdelete(ptptr->ptssem);
    semdelete(ptptr->ptrsem);
}
ptptr->ptstate = newstate;
return;
}

```

35

end

36