

Operating Systems

Inter-process Communication (IPC)

Deadlock

1

Coordination and Deadlock

- Process coordination can lead to *deadlock*
 - With blocking/synchronous application interfaces, processes may wait forever
- Obviously Process A can't wait on a message from B, while B waits on a message from A, or they will be waiting forever
 - This can occur with mutexes/semaphores as well

```
thread 1          thread 2
-----
. . .
pthread_mutex_lock(&a)  pthread_mutex_lock(&b)
pthread_mutex_lock(&b)  pthread_mutex_lock(&a)
. . .               . . .
```

2


Deadlock Definitions

- A set of processes is deadlocked when every process in the set is waiting for an event that can only be generated by some process in the set
- Livelock is similar but the states of the processes are changing
 - Still no process is actually progressing
- Contrast with starvation: a process waits indefinitely because some other process is using a resource
 - Deadlock and livelock are instances of starvation, but starvation is more general

3

Standard Problem – Dining Philosophers

- Acquire left and right forks/chopsticks to eat pasta/rice
- Get utensils, eat, drop utensils, think, repeat



4

Dining-Philosophers Problem (Cont.)

- The structure of Philosopher i :

```

While (true) {
    wait ( chopstick[i] );
    wait ( chopstick[ (i + 1) % 5] );

    // eat

    signal ( chopstick[i] );
    signal ( chopstick[ (i + 1) % 5] );

    // think
}

```

5

Dining Philosophers (3)

```

#define N      5          /* number of philosophers */
#define LEFT   (i+N-1)%N  /* number of i's left neighbor */
#define RIGHT  (i+1)%N    /* number of i's right neighbor */
#define THINKING 0        /* philosopher is thinking */
#define HUNGRY  1          /* philosopher is trying to get forks */
#define EATING  2          /* philosopher is eating */
typedef int semaphore;    /* semaphores are a special kind of int */
int state[N];             /* array to keep track of everyone's state */
semaphore mutex = 1;      /* mutual exclusion for critical regions */
semaphore s[N];           /* one semaphore per philosopher */

void philosopher(int i)    /* i: philosopher number, from 0 to N-1 */
{
    while (TRUE) {          /* repeat forever */
        think();           /* philosopher is thinking */
        take_forks(i);     /* acquire two forks or block */
        eat();             /* yum-yum, spaghetti */
        put_forks(i);      /* put both forks back on table */
    }
}

```

Solution to dining philosophers problem (part 1) -
Tannenbaum

6

Dining Philosophers (4)

```

void take_forks(int i)      /* i: philosopher number, from 0 to N-1 */
{
    down(&mutex);           /* enter critical region */
    state[i] = HUNGRY;      /* record fact that philosopher i is hungry */
    test(i);               /* try to acquire 2 forks */
    up(&mutex);             /* exit critical region */
    down(&s[i]);             /* block if forks were not acquired */
}

void put_forks(i)          /* i: philosopher number, from 0 to N-1 */
{
    down(&mutex);           /* enter critical region */
    state[i] = THINKING;    /* philosopher has finished eating */
    test(LEFT);            /* see if left neighbor can now eat */
    test(RIGHT);           /* see if right neighbor can now eat */
    up(&mutex);             /* exit critical region */
}

void test(i)               /* i: philosopher number, from 0 to N-1 */
{
    if (state[i] == HUNGRY && state[LEFT] != EATING && state[RIGHT] != EATING) {
        state[i] = EATING;
        up(&s[i]);
    }
}

```

Solution to dining philosophers problem (part 2)


7

Standard Problem – Dining Philosophers

- Acquire left and right forks/chopsticks to eat pasta/rice
 - Eat, drop utensils, think, repeat
- Solutions
 - Resource hierarchy – number “forks” and acquire in order
 - Arbitrator – acquire both “forks” at once or none at all (talking to a “waiter”, or using a mutex)
 - Communicate – indicate intent to acquire with a request or by setting state to “hungry”


8

Resource Allocation Graph



9

Wait-for Graph



10

Solutions to Deadlock

- The OS can't really do much
 - Application environments like databases can, however
- Some algorithms attempt to stay in a safe state by having processes declare their resource needs, and ensuring that at least one process can acquire all resources needed to complete
 - That there is a possible path to completion
 - Unrealistic to have processes declare their resource needs
- Can periodically test for cycles in the waiting graph
 - Kill all processes?
 - Unblock processes?
- A programmer can use non-blocking messaging interfaces, but care must be taken for synchronization

11

end

12