

## Operating Systems

### Interrupts

1

## Interrupts and Interrupt Handling

- Interrupt and exceptions concepts
- Interrupt implementation details
- Interrupt handling in the OS
- Motivation: separate I/O activities from processing. I/O is generally slow, but it can proceed in parallel with hardware that supports interrupting the CPU to signal completion of some I/O activity.

2

## Processing an Interrupt

- Jump to exception vector
- Disable further interrupts
  - Some counter-examples with interrupt priorities
- Save state to return to the code executing when the interrupt occurs
  - The processor will save only minimal values, leaving it to the OS to completely save state as necessary
- Jump to a specific memory location
  - Directly or indirectly arriving at the memory location with code to handle the specific interrupt
- Restore state and resume normal processing

3


## Vectored Interrupts

- Implementations vary, but in general an interrupt is identified by an Interrupt Request Number (IRQ)
  - Configured in different ways
- The IRQ is used as an index into a vector of function pointers
- `interrupt_vector[IRQ]` points to code to handle interrupt #IRQ

4

## Vectored Interrupts

Interrupt vector  
In memory



5

## Exceptions

- Exceptions and Interrupts are realized in essentially the same way
- Exceptions such as page faults, illegal instructions transfer control to specified handler code
- More generally, the interrupt vector is an exception vector
  - Particularly on x86; ARM has two levels as we will see

6

## Intel Exceptions and Interrupts

- The Intel Galileo has a unified vector of function pointers for both exceptions and IRQs
- Exception 0 indicates the “divide by zero” exception
- Hardware handles this just as if a device with IRQ 0 fired
  - Of course this means IRQ 0 can’t be used for a device

7

## Interrupts in ARM

- The ARM processor has various states including User mode, System mode and exception modes
- Distinct exception states
  - Software Interrupt
  - Reset
  - Undefined Instruction
  - Abort (Pre-fetch or Data)
  - Interrupt (IRQ and FIQ)
- Everything is mapped to one of these exception states
- Top level handlers are one of these exception states

8

## IRQ and FIQ

- ARM supports IRQs and Fast IRQs or FIQs
  - Xinu only supports IRQs
- Both IRQ and FIQ mode “bank” registers, meaning they are hidden until that mode is restored
  - r0 – r6 are visible across modes
  - IRQ mode banks the Stack Pointer (SP) and the Link Register (LR)
  - FIQ mode banks those as well as R8-R12
- FIQ mode can be used to optimize interrupt processing

9

## Interrupts in ARM

- When a IRQ is raised, the ARM processor:
  - Stops what it is doing, jumps to the IRQ handler location and switches to IRQ mode
  - Disables further IRQs (not FIQs)
  - Puts CPSR in SPSR
  - Puts the current PC in LR
  - Consults the exception vector and jumps to the appropriate handler
  - For an IRQ, it jumps to the exception handler for IRQs, which is itself a table of handlers

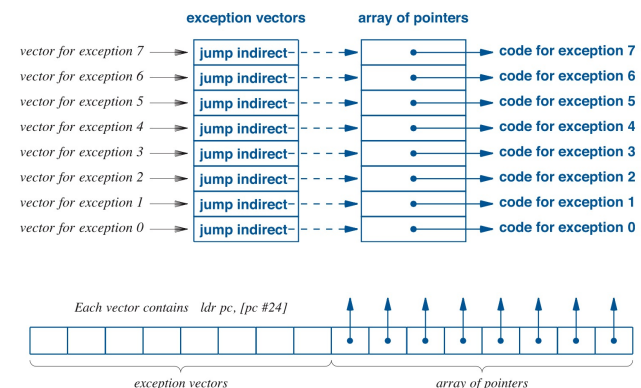
10

## ARM Exception Vectors

- While Intel’s exception vector contains pointers, the ARM exception vector contains instructions
- In practice, they are simply instructions that jump to an address but the implementation is different
  - ARM: PC = exception\_vector[i]
  - Intel: PC = \*exception\_vector[i]
- Each exception vector entry is one 32-bit instruction, so they are structured as indirect branch instructions

11

## Still an array of pointers in practice



12

## intr.S

```

/*intr.S - enable, disable, restore, halt, pause, irq_except(ARM) */

#include <armv7a.h>

.text
.globl disable
.globl restore
.globl enable
.globl pause
.globl halt
.globl irq_except
.globl irq_dispatch
.globl initvec
.globl expjpinstr

/*-----
 * disable - Disable interrupts and return the previous state
 *-----
 */

disable:
    mrs    r0, cpsr      /* Copy the CPSR into r0      */
    cpsid  i             /* Disable interrupts */
    mov    pc, lr        /* Return the CPSR   */

```

13

## intr.S

```

/*-----
 * restore - Restore interrupts to value given by mask argument
 *-----
 */

restore:
    push   {r1, r2}      /* Save r1, r2 on stack */
    mrs    r1, cpsr      /* Copy CPSR into r1   */
    ldr    r2, =0x01F00220
    and    r1, r1, r2     /* Extract flags and other important */
    bic    r0, r0, r2     /* bits from the mask */
    orr    r1, r1, r0
    msr    cpsr, r1       /* Restore the CPSR */
    pop    {r1, r2}      /* Restore r1, r2 */
    mov    pc, lr        /* Return to caller */

/*-----
 * enable - Enable interrupts
 *-----
 */

enable:
    cpsie  i             /* Enable interrupts */
    mov    pc, lr        /* Return */

```

14

## intr.S

```

/*-----
 * initvec - Initialize the exception vector
 *-----
 */

initvec:
    mrc    p15, 0, r0, c1, c0, 0 /* Read the c1-control register */
    bic    r0, r0, #ARMV7A_C1CTL_V /* V bit = 0, normal exp. base */
    mcr    p15, 0, r0, c1, c0, 0 /* Write the c1-control register */
    ldr    r0, =ARMV7A_EV_START /* Exception base address */
    mcr    p15, 0, r0, c12, c0, 0 /* Store expc. base addr. in c12 */
    ldr    r0, =ARMV7A_EV_START /* Start address of exp. vector */
    ldr    r1, =ARMV7A_EV_END /* End address of exp. vector */
    ldr    r2, =expjpinstr /* Copy the exp jump instr */
    ldr    r2, [r2] /* into register r2 */
expvect:
    str    r2, [r0] /* Store the jump instruction */
    add    r0, r0, #4 /* in the exception vector */
    cmp    r0, r1
    bne    expvec
    ldr    r0, =ARMV7A_EH_START /* Install the default exception */
    ldr    r1, =ARMV7A_EH_END /* handler for all exceptions */
    ldr    r2, =defexp_handler
    str    r2, [r0]
    add    r0, r0, #4
    cmp    r0, r1
    bne    exphnd
    ldr    r0, =ARMV7A_IRQH_ADDR /* Install the IRQ handler to */
    ldr    r1, =irq_except /* override the default */
    str    r1, [r0] /* exception handler */
    mov    pc, lr

```

15

## intr.S

```

/*-----
 * irq_except - Dispatch an IRQ exception to higher level IRQ dispatcher
 *-----
 */

irq_except:
    sub    lr, lr, #4 /* Correct the return address */
    srsdb  sp!, #19 /* Save return state on the supervisor */
    /* mode stack */
    cps    #19 /* Change to supervisor mode */
    push   {r0-r12, lr} /* Save all registers */
    bl     irq_dispatch /* Call IRQ dispatch */
    pop    {r0-r12, lr} /* Restore all registers */
    rfeia  sp! /* Return from the exception using info */
    /* stored on the stack */

/*-----
 * defexp_handler - Default Exception handler
 *-----
 */

defexp_handler:
    ldr    r0, =expmsg1
    mov    r1, lr
    bl     kprintf
    ldr    r0, =expmsg2
    bl     panic

```

16

## Assigning Device Numbers

- Manual configuration with jumpers and switches
  - This is where I came in. Error-prone even when you knew what you were doing...
- Automatic assignment at boot time
  - PCI brought the ability to assign IRQ numbers to devices after interrogating them
- Dynamic assignment for runtime pluggable devices
  - USB master driver loaded at initialization time
  - Gets an interrupt when a new device is connected
  - Interrogates device, loads appropriate driver
  - Vectors through another table for device interrupt handlers


17

## Interrupt Dispatching

- What is the division of labor between interrupt controller hardware and software?
- In software, how much is written in assembly and how much in a high-level language like C?
  - Obviously at least the register manipulation must be in assembly
- Differences in the BeagleBone and Galileo
  - The BB raises an IRQ exception and passes the IRQ number to the CPU
  - The Galileo interrupt controller holds the entire interrupt vector, calling the appropriate handler directly

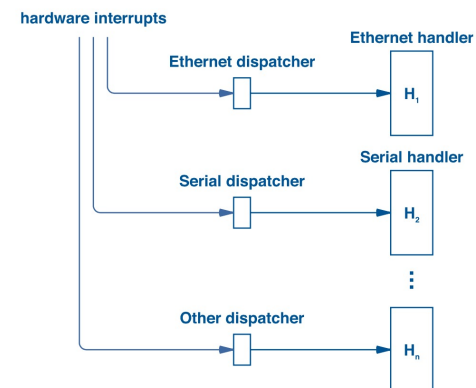
18

## Dispatching to Handlers



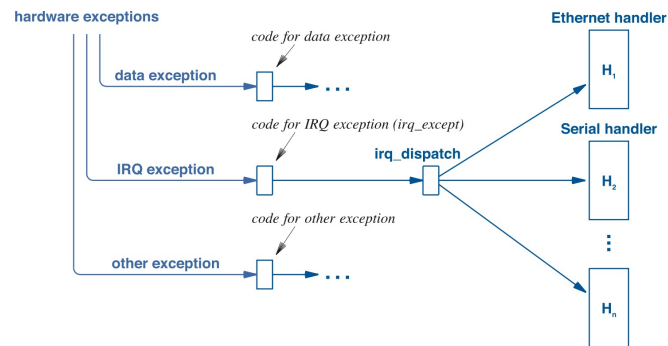
19

## Galileo



20

## BeagleBone



21

## irq\_dispatch

```

/*-----
 * irq_dispatch - call the handler for specific interrupt
 *-----
 */

void irq_dispatch(){

    struct intc_csreg *csrptr = (struct intc_csreg *)0x48200000;
    uint32 xnum;          /* Interrupt number of device */
    interrupt (*handler)(); /* Pointer to handler function */

    /* Get the interrupt number from the Interrupt controller */

    xnum = csrptr->sir_irq & 0x7F;

    /* If a handler is set for the interrupt, call it */

    if(intc_vector[xnum]) {
        handler = intc_vector[xnum];
        handler(xnum);
    }

    /* Acknowledge the interrupt */

    csrptr->control |= (INTC_CONTROL_NEWIRQAGR);
}

```

22

## Disabling Interrupts

- While an interrupt is being handled, further interrupts are disabled
- They remain disabled in the dispatcher and when the handler is called and returns
- Only re-enabled once the system is returning to user code
- This places constraints on how long the system can run with interrupts disabled

23

## Other constraints

- The interrupt handler can never invoke a function that will block
  - signal() but not wait()
  - The Null process must be current or ready
- Scheduling during an interrupt may be required
  - Maintain the scheduling invariant

24

## Exception Vectors (start.S)

```

/* ARM exception vector table. This is copied to address 0. See A2.6
 * "Exceptions" of the ARM Architecture Reference Manual. */
_vectors:
    ldr pc, reset_addr    /* Reset handler */
    ldr pc, undef_addr    /* Undefined instruction handler */
    ldr pc, swi_addr       /* Software interrupt handler */
    ldr pc, prefetch_addr /* Prefetch abort handler */
    ldr pc, abort_addr     /* Data abort handler */
    ldr pc, reserved_addr  /* Reserved */
    ldr pc, irq_addr       /* IRQ (Interrupt request) handler */
    ldr pc, fiq_addr       /* FIQ (Fast interrupt req) handler */

reset_addr:    .word reset_handler
undef_addr:    .word reset_handler
swi_addr:      .word reset_handler
prefetch_addr: .word reset_handler
abort_addr:    .word reset_handler
reserved_addr: .word reset_handler
irq_addr:      .word irq_handler
fiq_addr:      .word reset_handler

_endvectors:

```

25

end

26