

# Operating Systems

Xinu – Clocks and Timers

1

## Timed Events

- Many systems and applications need to make use of timed events
  - network protocols, user interfaces, interfaces to external devices
- Many systems are structured around an asynchronous event paradigm
  - programmer defines a set of handler functions that are invoked based on a given event
  - arguably the OS functions like this internally
- Other systems provide a synchronous approach
  - The Xinu system provides delay and the programmer creates processes to schedule events

2

## Clocks and Timers

- Systems contain various sorts of clocks and timers
  - Processor clock – regular pulses to drive processor operation
  - Real-time clock – independent of processor operation, pulses in some fraction of a second and generates an interrupt for each pulse
  - Time of day clock – once set (or reset), computes elapsed time, does not interrupt
  - Interval timer – real time clock and counter that is modified per pulse, interrupts when a target is reached

3

## Timers and Interrupts

- A real time clock interrupts regularly
  - Can have significant overhead but may be required for real time applications
- An interval timer interrupts after a specified delay
  - Can be used to emulate a real time clock
  - Can vary the granularity of clock ticks
  - Set timer and decrement until zero, interrupt
  - Reset the timer when the interrupt fires

4

## Hardware

- The BeagleBone has an interval timer
  - The Xinu code configures the interval timer to interrupt after 1 millisecond
  - The interval timer can be configured to reset itself when an interrupt occurs
- The Galileo has a real-time clock
  - Configured at startup to interrupt every millisecond
- This provides consistent behavior on both platforms, with regularly occurring (1ms) timer interrupts

5

## Regular (Real-time) Timer Interrupts

- A real-time clock interrupts regularly without accumulating interrupts
  - When using a timer to emulate a clock, the same is true
- Clock-related interrupts must be serviced quickly or subsequent interrupts can be lost
- The processor must be able to execute many instructions between timer interrupts and thus must operate faster than the clock
  - How many cycles between interrupts at 1 GHz?
- Preference is given to the clock interrupt over I/O device interrupts

6

## Timed Delay

- The OS does two main things with time
  - timed delay
  - preemption
- The timed delay mechanism allows a process to sleep for a specified amount of time
- Sleeping removes a process from the ready queue and restores it to the ready queue after the specified amount of time
  - Once back in the ready queue, it executes according to the scheduling policy

7

## Preemption

- The process manager uses preemption to return control to the operating system
  - Without preemption, an errant process can disrupt the system in an infinite loop
- Time slicing – the OS may switch to another runnable process once the running process has used its share of time
- When there are several processes of the same priority, a call to `resched()` will place the current process at the end of the set in the ready list, and switches to the first process on the list
- Thus, processes of the same priority get scheduled round-robin

8

### Time Slicing

- Define a maximum time that a process may execute without allowing other processes to execute – a time *quantum*
- Short time slice values enable more sharing, but increase overhead
- All processes could potentially execute for their entire time slice, but generally few processes do, but rather get descheduled due to e.g. I/O

9

### Preemption Implementation

- *QUANTUM* specifies the number of clock ticks in a timeslice
- When a process is scheduled, *preempt* is set to *QUANTUM*
- When the clock tick occurs, *preempt* is decremented
- When it reaches 0, reset *preempt* to *QUANTUM* and call *resched()*
  - Two possibilities for resetting *preempt*

10

### Delay for Processes

- Maintain a list of processes that have requested delay
- Xinu uses a “delta list” to avoid searching the list for processes to wake
- A list in the *queuetab* ordered by time at which the process should be made ready to run again
- Rather than keeping absolute times, the list simply stores the difference between process wakeup times
- Only the first item of the list needs to be updated at each clock tick

11


### Delta List

- The key of the first process on a delta list specifies the number of clock ticks a process must delay beyond the current time; the key of each other process on a delta list specifies the number of clock ticks the process must delay beyond the preceding process on the list.

12

## Delta List

- Processes A, B, C and D requesting delays of 6, 12, 27 and 50 ticks (respectively) at roughly the same time (within one clock tick)
- These result in a delta list like this:



13

## Delta List Implementation

- The global variable *sleepq* points to the queue ID of the delta list
- When the clock ticks, the clock interrupt handler decrements the key of the first item in the list (if non-empty)
- When the key == 0, it is time for the process to be awakened
- The handler calls function *wakeup()*

14

## Delta List Implementation

- To maintain the list, *insertd()* takes a PID, Queue ID, and a delay in the argument *key*
- The key can be directly compared to the first key in the list
- But not to successive nodes as they specify delays relative to their predecessor
- insertd* subtracts the relative delays from key at each step so that it is comparable to *queuetab[next].qkey*
- Must also decrement the next key in the list by the key value being inserted

15

## insertd.c

```

/* insertd.c - insertd */
#include <xinu.h>

/*-----
 * insertd - Insert a process in delta list using delay as the key
 *----- */

status insertd(
    pid32    pid,          /* ID of process to insert */
    qid16    q,            /* ID of queue to use */
    int32    key,          /* Delay from "now" (in ms.) */
    int32    next,         /* Follows next through the list */
    int32    prev)         /* Runs through the delta list */
{
    if (isbadqid(q) || isbadpid(pid)) {
        return SYSERR;
    }

    prev = queuehead(q);
    next = queuetab[queuehead(q)].qnext;

```

16

## insertd.c

```
while ((next != queuehead(q)) && (queuehead[next].qkey <= key)) {
    key -= queuehead[next].qkey;
    prev = next;
    next = queuehead[next].qnext;
}

/* Insert new node between prev and next nodes */

queuehead[pid].qnext = next;
queuehead[pid].qprev = prev;
queuehead[pid].qkey = key;
queuehead[prev].qnext = pid;
queuehead[next].qprev = pid;

if (next != queuehead(q)) {
    queuehead[next].qkey -= key;
}

return OK;
}
```

17

## Putting a Process to Sleep

- Processes use the sleep() or sleepms() calls
  - sleep() is in seconds
  - sleepms() is in milliseconds
- Single implementation in which sleep() multiplies argument and calls sleepms()
  - Matches the clock interrupt
- Sleeping is a distinct state, PR\_SLEEP

18

## sleep.c

```
/* sleep.c - sleep sleeps */

#include <xinu.h>

#define MAXSECONDS 4294967 /* Max seconds per 32-bit msec */

/*
 * sleep - Delay the calling process n seconds
 */
syscall sleep(
    uint32 delay /* Time to delay in seconds */
)
{
    if (delay > MAXSECONDS) {
        return SYSERR;
    }
    sleepms(1000*delay);
    return OK;
}
```

19

## sleep.c

```
/*
 * sleepms - Delay the calling process n milliseconds
 */
syscall sleepms(
    uint32 delay /* Time to delay in msec. */
)
{
    intmask mask; /* Saved interrupt mask */

    mask = disable();
    if (delay == 0) {
        yield();
        restore(mask);
        return OK;
    }

    /* Delay calling process */

    if (insertd(currpid, sleepq, delay) == SYSERR) {
        restore(mask);
        return SYSERR;
    }

    proctab[currpid].prstate = PR_SLEEP;
    resched();
    restore(mask);
    return OK;
}
```

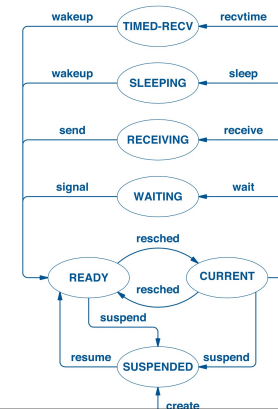
20

## Timed Message Receive

- Useful in networked applications and network protocols
  - Wake for lost message recovery
- Block for a message for some time and then unblock
  - Receive() OR Sleep(x)
- Place process in the sleep queue but in state TIMED-RECV
- If a message arrives, remove the process from the sleep queue
  - Detected by the send() call, which calls unsleep()

21

## Complete State Diagram



22

## unsleep.c

```

/* unsleep.c - unsleep */
#include <xinu.h>

/*
 * unsleep - Internal function to remove a process from the sleep
 * queue prematurely. The caller must adjust the delay
 * of successive processes.
 */
status unsleep(
    pid32 pid /* ID of process to remove */
)
{
    intmask mask; /* Saved interrupt mask */
    struct procent *prptr; /* Ptr to process' table entry */

    pid32 pidnext; /* ID of process on sleep queue */
    /* that follows the process */
    /* which is being removed */

    mask = disable();

    if (isbadpid(pid)) {
        restore(mask);
        return SYSERR;
    }

    /* Verify that candidate process is on the sleep queue */

```

23

## unsleep.c

```

/* Verify that candidate process is on the sleep queue */
prptr = &proctab[pid];
if ((prptr->prstate!=PR_SLEEP) && (prptr->prstate!=PR_RECTIM)) {
    restore(mask);
    return SYSERR;
}

/* Increment delay of next process if such a process exists */

pidnext = queuetab[pid].qnext;
if (pidnext < NPROC) {
    queuetab[pidnext].qkey += queuetab[pid].qkey;
}

getitem(pid); /* Unlink process from queue */
restore(mask);
return OK;
}

```

24

## wakeup.c

```
/* wakeup.c - wakeup */
#include <xinu.h>

/*-----
 * wakeup - Called by clock interrupt handler to awaken processes
 *-----
 */
void wakeup(void)
{
    /* Awaken all processes that have no more time to sleep */

    resched_cntl(DEFER_START);
    while (nonempty(sleepq) && (firstkey(sleepq) <= 0)) {
        ready(dequeue(sleepq));
    }

    resched_cntl(DEFER_STOP);
    return;
}
```

25

## clkhandler.c

```
/* clkhandler.c - clkhandler */
#include <xinu.h>

/*-----
 * clkhandler - high level clock interrupt handler
 *-----
 */
void clkhandler()
{
    static uint32 count1000 = 1000; /* variable to count 1000ms */
    volatile struct am335x_timer1ms *csrptr = 0x44E31000;
    /* Pointer to timer CSR */

    /* If there is no interrupt, return */

    if((csrptr->tisr & AM335X_TIMER1MS_TISR_OVF_IT_FLAG) == 0) {
        return;
    }

    /* Acknowledge the interrupt */

    csrptr->tisr = AM335X_TIMER1MS_TISR_OVF_IT_FLAG;

    /* Decrement 1000ms counter */

    count1000--;
```

26

## clkhandler.c

```
/* After 1 sec, increment clktime */

if(count1000 == 0) {
    clktime++;
    count1000 = 1000;
}

/* check if sleep queue is empty */

if(!isempty(sleepq)) {

    /* sleepq nonempty, decrement the key of */
    /* topmost process on sleepq */

    if((--queuetab[firstid(sleepq)].qkey) == 0) {
        wakeup();
    }
}

/* Decrement the preemption counter */
/* Reschedule if necessary */

if((--preempt) == 0) {
    preempt = QUANTUM;
    resched();
}
}
```

27

## resched.c

```
/* resched.c - resched, resched_cntl */
#include <xinu.h>

struct defer Defer;

/*-----
 * resched - Reschedule processor to highest priority eligible process
 *-----
 */
void resched(void) /* Assumes interrupts are disabled */
{
    struct proctent *ptold; /* Ptr to table entry for old process */
    struct proctent *ptnew; /* Ptr to table entry for new process */

    /* If rescheduling is deferred, record attempt and return */

    if (Defer.ndefers > 0) {
        Defer.attempt = TRUE;
        return;
    }

    /* Point to process table entry for the current (old) process */

    ptold = &proctab[currpid];
```

28

## resched.c

```
if (ptold->prstate == PR_CURR) { /* Process remains eligible */
    if (ptold->prprio > firstkey(readylist)) {
        return;
    }

    /* Old process will no longer remain current */

    ptold->prstate = PR_READY;
    insert(currpid, readylist, ptold->prprio);
}

/* Force context switch to highest priority ready process */

currpri = dequeue(readylist);
ptnew = &proctab[currpri];
ptnew->prstate = PR_CURR;
preempt = QUANTUM; /* Reset time slice for process */
ctxsw(&ptold->prstkptr, &ptnew->prstkptr);

/* Old process returns here when resumed */

return;
}
```

29

## clkinit.c

```
/* Set interrupt vector for clock to invoke clkhandler */
set_evec(AM335X_TIMER1MS_IRQ, (uint32)clkhandler);

sleepq = newqueue(); /* Allocate a queue to hold the delta */
/* list of sleeping processes */

preempt = QUANTUM; /* Set the preemption time */

clktime = 0; /* Start counting seconds */
```

30

## Summary

- Clock interrupts are critical to preemptive operating systems
- Timers exist in various forms but generally OSes use regular interrupts and construct timed events with software

31