

## Operating Systems

I/O

1

## Overview

- Devices and Controllers
- I/O Subsystem
- Device Drivers

2

## I/O Controllers

- A device controller is attached to the system or integrated into the motherboard or SoC
- The peripheral itself attaches to the controller
  - RS-232, SCSI, SATA, SAS, USB
- Convert a string of bits into bytes or blocks of bytes
  - Even disks are strings of bits
- The controller has registers mapped into memory
  - Read and written to control and check status

3

## I/O Ports vs Memory Mapping

- I/O ports in a dedicated namespace
- Accessed with special I/O instructions
  - `outb %al,$18`
- As opposed to memory space which is accessed with standard load, store, move

```
$ cat /proc/ioports
0000-0cf7 :PCI Bus
0000:00 0000-001f : dma1
0020-0021 : pic1
0040-0043 : timer0
0050-0053 : timer1
0060-0060 : keyboard
0064-0064 : keyboard
0070-0071 : rtc0
0080-008f : dma page reg
00a0-00a1 : pic2
00c0-00df : dma2
00f0-00ff : fpu
02f8-02ff : serial
0378-037a : parport0
03c0-03df : vga+
03f2-03f2 : floppy
```

4

## The I/O Subsystem

- Devices have a complicated low-level interface
  - Control and data registers mapped into memory
  - Header files and documentation to understand registers and bit fields
- Goal: provide a high-level interface so that programs don't have to be rewritten
  - I/O devices do mostly the same thing – they input, they output
  - Design: a small set of of abstract routines can encapsulate various devices

5

## I/O Interfaces

- Another purpose of the I/O interface is to protect shared I/O resources (devices, buffers)
- Safe and fair access
- Policies can be applied at the high-level interfaces and can be generalized over various devices
- The Unix abstraction is that “everything is a file”
  - Provides a namespace and an authorization mechanism

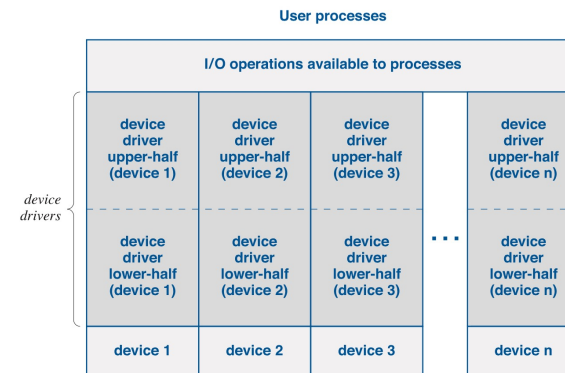
6

## Organization of the I/O Subsystem

- High level I/O functions to abstract the details of hardware and provide general entry points
  - Design challenge is to capture diversity of devices with a generic interface
- Device drivers interact with specific devices
- Drivers have an upper and lower half
- The upper half interacts with process requests
- The lower half responds to interrupts with handler functions
  - service interrupts, initiate new operations as necessary

7

## Organization of the I/O Subsystem



8

## I/O and Driver Abstractions

- Synchronous vs Asynchronous
  - Synchronous: the requesting process is blocked until I/O completes
    - easier to program
  - Asynchronous: the process can continue to execute – more control of overlap of communication and computation
- Asynchronous I/O interfaces must notify the process
  - Deliver a signal
  - Spawn a thread
  - Check (poll) the status of an I/O action, or read from a queue of completion actions

9

## POSIX AIO

- Allows applications to initiate one or more asynchronous I/O operations
  - signal, instantiate a thread, no notification
- `aio_read(struct aiocb *aiocbp), aio_write()`
- `aio_return()` – to check the return status of an AIO operation

```
struct aiocb { /* The order of these fields is implementation-dependent */
    int aio_fildes; /* File descriptor */
    off_t aio_offset; /* File offset */
    volatile void *aio_buf; /* Location of buffer */
    size_t aio_nbytes; /* Length of transfer */
    int aio_reqprio; /* Request priority */
    struct sigevent aio_sigevent; /* Notification method */
    .....
};
```

10

## I/O and Driver Abstractions

- Format and size of data transfers
  - Bytes, strings, blocks
- Block vs character interfaces
  - Look in `/dev` on your favorite \*nix system
  - The question is whether chunks of data (blocks) are independently addressable
- Buffering can be used to adapt between the two
- How much state is preserved between requests?
  - Specify a starting point and read successive blocks or specify a block with each read?

11

## Abstract I/O Interface

Operation	Purpose
<code>close</code>	Terminate use of a device
<code>control</code>	Perform operations other than data transfer
<code>getc</code>	Input a single byte of data
<code>init</code>	Initialize the device at system startup
<code>open</code>	Prepare the device for use
<code>putc</code>	Output a single byte of data
<code>read</code>	Input multiple bytes of data
<code>seek</code>	Move to specific data (usually a disk)
<code>write</code>	Output multiple bytes of data

12

## Open, Read, Write, Close

- Common paradigm (Xinu, Unix, Windows)
- Before a process can use a device, it must open it
  - Manage exclusive access
  - Check permissions
  - Set up state in system data structures
- Close when finished
  - Clean up state
  - The device could be powered down

13

## Control

- Control interface allows for configuration of device driver parameters
- Can also manage device-specific interactions that are not possible with the standard interfaces
  - Buffering or caching behavior
- `ioctl()` on Unix

14

## Binding Operations and Devices

- Abstract interfaces need to act on specific devices
- Must be mapped to device driver functions
- The OS provides a virtual I/O environment, passing operations through to devices via drivers
- Unix embeds devices in the filesystem, providing names to specific devices
- General-purpose OSes construct this dynamically, but embedded systems often statically configure it

15

## Device Names in Xinu

- Specify a set of devices when the system is configured
- Assign an integer device descriptor
- For instance, `CONSOLE` is device 0
- Programs don't need to be rewritten when devices change, but the system does need to be reconfigured and recompiled

16

## Xinu's Device Switch Table

- The OS must forward I/O operations to the correct driver function
- The device ID is used as an index into a table of device-specific functions
- Each entry in the table contains information about the device and function pointers to functions that implement operations
- To write to a device, find the device entry and invoke the specific write function
- Xinu's approach is simple but is fundamentally the same as e.g. Unix

17

## Xinu Example

	open	close	read	write	getc	
CONSOLE	conopen	conclose	conread	conwrite	congetc	
ETHER	ethopen	ethclose	ethread	ethwrite	ethgetc	...
DISK	dskopen	dskclose	dskread	dskwrite	dskgetc	
			⋮			

- Uniform interface hiding the differences of underlying hardware

18

## Multiple Instances of a Device

- Multiple instances of a device can share a driver
- Multiple instances in the device table that are largely the same, differing in only a few aspects
- Each instance will have its own control and status registers
- Can also be distinguished by the “minor” device number

19

## Device Table Entry

```
/* Device table entry */
struct dentry {
    int32  dvnum;
    int32  dvminor;
    char   *dvname;
    devcall (*dvinit) (struct dentry *);
    devcall (*dvopen) (struct dentry *, char *, char *);
    devcall (*dvclose) (struct dentry *);
    devcall (*dvread) (struct dentry *, void *, uint32);
    devcall (*dvwrite) (struct dentry *, void *, uint32);
    devcall (*dvseek) (struct dentry *, int32);
    devcall (*dvgetc) (struct dentry *);
    devcall (*dvputc) (struct dentry *, char);
    devcall (*dvcntl) (struct dentry *, int32, int32, int32);
    void    *dvcsr;
    void    (*dvintr) (void);
    byte    dvirq;
};
```

20

## Some Devices

```
extern struct dentry devtab[]; /* one entry per device */

/* Device name definitions */

#define CONSOLE 0 /* type tty */
#define NULLDEV 1 /* type null */
#define ETHER0 2 /* type eth */
#define NAMESPACE 3 /* type nam */
#define RDISK 4 /* type rds */
#define RAM0 5 /* type ram */
#define RFILESYS 6 /* type rfs */
#define RFILE0 7 /* type rfl */
#define RFILE1 8 /* type rfl */
#define RFILE2 9 /* type rfl */
#define RFILE3 10 /* type rfl */
#define RFILE4 11 /* type rfl */
#define RFILE5 12 /* type rfl */
#define RFILE6 13 /* type rfl */
```

21

## read()

```
syscall read(
    did32 descrp, /* Descriptor for device */
    char *buffer, /* Address of buffer */
    uint32 count /* Length of buffer */
)
{
    intmask mask; /* Saved interrupt mask */
    struct dentry *devptr; /* Entry in device switch table */
    int32 retval; /* Value to return to caller */

    mask = disable();
    if (isbaddev(descrp)) {
        restore(mask);
        return SYSERR;
    }
    devptr = (struct dentry *) &devtab[descrp];
    retval = (*devptr->dread) (devptr, buffer, count);
    restore(mask);
    return retval;
}
```

22

## control()

```
syscall control(
    did32 descrp, /* Descriptor for device */
    int32 func, /* Specific control function */
    int32 arg1, /* Specific argument for func */
    int32 arg2 /* Specific argument for func */
)
{
    intmask mask; /* Saved interrupt mask */
    struct dentry *devptr; /* Entry in device switch table */
    int32 retval; /* Value to return to caller */

    mask = disable();
    if (isbaddev(descrp)) {
        restore(mask);
        return SYSERR;
    }
    devptr = (struct dentry *) &devtab[descrp];
    retval = (*devptr->dvcntl) (devptr, func, arg1, arg2);
    restore(mask);
    return retval;
}
```

23

## open() and close()

- Implemented identically to read() and control()
- Explicit opening and closing allows the system to maintain a *reference count* of processes using a device
  - Again, the system may power a device down when not in use

24

## Null Entries in Devtab

- Note that each of the high-level functions calls the device-specific function without checking its validity
- Not all operations make sense on all devices
  - You can't seek on the console, or getc() on a network device
- ionull() returns OK
- ioerr() returns SYSERR

25

## Initialization

- General operating systems can dynamically initialize devices
  - recall the discussion of USB devices
- Embedded systems like Xinu use static configuration
- Xinu specifies devices and functions in a file called Configuration, and generates a C file and a header with appropriate values

26

## Configuration

```
/* Entries for a device specify the functions that handle each of the */
/* high-level I/O operations as follows: */
/* */
/* -i init -o open -c close */
/* -r read -w write -s seek */
/* -g getc -p putc -n control */
/* -intr int_hdlr -csr csr -irq irq */
/* */
/* ===== */

/* Type Declarations for both real- and pseudo- device types */

/* type of a null device */
null:
on nothing
-i ionull -o ionull -c ionull
-r ionull -g ionull -p ionull
-w ionull -s ioerr
```

27

## Configuration

```
/* type of a tty device */
tty:
on uart
-i ttyinit -o ionull -c ionull
-r ttyread -g ttygetc -p ttyputc
-w ttywrite -s ioerr -n ttycontrol
-intr ttyhandler

/* type of a ethernet device */
eth:
on am335x_eth
-i ethinit -o ioerr -c ioerr
-r ethread -g ioerr -p ioerr
-w ethwrite -s ioerr -n ethcontrol
-intr ethhandler
```

28

## conf.c

```
struct dentry devtab[NDEVS] =
{
/**
 * Format of entries is:
 * dev-number, minor-number, dev-name,
 * init, open, close,
 * read, write, seek,
 * getc, putc, control,
 * dev-csr-address, intr-handler, irq
 */

/* CONSOLE is tty */
{ 0, 0, "CONSOLE",
  (void *)ttyinit, (void *)ionull, (void *)ionull,
  (void *)ttyread, (void *)ttywrite, (void *)ioerr,
  (void *)ttygetc, (void *)ttyputc, (void *)ttycontrol,
  (void *)0x44e09000, (void *)ttyhandler, 72 },
```

29

## conf.c

```
/* CONSOLE is tty */
{ 0, 0, "CONSOLE",
  (void *)ttyinit, (void *)ionull, (void *)ionull,
  (void *)ttyread, (void *)ttywrite, (void *)ioerr,
  (void *)ttygetc, (void *)ttyputc, (void *)ttycontrol,
  (void *)0x44e09000, (void *)ttyhandler, 72 },

/* NULLDEV is null */
{ 1, 0, "NULLDEV",
  (void *)ionull, (void *)ionull, (void *)ionull,
  (void *)ionull, (void *)ionull, (void *)ioerr,
  (void *)ionull, (void *)ionull, (void *)ioerr,
  (void *)0x0, (void *)ioerr, 0 },

/* ETHER0 is eth */
{ 2, 0, "ETHER0",
  (void *)ethinit, (void *)ioerr, (void *)ioerr,
  (void *)ethread, (void *)ethwrite, (void *)ioerr,
  (void *)ioerr, (void *)ioerr, (void *)ethcontrol,
  (void *)0x0, (void *)ethhandler, 0 },
```

30

## Embedded Linux - DeviceTree

- Embedded Linux uses something called DeviceTree
- A Flattened Device Tree (FDT) is shipped so that a kernel image can be configured appropriately for hardware
- Required for Linux on new ARM SoCs
- Identifies the type of CPU and describes devices in the system very similarly to what we have discussed

31

## Summary

- Complex operating systems have more functionality between the abstract interfaces and the device drivers
- Caching
- Security and policy
- This same abstraction / indirection mechanism forms the basis

32