

Operating Systems

DMA Devices and the Ethernet Driver

Device Drivers and DMA

- We have seen the structure of device drivers, with upper and lower halves interfacing with shared data structures
- Next we examine devices that can transfer data to memory directly, and look at their device drivers
 - DMA – Direct Memory Access
- One necessity from the driver perspective is separation of request blocks and their buffers

DMA and Buffers

- The goal of DMA is to reduce interrupts and allow devices to transfer data to memory while the CPU is busy with other tasks
 - Reduction of overhead
 - Concurrent activities
- A bus can only transfer a bus-width of data at a time
 - One word or a double word
 - One byte a bit at a time for serial interfaces
- Block or frame oriented devices need to transfer multiple words for every transaction
- DMA allows the transfer of an entire chunk, or multiple chunks, before generating an interrupt

DMA and Buffers

- For a write operation, place data in a buffer and create a write request to pass to the device
 - The request contains a pointer to the buffer to be written
- For read, post a read request to the device, with a pointer to an empty buffer and
 - For a disk, specify the block to be read
 - For a network device, the read request is left pending until the device produces data

DMA and Buffers

- While system software is processing a data element, the device should be able to operate on others
 - Multiple data buffers posted
 - This gives rise to a
- Producer / consumer problem with a familiar solution
 - Have a logical ring of buffer elements
- Control access to the buffers based on whether they are empty or full

DMA Operations

- Post a linked list of operations
 - The device can follow the linked list of operation requests
- Generates an interrupt on completion of one operation, but continues to the next without waiting for the processor
- As long as requests remain on the list, the device can continue to operate
 - Multiple posted receive operations allow the network to continue to receive packets and place them in the ring buffer

DMA Operations

- Must have an EMPTY/FULL status bit that the hardware understands
- Hardware traverses the list, marking FULL or EMPTY
 - Depending on ingress or egress
- On input, if the hardware goes around the ring and all buffers are FULL, set an overflow bit
- On output, consume and transmit all FULL buffers

Ethernet Driver with DMA

- Separate DMA engines for input and output
- As in the serial case, create input and output ring buffers
- Separate input and output control registers get pointers to input and output request lists
- Mark all buffers as empty, post receive operations for input buffers
- The Galileo Ethernet device has a single interrupt vector, so as in the serial case, the handler must determine whether the interrupts indicates and input or output operation

Ethernet Hardware

- Ethernet defines a Media Access Control (MAC) layer
 - The lower half of the “Data Link” layer in the OSI network model
 - The upper half is the Logical Link Control (LLC) layer
- Below the MAC is the Physical layer, or PHY
- The MAC and the PHY are decoupled in various ways so that the MAC operation is independent of e.g. copper or fiber optic physical layers

Ethernet Hardware

- The system interfaces with the MAC using the Media-Independent Interface (MII)
- The Intel Galileo (Quark SoC) exposes MII operations with the General MII (GMII) interface
 - Specifically a simplified subset known as the Reduced MII (RMII) interface
- The BeagleBone Black exposes a Management Data Input/Output (MDIO) interface
- Communication occurs from the Serial Management Interface (SMI) to the PHY chip
- These details matter in the driver, as the driver must understand the Control/Status registers (CSR) of the devices

quark_eth.h

```
/* quark_eth.h */
```

```
/* Definitions for Intel Quark Ethernet */
```

```
#define INTEL_ETH_QUARK_PCI_DID 0x0937 /* MAC PCI Device ID */
```

```
#define INTEL_ETH_QUARK_PCI_VID 0x8086 /* MAC PCI Vendor ID */
```



```
struct eth_q_csreg {  
    uint32_t maccr; /* MAC Configuration Register */  
    uint32_t macff; /* MAC Frame Filter Register */  
    uint32_t hthr; /* Hash Table High Register */  
    uint32_t htlr; /* Hash Table Low Register */  
    uint32_t gmiar; /* GMII Address Register */  
    uint32_t gmidr; /* GMII Data Register */  
    uint32_t fcr; /* Flow Control Register */  
    uint32_t vlantag; /* VLAV Tag Register */  
    uint32_t version; /* Version Register */  
    uint32_t debug; /* Debug Register */  
    uint32_t res1[4]; /* Skipped Addresses */  
    uint32_t ir; /* Interrupt Register */  
};
```

quark_eth.h

```
uint32_t imr;      /* Interrupt Mask Register      */
uint32_t macaddr0h; /* MAC Address0 High Register    */
uint32_t macaddr0l; /* MAC Address0 Low Register     */
uint32_t res2[46];
uint32_t mmccr;    /* MAC Management Counter Cntl Register */
uint32_t mmcrvcir; /* MMC Receive Interrupt Register */
uint32_t mmctxir;  /* MMC Transmit Interrupt Register */
uint32_t res3[957]; /* Skipped Addresses              */
uint32_t bmr;      /* Bus Mode Register              */
uint32_t tpdr;     /* Transmit Poll Demand Register  */
uint32_t rpdr;     /* Receive Poll Demand Register   */
uint32_t rdla;     /* Receive Descriptor List Addr   */
uint32_t tdla;     /* Transmit Descriptor List Addr  */
uint32_t sr;       /* Status Register                */
uint32_t omr;      /* Operation Mode Register        */
uint32_t ier;      /* Interrupt Enable Register      */
};
```


Rings and Buffers

- With separate request descriptors and buffers, the “rings” consist of the descriptors
 - The serial driver combined them
- The input and output rings are linked lists of descriptors with a status word indicating whether the associated buffer is empty or full



Descriptor Ring

- Each node contains
 - A pointer to the buffer
 - Status word
 - A pointer to the next node in the list
- The driver uses `getmem()` to allocate contiguous memory and links the nodes together



quark_eth.h

```
/* Quark Ethernet Transmit Descriptor */
```

```
struct eth_q_tx_desc {  
    uint32  ctrlstat; /* Control and status */  
    uint16  buf1size; /* Size of buffer 1 */  
    uint16  buf2size; /* Size of buffer 2 */  
    uint32  buffer1;  /* Address of buffer 1  */  
    uint32  buffer2;  /* Address of buffer 2  */  
};
```

```
#define ETH_QUARK_TDCS_OWN  0x80000000 /* Descrip. owned by DMA */  
#define ETH_QUARK_TDCS_IC  0x40000000 /* Int on Completion  */  
#define ETH_QUARK_TDCS_LS  0x20000000 /* Last Segment      */  
#define ETH_QUARK_TDCS_FS  0x10000000 /* First Segment     */  
#define ETH_QUARK_TDCS_TER  0x00200000 /* Transmit End of Ring */  
#define ETH_QUARK_TDCS_ES  0x00008000 /* Error Summary     */
```

quark_eth.h

```
/* Quark Ethernet Receive Descriptor */
```

```
struct eth_q_rx_desc {  
    uint32  status; /* Desc status word */  
    uint16  buf1size; /* Size of buffer 1 */  
    uint16  buf2size; /* Size of buffer 2 */  
    uint32  buffer1; /* Address of buffer 1 */  
    uint32  buffer2; /* Address of buffer 2 */  
};  
  
#define rdctl1  buf1size /* Buffer 1 size field has control bits too */  
#define rdctl2  buf2size /* Buffer 2 size field has control bits too */  
  
#define ETH_QUARK_RDST_OWN  0x80000000 /* Descrip. owned by DMA*/  
#define ETH_QUARK_RDST_ES  0x00008000 /* Error Summary */  
#define ETH_QUARK_RDST_FS  0x00000200 /* First Segment */  
#define ETH_QUARK_RDST_LS  0x00000100 /* Last segment */  
#define ETH_QUARK_RDST_FTETH 0x00000020 /* Frame Type = Ethernet*/
```


Ethernet Driver Control Block (ether.h)

```
/* ether.h */
```

```
/* Ethernet packet format:
```

```

+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| Dest. MAC (6) | Src. MAC (6) | Type (2) |      Data (46-1500)...      |
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
*/
```

```
#define ETH_ADDR_LEN 6 /* Len. of Ethernet (MAC) addr. */
```

```
typedef unsigned char Eaddr[ETH_ADDR_LEN]; /* Physical Ethernet address*/
```

```
/* Ethernet packet header */
```

```
struct etherPkt {
    byte dst[ETH_ADDR_LEN]; /* Destination Mac address */
    byte src[ETH_ADDR_LEN]; /* Source Mac address */
    uint16 type; /* Ether type field */
    byte data[1]; /* Packet payload */
};
```

```
#define ETH_HDR_LEN 14 /* Length of Ethernet packet */
/* header */
```

ether.h

```
/* Ethernet DMA buffer sizes */
```

```
#define ETH_MTU      1500 /* Maximum transmission unit */  
#define ETH_VLAN_LEN  4 /* Length of Ethernet vlan tag */  
#define ETH_CRC_LEN   4 /* Length of CRC on Ethernet */  
    /*      frame      */
```

```
#define ETH_MAX_PKT_LEN ( ETH_HDR_LEN + ETH_VLAN_LEN + ETH_MTU )
```

```
#define ETH_BUF_SIZE    2048 /* A multiple of 16 greater */  
    /*      than the max packet */  
    /*      size (for cache alignment) */
```

```
/* State of the Ethernet interface */
```

```
#define ETH_STATE_FREE    0 /* Control block is unused */  
#define ETH_STATE_DOWN    1 /* Interface is inactive */  
#define ETH_STATE_UP      2 /* Interface is currently active*/
```

```
/* Ethernet device control functions */
```

```
#define ETH_CTRL_GET_MAC    1 /* Get the MAC for this device */  
#define ETH_CTRL_ADD_MCAST  2 /* Add a multicast address */  
#define ETH_CTRL_REMOVE_MCAST 3 /* Remove a multicast address */
```

Ethernet Driver Control Block (ether.h)

```
/* Control block for Ethernet device */
```

```
struct ethcblk {  
    byte  state;      /* ETH_STATE_... as defined above */  
    struct dentry  *phy; /* physical eth device for Tx DMA */  
    byte  type;       /* NIC type_... as defined above */  
  
    /* Pointers to associated structures */  
  
    struct dentry  *dev; /* Address in device switch table */  
    void  *csr; /* Control and status register address */  
    uint32  pcidev; /* PCI device number */  
    uint32  iobase; /* I/O base from config */  
    uint32  flashbase; /* Flash base from config */  
    uint32  membase; /* Memory base for device from config */  
  
    void  *rxRing; /* Ptr to array of recv ring descriptors */  
    void  *rxBufs; /* Ptr to Rx packet buffers in memory */  
    uint32  rxHead; /* Index of current head of Rx ring */  
    uint32  rxTail; /* Index of current tail of Rx ring */  
    uint32  rxRingSize; /* Size of Rx ring descriptor array */  
    uint32  rxIrq; /* Count of Rx interrupt requests */  
}
```



```

void    *txRing;    /* Ptr to array of xmit ring descriptors*/
void    *txBufs;    /* Ptr to Tx packet buffers in memory */
uint32  txHead;     /* Index of current head of Tx ring */
uint32  txTail;     /* Index of current tail of Tx ring */
uint32  txRingSize; /* Size of Tx ring descriptor array */
uint32  txIrq;      /* Count of Tx interrupt requests      */

uint8 devAddress[ETH_ADDR_LEN]; /* MAC address      */

uint8 addrLen; /* Hardware address length      */
uint16 mtu;     /* Maximum transmission unit (payload) */

uint32  errors; /* Number of Ethernet errors      */
uint32  isem;   /* Semaphore for Ethernet input    */
uint32  osem;   /* Semaphore for Ethernet output   */
uint16  istart; /* Index of next packet in the ring */

int16  inPool; /* Buffer pool ID for input buffers */
int16  outPool; /* Buffer pool ID for output buffers */

int16  proms; /* Nonzero => promiscuous mode */

int16  ed_mcset; /* Nonzero => multicast reception set */
int16  ed_mcc; /* Count of multicast addresses */
uint8 ed_mca[ETH_NUM_MCAST]; /* Array of multicast addrs */

```

```
};
```

Reading from the Ethernet

- The DMA engine is configured to fill input buffers
- While the DMA engine can read the input descriptor ring, it cannot manipulate semaphores
- To read a packet, a process waits on a semaphore
 - Initial value is 0, so the first read will block
- When a packet is available, the interrupt handler will signal the semaphore and the process can read from the buffer

ethread.c

```
/*-----  
 * ethread - Read an incoming packet on Intel Quark Ethernet  
 *-----  
 */  
  
devcall ethread (  
    struct dentry *devptr, /* Entry in device switch table */  
    char *buf, /* Buffer for the packet */  
    int32 len /* Size of the buffer */  
)  
{  
    struct ethcblk *ethptr; /* Ethertab entry pointer */  
    struct eth_q_rx_desc *rdescptr; /* Pointer to the descriptor */  
    struct netpacket *pktptr; /* Pointer to packet */  
    uint32 framelen = 0; /* Length of the incoming frame */  
    bool8 valid_addr;  
    int32 i;  
  
    ethptr = &ethertab[devptr->dvminor];  
  
    while(1) {  
  
        /* Wait until there is a packet in the receive queue */  
  
        wait(ethptr->isem);  

```

ethread.c

```
/* Point to the head of the descriptor list */

rdescptr = (struct eth_q_rx_desc *)ethptr->rxRing +
            ethptr->rxHead;
pktptr = (struct netpacket*)rdescptr->buffer1;

/* See if destination address is our unicast address */

if(!memcmp(pktptr->net_ethdst, ethptr->devAddress, 6)) {
    valid_addr = TRUE;

/* See if destination address is the broadcast address */

} else if(!memcmp(pktptr->net_ethdst,
                  NetData.ethbcast,6)) {
    valid_addr = TRUE;

/* For multicast addresses, see if we should accept */

    } else {
valid_addr = FALSE;
for(i = 0; i < (ethptr->ed_mcc); i++) {
    if(memcmp(pktptr->net_ethdst,
              ethptr->ed_mca[i], 6) == 0){
        valid_addr = TRUE;
        break;
    }
}
}
}
```

ethread.c

```
if(valid_addr == TRUE){ /* Accept this packet */

    /* Get the length of the frame */

    framelen = (rdescptr->status >> 16) & 0x00003FFF;

    /* Only return len characters to caller */

    if(framelen > len) {
        framelen = len;
    }

    /* Copy the packet into the caller's buffer */

    memcpy(buf, (void*)rdescptr->buffer1, framelen);
}

/* Increment the head of the descriptor list */

ethptr->rxHead += 1;
if(ethptr->rxHead >= ETH_QUARK_RX_RING_SIZE) {
    ethptr->rxHead = 0;
}
```


ethread.c

```
/* Reset the descriptor to max possible frame len */
```

```
rdescptr->buf1size = sizeof(struct netpacket);
```

```
/* If we reach the end of the ring, mark the descriptor */
```

```
if(ethptr->rxHead == 0) {  
    rdescptr->rdctl1 |= (ETH_QUARK_RDCTL1_RER);  
}
```

```
/* Indicate that the descriptor is ready for DMA input */
```

```
rdescptr->status = ETH_QUARK_RDST_OWN;
```

```
if(valid_addr == TRUE) {  
    break;  
}  
}
```

```
/* Return the number of bytes returned from the packet */
```

```
return framelen;
```

```
}
```

ethwrite.c

```
/*-----  
 * ethwrite - enqueue packet for transmission on Intel Quark Ethernet  
 *-----  
 */  
  
devcall ethwrite (  
    struct dentry *devptr, /* Entry in device switch table */  
    char *buf, /* Buffer that holds a packet */  
    int32 len /* Length of the packet */  
)  
{  
    struct ethcblk *ethptr; /* Pointer to control block */  
    struct eth_q_csreg *csrptr; /* Address of device CSRs */  
    volatile struct eth_q_tx_desc *descptr; /* Ptr to descriptor */  
    uint32 i; /* Counts bytes during copy */  
  
    ethptr = &ethertab[devptr->dvminor];  
  
    csrptr = (struct eth_q_csreg *)ethptr->csr;  
  
    /* Wait for an empty slot in the transmit descriptor ring */  
  
    wait(ethptr->osem);
```

ethwrite.c

```
/* Point to the tail of the descriptor ring */
```

```
descptr = (struct eth_q_tx_desc *)ethptr->txRing + ethptr->txTail;
```

```
/* Increment the tail index and wrap, if needed */
```

```
ethptr->txTail += 1;  
if(ethptr->txTail >= ethptr->txRingSize) {  
    ethptr->txTail = 0;  
}
```

```
/* Add packet length to the descriptor */
```

```
descptr->buf1size = len;
```

```
/* Copy packet into the buffer associated with the descriptor */
```

```
for(i = 0; i < len; i++) {  
    *((char *)descptr->buffer1 + i) = *((char *)buf + i);  
}
```

ethwrite.c

```
/* Mark the descriptor if we are at the end of the ring */
```

```
if(ethptr->txTail == 0) {  
    descptr->ctrlstat = ETH_QUARK_TDCS_TER;  
} else {  
    descptr->ctrlstat = 0;  
}
```

```
/* Initialize the descriptor */
```

```
descptr->ctrlstat |=  
    (ETH_QUARK_TDCS_OWN | /* The desc is owned by DMA */  
     ETH_QUARK_TDCS_IC  | /* Interrupt after transfer */  
     ETH_QUARK_TDCS_LS  | /* Last segment of packet   */  
     ETH_QUARK_TDCS_FS); /* First segment of packet  */
```

```
/* Un-suspend DMA on the device */
```

```
csrptr->tpdr = 1;
```

```
return OK;
```

```
}
```

ethhandler.c

```
/*-----  
 * ethhandler - Interrupt handler for Intel Quark Ethernet  
 *-----  
 */  
  
interrupt ethhandler(void)  
{  
    struct ethcbk *ethptr; /* Ethertab entry pointer */  
    struct eth_q_csreg *csrptr; /* Pointer to Ethernet CRSs */  
    struct eth_q_tx_desc *tdescptr; /* Pointer to tx descriptor */  
    struct eth_q_rx_desc *rdescptr; /* Pointer to rx descriptor */  
    volatile uint32 sr; /* Copy of status register */  
    uint32 count; /* Variable used to count pkts */  
  
    ethptr = &ethertab[devtab[ETHER0].dvminor];  
  
    csrptr = (struct eth_q_csreg *)ethptr->csr;  
  
    /* Copy the status register into a local variable */  
  
    sr = csrptr->sr;  
  
    /* If there is no interrupt pending, return */  
  
    if((csrptr->sr & ETH_QUARK_SR_NIS) == 0) {  
        return;  
    }  
}
```

```
/* Acknowledge the interrupt */
```

ethhandler.c

```
csrptr->sr = sr;
```

```
/* Check status register to figure out the source of interrupt */
```

```
if (sr & ETH_QUARK_SR_TI) { /* Transmit interrupt */
```

```
    /* Pointer to the head of transmit desc ring */
```

```
    tdescptr = (struct eth_q_tx_desc *)ethptr->txRing +  
                ethptr->txHead;
```

```
    /* Start packet count at zero */
```

```
    count = 0;
```

```
    /* Repeat until we process all the descriptor slots */
```

```
    while(ethptr->txHead != ethptr->txTail) {
```

```
        /* If the descriptor is owned by DMA, stop here */
```

```
        if(tdescptr->ctrlstat & ETH_QUARK_TDCS_OWN) {  
            break;  
        }
```

ethhandler.c

```
/* Descriptor was processed; increment count */
```

```
count++;
```

```
/* Go to the next descriptor */
```

```
tdescptr += 1;
```

```
/* Increment the head of the transmit desc ring */
```

```
ethptr->txHead += 1;
```

```
if(ethptr->txHead >= ethptr->txRingSize) {
```

```
    ethptr->txHead = 0;
```

```
    tdescptr = (struct eth_q_tx_desc *)
```

```
        ethptr->txRing;
```

```
}
```

```
}
```

```
/* 'count' packets were processed by DMA, and slots are */
```

```
/* now free; signal the semaphore accordingly */
```

```
signaln(ethptr->osem, count);
```

```
}
```

ethhandler.c

```
if(sr & ETH_QUARK_SR_RI) { /* Receive interrupt */

    /* Get the pointer to the tail of the receive desc list */

    rdescptr = (struct eth_q_rx_desc *)ethptr->rxRing +
                ethptr->rxTail;

    count = 0; /* Start packet count at zero */

    /* Repeat until we have received      */
    /* maximum no. packets that can fit in queue */

    while(count <= ethptr->rxRingSize) {

        /* If the descriptor is owned by the DMA, stop */

        if(rdescptr->status & ETH_QUARK_RDST_OWN) {
            break;
        }
    }
}
```


ethhandler.c

```
/* Descriptor was processed; increment count */
```

```
count++;
```

```
/* Go to the next descriptor */
```

```
rdescptr += 1;
```

```
/* Increment the tail index of the rx desc ring */
```

```
ethptr->rxTail += 1;
```

```
if(ethptr->rxTail >= ethptr->rxRingSize) {
```

```
    ethptr->rxTail = 0;
```

```
    rdescptr = (struct eth_q_rx_desc *)
```

```
        ethptr->rxRing;
```

```
}
```

```
}
```

```
/* 'count' packets were received and are available, */
```

```
/*   so signal the semaphore accordingly   */
```

```
signaln(ethptr->isem, count);
```

```
}
```

```
return;
```

```
}
```

ethhandler.c (BBB)

```
/*-----  
 * ethhandler - TI AM335X Ethernet Interrupt Handler  
 *-----  
 */  
  
interrupt ethhandler (  
    uint32 xnum /* IRQ number */  
)  
{  
    struct eth_a_csreg *csrptr; /* Ethernet CSR pointer */  
    struct eth_a_tx_desc *tdescptr; /* Tx desc pointer */  
    struct eth_a_rx_desc *rdescptr; /* Rx desc pointer */  
    struct ethcblk *ethptr = &ethertab[0]; /* Ethernet ctl blk ptr */  
  
    csrptr = (struct eth_a_csreg *)ethptr->csr;  
  
    if(xnum == ETH_AM335X_TXINT) { /* Transmit interrupt */  
  
        /* Get pointer to first desc in queue */  
  
        tdescptr = (struct eth_a_tx_desc *)ethptr->txRing +  
            ethptr->txHead;  
  
        /* Defer scheduling until all descs are processed */  
  
        resched_cntl(DEFER_START);  
    }  
}
```

```
while(semcount(ethptr->osem) < (int32)ethptr->txRingSize) {
```

ethhandler.c (BBB)

```
/* If desc owned by DMA, check if we need to */  
/* Restart the transmission */
```

```
if(tdescptr->stat & ETH_AM335X_TDS_OWN) {  
    if(csrptr->stateram->tx_hdp[0] == 0) {  
        csrptr->stateram->tx_hdp[0] =  
            (uint32)tdescptr;  
    }  
    break;  
}
```

```
/* Acknowledge the interrupt */
```

```
csrptr->stateram->tx_cp[0] = (uint32)tdescptr;
```

```
/* Increment the head index of the queue */  
/* And go to the next descriptor in queue */
```

```
ethptr->txHead++;  
tdescptr++;  
if(ethptr->txHead >= ethptr->txRingSize) {  
    ethptr->txHead = 0;  
    tdescptr = (struct eth_a_tx_desc *)  
        ethptr->txRing;  
}
```

ethhandler.c (BBB)

```
/* Signal the output semaphore */
```

```
    signal(ethptr->osem);  
}
```

```
/* Acknowledge the transmit interrupt */
```

```
csrptr->cpdma->eoi_vector = 0x2;
```

```
/* Resume rescheduling */
```

```
resched_cntl(DEFER_STOP);  
}
```

ethhandler.c (BBB)

```
else if(xnum == ETH_AM335X_RXINT) { /* Receive interrupt */

    /* Get the pointer to last desc in the queue */

    rdescptr = (struct eth_a_rx_desc *)ethptr->rxRing +
                ethptr->rxTail;

    /* Defer scheduling until all descriptors are processed */

    resched_cntl(DEFER_START);

    while(semcount(ethptr->isem) < (int32)ethptr->rxRingSize) {

        /* Check if we need to restart the DMA */

        if(rdescptr->stat & ETH_AM335X_RDS_OWN) {
            if(csrrptr->stateram->rx_hdp[0] == 0) {
                csrrptr->stateram->rx_hdp[0] =
                    (uint32)rdescptr;
            }
            break;
        }
    }
```

ethhandler.c (BBB)

```
/* Acknowledge the interrupt */
```

```
csrptr->stateram->rx_cp[0] = (uint32)rdescptr;
```

```
/* Increment the tail index of the queue */
```

```
/* And go to the next descriptor in the queue */
```

```
ethptr->rxTail++;
```

```
rdescptr++;
```

```
if(ethptr->rxTail >= ethptr->rxRingSize) {
```

```
    ethptr->rxTail = 0;
```

```
    rdescptr = (struct eth_a_rx_desc *)
```

```
        ethptr->rxRing;
```

```
}
```

```
/* Signal the input semaphore */
```

```
signal(ethptr->isem);
```

```
}
```

ethhandler.c (BBB)

```
/* Acknowledge the receive interrupt */
```

```
csrptr->cpdma->eoi_vector = 0x1;
```

```
/* Resume rescheduling */
```

```
resched_cntl(DEFER_STOP);
```

```
}
```

```
}
```

Improving Performance

- Interrupt Coalescing
- Interrupt Affinity
- Large Send Offload
- Reduce Data Copying

Improving Performance

