

Operating Systems

Filesystem Examples and Case Studies

Filesystem Examples

- Ext2
- Ext3
- Ext4
- XFS
- JFFS
- WAFL
- ZFS

Ext2

- The Second Extended Filesystem
 - Default Linux filesystem for 7 years
 - Written in 1993 to replace the “Extended Filesystem”
 - Borrows heavily from the Berkeley Fast File System

Ext2 Features

- Variable Block Size(.5k, 1k, 2k and 4k)
- 2TB maximum file size
- 16TB maximum filesystem size

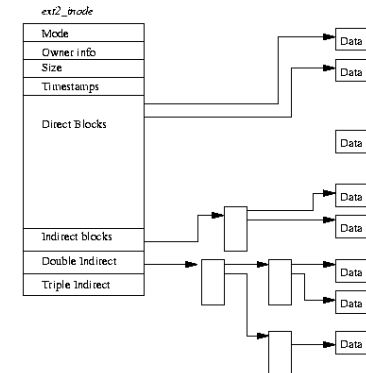
Ext2 On-Disk Layout

- Block Groups
 - Disk broken up into groups to keep related data “close”
 - Contain copies of “filesystem control information”

Super Block	Group Desc	Block Bitmap	Inode Bitmap	Inodes	Data
-------------	------------	--------------	--------------	--------	------

Ext2 On-Disk Layout

- Inodes



Ext2 On-Disk Layout

- Directories
 - Linked list of directory entries
 - Directory entries can be variable length
 - allows for long filenames

Inode Number	entry length	filename length	filename
--------------	--------------	-----------------	----------

Ext2 Performance Optimizations

- Allocates data blocks in same block group as the inode that refers to them
- Preallocates blocks on write

Ext2 Limitations

- No filesystem reliability
 - file system checks unwieldy with larger drives
- Directory scalability
 - $O(n)$ time to find a file in the directory

Ext3

- Adds journaling capabilities to ext2
 - removes the need to perform file system checks if one shuts down improperly
- Maintains backward compatible with Ext2

Ext3 – Journaling

- Adds a file called .journal to the ext2 filesystem
 - Allows for easy conversion from ext2 to ext3 and vice versa
- When the filesystem is mounted as ext3, the .journal file is used for the transaction log
- When the filesystem is unmounted properly, it looks just like an ext2 filesystem.
- When not unmounted properly, it is marked so that ext2 knows not to mount it.

Ext3 Features

- Multiple types of journaling
 - Data Journaling – Data blocks are written to the journal so data is written twice. High overhead, high reliability
 - Ordered – Data updates are flushed before metadata transactions are committed
 - Writeback – Data updates are left to the “sync” mechanism so files can contain garbage after a crash
 - Allows administrators to tune the reliability of their filesystem

Ext3 – Performance Improvements

- Directory Scalability
 - HTree – similar to B-tree with low depth and high fanout
 - Directories are viewed as a hash table instead of linked list
 - To find a directory entry, simply hash the name and look at that location
 - Backwards compatible with ext2
 - Creating 100,000 files
 - 38 minutes with htree disabled
 - 11 seconds with htree enabled

Ext3 – Performance Improvements

- Improved Directory Allocation
 - Borrowed from the BSD world
 - Allocate subdirectories in the same block group as their parent
 - Allocate space for children when a directory is allocated
 - Put unrelated directories in different block groups (if possible)
 - Significantly improves data locality
 - As always, backwards compatible

Ext3 – Limitations

- Journaling means data gets written to disk twice
 - slower than ext2
- Backward compatibility with ext2 means that it can't make use of a number of newer filesystem features (extents, tail packing, etc)

Ext4

- Forked from Ext3 to address
 - 16TB filesystem limit due to 32-bit block numbers
 - 32K entry limit on subdirectories
 - timestamp resolution (seconds)
 - performance

Ext4

- 48-bit block numbers -> 1EB filesystem
- Indirect blocks replaced with extents
 - A single descriptor for a range of contiguous blocks
 - Logical, Length, Physical

```
struct ext4_extent {  
    __le32 ee_block;    /* first logical block extent covers */  
    __le16 ee_len;      /* number of blocks covered by extent */  
    __le16 ee_start_hi; /* high 16 bits of physical block */  
    __le32 ee_start;    /* low 32 bits of physical block */  
};
```

- 16TB file (32 bit logical block number)
- Max extent 128MB (15 bits)

Ext4

- Persistent pre-allocation – space is reserved and likely contiguous
 - fallocate() call
- Nanosecond timestamps
- Journal checksumming

XFS

- Written in 1996 as a replacement for EFS on Irix (The Silicon Graphics UNIX system)
 - EFS was not supporting the I/O rates that SGI customers needed
 - EFS didn't support filesystems that were greater than 8GB in size, or files greater than 2GB in size

XFS Features

- 9EB filesystems
- 9EB files
- Journaled
- Variable block size
- Good sparse file support

XFS Features

- Implementing Files
 - Inodes contain links to extents
 - If more extents are needed than fit in the inode, a B-tree is used to manage the list of extents
 - Permits efficient implementation of sparse files
- Pre-allocation – file space allocated, but nothing written (like Ext4)

XFS Features

- Support for Large Numbers of Files
 - Inodes are generated on the fly
 - A B-Tree keeps track of the locations of allocated inodes
 - Use 64-bit inode numbers
 - Much more complex than the simple methods used by filesystems like ext2 and ext3

XFS Features

- Implementing Directories
 - Entries are stored in a B-tree indexed on the filename
 - Filenames are hashed to make the B-tree implementation faster

XFS Features

- Managing Free Space
 - 2 B-trees used
 - One uses the extent location as the index
 - The other uses the length of the extent as the index
 - Allows for free space to be found quickly

XFS Features

- Delayed Allocation (Allocate-on-Flush)
 - Don't allocate space for files until you actually write them out
 - Much more likely to have contiguous allocation
 - Temporary files will never go to the disk

JFFS

- Journaling Flash File System
- Written by Axis Communications AB in Sweden
- A filesystem for embedded systems
 - Intended to be used on Flash memory devices

JFFS

- Flash Characteristics
 - Divided into blocks of a given size
 - Writes to a block require writing the entire block
 - Lifetime of Flash is measured in “erase” cycles
 - Hundreds of thousands

JFFS

- Configured as a Log-Structured Filesystem
 - Modifications write nodes to a Flash block
 - When the block fills, the FS consults a list of unused blocks to find the next block to write to
 - When the unused blocks list hits a threshold, garbage collection is performed

JFFS

- Garbage Collection
 - Selects a block from the dirty list or the clean list
 - Writes any valid nodes in it to the tail of the transaction log
 - Ensures an even distribution of writes to blocks


JFFS

- How do reads occur?
 - On boot, the entire medium is scanned
 - The most recent version of each inode is brought into memory


WAFL File System

- Used on Network Appliance's distributed file system appliances
- “Write-anywhere file layout”
- Serves NFS, CIFS, http, ftp
- Random I/O optimized, write optimized
 - NVRAM for write caching
- Similar to Berkeley Fast File System, with extensive modifications
- Uses novel approach for writing data to disk
 - Ensures that the filesystem is always consistent
 - Allows for easy creation of snapshots


The WAFL File Layout




Snapshots in WAFL



(a) Before a snapshot.

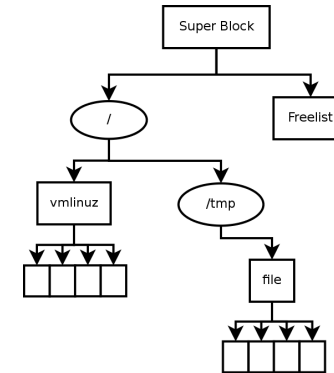


(b) After a snapshot, before any blocks change.




(c) After block D has changed to D'.

WAFL



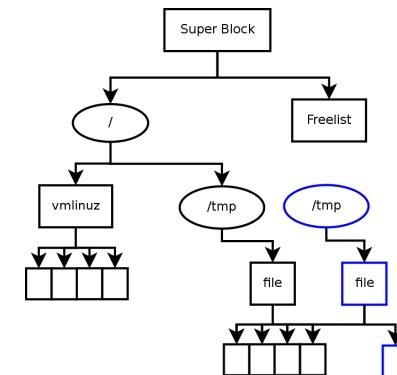
Initial Filesystem State

WAFL



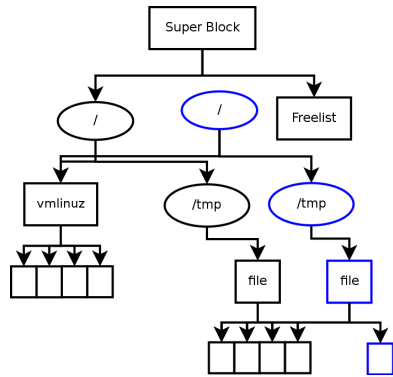
Create a new "file" inode

WAFL



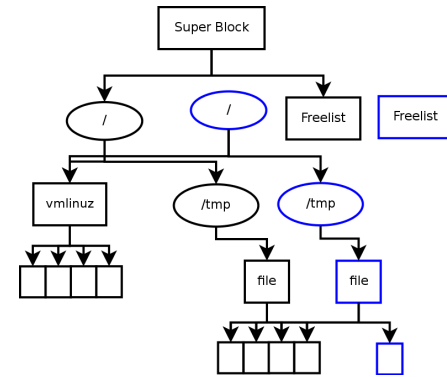
Create a new "/tmp" directory inode

WAFL



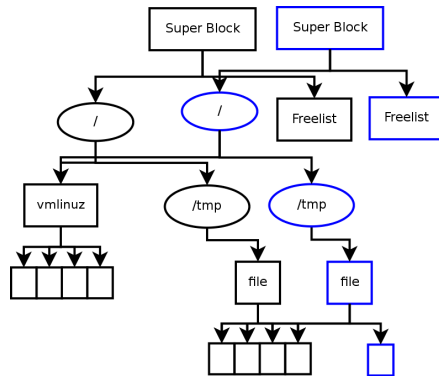
Create a new root directory inode

WAFL



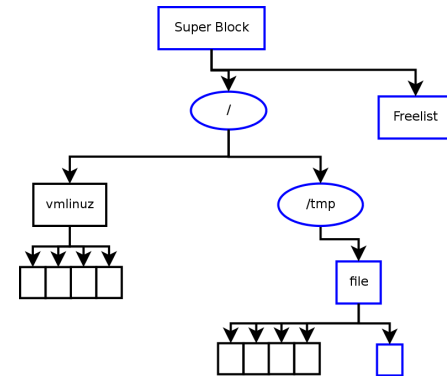
Update a copy of the freelist

WAFL



Create a new Super Block

WAFL



Overwrite the Old Superblock with the New One

ReiserFS4

- Attempted to unify the concepts of files and directories
 - Multics had directories as special types of file
 - ReiserFS4 defines files and directories as the same thing
 - A file can be accessed using open or opendir
 - File attributes implemented as “subfiles” of a given file
- Uses a plug-in architecture
 - allows new kinds of “files” to be implemented via a plug-in
 - very extensible

NTFS

- Allows files to have multiple data streams
 - Files are not just a sequence of bytes
- Case sensitive/insensitive filenames
- Per File/Directory Encryption
 - handled outside of the filesystem

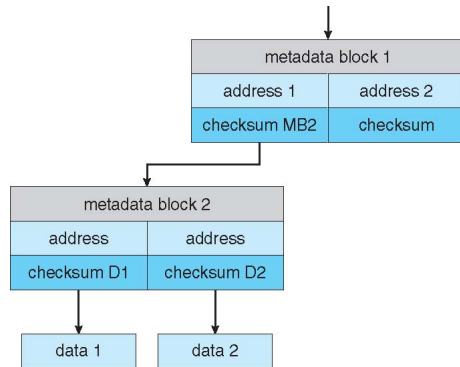
ZFS

- Developed by Sun
 - Open source
 - Originally stood for Zettabyte File System
- Performs copy on write similar to WAFL
 - Adds the ability to keep snapshots of filesystem state
- Includes an “intent log” to describe operations and aid in recovery
- Uses variable block sizes


ZFS

- ZFS includes checksums of all data and metadata
 - Checksums kept with pointer to object, to detect if object is the right one and whether it changed
 - Can detect and correct data and metadata corruption
- ZFS also removes the notion of volumes, partitions
 - Disks allocated in pools
 - Filesystems with a pool share that pool, use and release space like “malloc” and “free” memory allocate / release calls


ZFS Checksums All Metadata and Data



Traditional and Pooled Storage in ZFS



(a) Traditional volumes and file systems.



(b) ZFS and pooled storage.

FUSE

- Filesystem in Userspace
- Loadable module that sends file I/O to a filesystem implementation in userspace

