

Operating Systems

Network Stack

Network Communication

- Key functionality in an OS
 - Drove the development of BSD and Xinu
 - Everything interesting communicates in some way
- The Internet Protocol suite enables the internetworking of different network types and technologies
 - Even though Ethernet is now ubiquitous, the addressing mechanism does not allow location-based routing
- Layers of network functionality and protocols give rise to the network stack

Internet Protocols

- IP – Internet Protocol, Layer 3
 - Operates on packets, or datagrams, generally embedded in one or more Data Link (Ethernet) frames
- UDP – User Datagram protocol
 - Minimal Transport (Layer 4) protocol which includes port numbers to demultiplex inside a host
- ARP – Address Resolution Protocol
 - Binds global Layer 3 addresses to local Layer 2 addresses

Internet Protocols


- DHCP – Dynamic Host Configuration Protocol
 - Allows a server to provide an IP address, default gateway, nameserver, etc.
- ICMP – Internet Control Message Protocol
 - Error and informational messages
 - Useful for debugging
 - ping uses Echo Request and Echo Response ICMP messages

Xinu Network Processes

- Single network input process called *netin*
- IP output process called *ipout*
- The software uses the timed message receipt call *recvtime*
- When a process sends a network message, the *netin* process sends an internal message (with *send*) when the response is received
- If the timer expires first, *recvtime* returns the message *TIMEOUT*

Xinu Network Processes

- Coordination between the *netin* process and other waiting processes occurs via UDP-specific queues or ARP table entries



Kernel Threads


- The Xinu network stack is realized in terms of processes, which we have considered to be application programs
- Unix/Linux does essentially the same thing with threads running in kernel space – kernel threads
 - Show up in *ps* output surrounded by [brackets]
- Background system-level tasks like protocol processing, buffer management are implemented this way
 - A pure microkernel approach would put them in userspace, interacting with messages

Kernel Threads

- One of the Linux kernel threads is called *ksoftirqd*
- Soft IRQs are raised by interrupt handlers to request finishing bottom half processing tasks that can be deferred
 - The goal is to minimize processing time while interrupts are deferred
- Another is *events/0*, which provides a work queue for other kernel tasks
- Either could be used to restructure *netin* and *ipout* in Xinu

Xinu Network Processes

- netin is the only process that reads network packets – it cannot block waiting for a packet
- Incoming IP packets may require the generation of a response (e.g. ping)
- Transmission of outgoing IP packets may require an ARP exchange
- The processes are decoupled and utilize a queue



arp.h

```

/* Items related to ARP - definition of cache and the packet format */

#define ARP_HALEN 6 /* Size of Ethernet MAC address */
#define ARP_PALEN 4 /* Size of IP address */

#define ARP_HTYPE 1 /* Ethernet hardware type */
#define ARP_PTYPE 0x0800 /* IP protocol type */

#define ARP_OP_REQ 1 /* Request op code */
#define ARP_OP_RPLY 2 /* Reply op code */

#define ARP_SIZ 16 /* Number of entries in a cache */

#define ARP_RETRY 3 /* Num. retries for ARP request */

#define ARP_TIMEOUT 300 /* Retry timer in milliseconds */

/* State of an ARP cache entry */

#define AR_FREE 0 /* Slot is unused */
#define AR_PENDING 1 /* Resolution in progress */
#define AR_RESOLVED 2 /* Entry is valid */
  
```

arp.h

```

#pragma pack(2)
struct arppacket { /* ARP packet for IP & Ethernet */
    byte arp_ethdst[ETH_ADDR_LEN]; /* Ethernet dest. MAC addr */
    byte arp_ethsrc[ETH_ADDR_LEN]; /* Ethernet source MAC address */
    uint16 arp_etype; /* Ethernet type field */
    uint16 arp_htype; /* ARP hardware type */
    uint16 arp_ptype; /* ARP protocol type */
    byte arp_hlen; /* ARP hardware address length */
    byte arp_plen; /* ARP protocol address length */
    uint16 arp_op; /* ARP operation */
    byte arp_sndha[ARP_HALEN]; /* ARP sender's Ethernet addr */
    uint32 arp_sndpa; /* ARP sender's IP address */
    byte arp_tarha[ARP_HALEN]; /* ARP target's Ethernet addr */
    uint32 arp_tarpa; /* ARP target's IP address */
};
#pragma pack()

struct arprent { /* Entry in the ARP cache */
    int32 arstate; /* State of the entry */
    uint32 arpaddr; /* IP address of the entry */
    pid32 arpid; /* Waiting process or -1 */
    byte arhaddr[ARP_HALEN]; /* Ethernet address of the entry */
};
  
```

arp.c

```

struct arprent arpcache[ARP_SIZ]; /* ARP cache */

/*-----
 * arp_init - Initialize ARP cache for an Ethernet interface
 *-----
 */
void arp_init(void)
{
    int32 i; /* ARP cache index */

    for (i=1; i<ARP_SIZ; i++) { /* Initialize cache to empty */
        arpcache[i].arstate = AR_FREE;
    }
}
  
```

arp.c

```

/*-----
 * arp_resolve - Use ARP to resolve an IP address to an Ethernet address
 *-----
 */
status arp_resolve (
    uint32 nxthop, /* Next-hop address to resolve */
    byte mac[ETH_ADDR_LEN] /* Array into which Ethernet */
) /* address should be placed */
{
    intmask mask; /* Saved interrupt mask */
    struct arppacket apkt; /* Local packet buffer */
    int32 i; /* Index into arpcache */
    int32 slot; /* ARP table slot to use */
    struct arprent *arptr; /* Ptr to ARP cache entry */
    int32 msg; /* Message returned by recvtime */

    /* Use MAC broadcast address for IP limited broadcast */

    if (nxthop == IP_BCAST) {
        memcpy(mac, NetData.ethbcast, ETH_ADDR_LEN);
        return OK;
    }
}

```

arp.c

```

/* Use MAC broadcast address for IP network broadcast */

if (nxthop == NetData.ipbcast) {
    memcpy(mac, NetData.ethbcast, ETH_ADDR_LEN);
    return OK;
}

/* Ensure only one process uses ARP at a time */

mask = disable();

/* See if next hop address is already present in ARP cache */

for (i=0; i<ARP_SIZ; i++) {
    arptr = &arpcache[i];
    if (arptr->arstate == AR_FREE) {
        continue;
    }
    if (arptr->arpaddr == nxthop) { /* Address is in cache */
        break;
    }
}

```

arp.c

```

if (i < ARP_SIZ) { /* Entry was found */

    /* If entry is resolved - handle and return */

    if (arptr->arstate == AR_RESOLVED) {
        memcpy(mac, arptr->arhaddr, ARP_HALEN);
        restore(mask);
        return OK;
    }

    /* Entry is already pending - return error because */
    /* only one process can be waiting at a time */

    if (arptr->arstate == AR_PENDING) {
        restore(mask);
        return SYSERR;
    }
}

/* IP address not in cache - allocate a new cache entry and */
/* send an ARP request to obtain the answer */

slot = arp_alloc();
if (slot == SYSERR) {
    restore(mask);
    return SYSERR;
}

```

arp.c

```

arptr = &arpcache[slot];
arptr->arstate = AR_PENDING;
arptr->arpaddr = nxthop;
arptr->arpid = currid;

/* Hand-craft an ARP Request packet */

memcpy(apkt.arp_ethdst, NetData.ethbcast, ETH_ADDR_LEN);
memcpy(apkt.arp_ethsrc, NetData.ethucast, ETH_ADDR_LEN);
apkt.arp_ethtype = ETH_ARP; /* Packet type is ARP */
apkt.arp_hatype = ARP_ETYPE; /* Hardware type is Ethernet */
apkt.arp_ptype = ARP_PTYPE; /* Protocol type is IP */
apkt.arp_hlen = 0xff & ARP_HALEN; /* Ethernet MAC size in bytes */
apkt.arp_plen = 0xff & ARP_PALEN; /* IP address size in bytes */
apkt.arp_op = 0xffff & ARP_OP_REQ; /* ARP type is Request */
memcpy(apkt.arp_sndha, NetData.ethucast, ARP_HALEN);
apkt.arp_sndpa = NetData.ipucast; /* IP address of interface */
memset(apkt.arp_tarha, '\0', ARP_HALEN); /* Target HA is unknown */
apkt.arp_tarpa = nxthop; /* Target protocol address */

```

arp.c

```

/* Convert ARP packet from host to net byte order */
arp_hton(&apkt);

/* Convert Ethernet header from host to net byte order */
eth_hton((struct netpacket *)&apkt);

/* Send the packet ARP_RETRY times and await response */
msg = recvclr();
for (i=0; i<ARP_RETRY; i++) {
    write(ETHER0, (char *)&apkt, sizeof(struct arppacket));
    msg = recvtime(ARP_TIMEOUT);
    if (msg == TIMEOUT) {
        continue;
    } else if (msg == SYSERR) {
        restore(mask);
        return SYSERR;
    } else { /* entry is resolved */
        break;
    }
}

```

arp.c

```

/* If no response, return TIMEOUT */

if (msg == TIMEOUT) {
    arptr->arstate = AR_FREE; /* Invalidate cache entry */
    restore(mask);
    return TIMEOUT;
}

/* Return hardware address */

memcpy(mac, arptr->arhaddr, ARP_HALEN);
restore(mask);
return OK;
}

```

arp.c

```

/*-----
 * arp_in - Handle an incoming ARP packet
 *-----
 */
void arp_in (
    struct arppacket *pktptr /* Ptr to incoming packet */
)
{
    intmask mask; /* Saved interrupt mask */
    struct arppacket apkt; /* Local packet buffer */
    int32 slot; /* Slot in cache */
    struct arprent *arptr; /* Ptr to ARP cache entry */
    bool8 found; /* Is the sender's address in */
    /* the cache? */

    /* Convert packet from network order to host order */
    arp_ntoh(pktptr);

    /* Verify ARP is for IPv4 and Ethernet */

    if ( (pktptr->arp_hatype != ARP_HATYPE) ||
        (pktptr->arp_ptype != ARP_PTYPE) ) {
        freebuf((char *)pktptr);
        return;
    }
}

```

arp.c

```

/* Ensure only one process uses ARP at a time */

mask = disable();

/* Search cache for sender's IP address */

found = FALSE;

for (slot=0; slot < ARP_SIZ; slot++) {
    arptr = &arp_cache[slot];

    /* Skip table entries that are unused */

    if (arptr->arstate == AR_FREE) {
        continue;
    }

    /* If sender's address matches, we've found it */

    if (arptr->arpaddr == pktptr->arp_sndpa) {
        found = TRUE;
        break;
    }
}

```


arp.c

```
if (found) {
    /* Update sender's hardware address */
    memcpy(arptr->arhaddr, pktptr->arp_sndha, ARP_HALEN);

    /* If a process was waiting, inform the process */

    if (arptr->arstate == AR_PENDING) {
        /* Mark resolved and notify waiting process */
        arptr->arstate = AR_RESOLVED;
        send(arptr->arpid, OK);
    }
}

/* For an ARP reply, processing is complete */

if (pktptr->arp_op == ARP_OP_RPLY) {
    freebuf((char *)pktptr);
    restore(mask);
    return;
}
```

arp.c

```
/* The following is for an ARP request packet: if the local */
/* machine is not the target or the local IP address is not */
/* yet known, ignore the request (i.e., processing is complete)*/

if (!NetData.ipvalid) ||
    (pktptr->arp_tarpa != NetData.ipucast)) {
    freebuf((char *)pktptr);
    restore(mask);
    return;
}

/* Request has been sent to the local machine's address. So, */
/* add sender's info to cache, if not already present */

if (!found) {
    slot = arp_alloc();
    if (slot == SYSERR) { /* Cache is full */
        kprintf("ARP cache overflow on interface\n");
        freebuf((char *)pktptr);
        restore(mask);
        return;
    }
    arptr = &arp_cache[slot];
    arptr->arpaddr = pktptr->arp_sndpa;
    memcpy(arptr->arhaddr, pktptr->arp_sndha, ARP_HALEN);
    arptr->arstate = AR_RESOLVED;
}
```

arp.c

```
/* Hand-craft an ARP reply packet and send back to requester */

memcpy(apkt.arp_ethdst, pktptr->arp_sndha, ARP_HALEN);
memcpy(apkt.arp_ethsrc, NetData.ethucast, ARP_HALEN);
apkt.arp_ethtype = ETH_ARP; /* Frame carries ARP */
apkt.arp_htype = ARP_HTYPE; /* Hardware is Ethernet */
apkt.arp_ptype = ARP_PTYPE; /* Protocol is IP */
apkt.arp_hlen = ARP_HALEN; /* Ethernet address size */
apkt.arp_plen = ARP_PALEN; /* IP address size */
apkt.arp_op = ARP_OP_RPLY; /* Type is Reply */

/* Insert local Ethernet and IP address in sender fields */

memcpy(apkt.arp_sndha, NetData.ethucast, ARP_HALEN);
apkt.arp_sndpa = NetData.ipucast;

/* Copy target Ethernet and IP addresses from request packet */

memcpy(apkt.arp_tarha, pktptr->arp_sndha, ARP_HALEN);
apkt.arp_tarpa = pktptr->arp_sndpa;
```

arp.c

```
/* Convert ARP packet from host to network byte order */

arp_hton(&apkt);

/* Convert the Ethernet header to network byte order */

eth_hton((struct netpacket *)&apkt);

/* Send the reply */

write(ETHER0, (char *)&apkt, sizeof(struct arppacket));
freebuf((char *)pktptr);
restore(mask);
return;
}
```

arp.c

```
/*-----  
 * arp_alloc - Find a free slot or kick out an entry to create one  
 *-----  
 */  
int32 arp_alloc ()  
{  
    int32 slot; /* Slot in ARP cache */  
  
    /* Search for a free slot */  
  
    for (slot=0; slot < ARP_SIZ; slot++) {  
        if (arpcache[slot].arstate == AR_FREE) {  
            memset((char *)&arpcache[slot],  
                NULLCH, sizeof(struct arprent));  
            return slot;  
        }  
    }  
}
```

arp.c

```
/* Search for a resolved entry */  
  
for (slot=0; slot < ARP_SIZ; slot++) {  
    if (arpcache[slot].arstate == AR_RESOLVED) {  
        memset((char *)&arpcache[slot],  
            NULLCH, sizeof(struct arprent));  
        return slot;  
    }  
}  
  
/* At this point, all slots are pending (should not happen) */  
  
kprintf("ARP cache size exceeded\n");  
  
return SYSERR;  
}
```

arp.c

```
/*-----  
 * arp_ntoh - Convert ARP packet fields from net to host byte order  
 *-----  
 */  
void arp_ntoh(  
    struct arppacket *pktptr  
)  
{  
    pktptr->arp_htype = ntohs(pktptr->arp_htype);  
    pktptr->arp_ptype = ntohs(pktptr->arp_ptype);  
    pktptr->arp_op = ntohs(pktptr->arp_op);  
    pktptr->arp_sndpa = ntohl(pktptr->arp_sndpa);  
    pktptr->arp_tarpa = ntohl(pktptr->arp_tarpa);  
}  
  
/*-----  
 * arp_hton - Convert ARP packet fields from net to host byte order  
 *-----  
 */  
void arp_hton(  
    struct arppacket *pktptr  
)  
{  
    pktptr->arp_htype = htons(pktptr->arp_htype);  
    pktptr->arp_ptype = htons(pktptr->arp_ptype);  
    pktptr->arp_op = htons(pktptr->arp_op);  
    pktptr->arp_sndpa = htonl(pktptr->arp_sndpa);  
    pktptr->arp_tarpa = htonl(pktptr->arp_tarpa);  
}
```

net.h

```
#define NETSTK 8192 /* Stack size for network setup */  
#define NETPRIO 500 /* Network startup priority */  
#define NETBOOTFILE 128 /* Size of the netboot filename */  
  
/* Constants used in the networking code */  
  
#define ETH_ARP 0x0806 /* Ethernet type for ARP */  
#define ETH_IP 0x0800 /* Ethernet type for IP */  
#define ETH_IPv6 0x86DD /* Ethernet type for IPv6 */  
  
/* Format of an Ethernet packet carrying IPv4 and UDP */  
  
#pragma pack(2)  
struct netpacket {  
    byte net_ethdst[ETH_ADDR_LEN]; /* Ethernet dest. MAC address */  
    byte net_ethsrc[ETH_ADDR_LEN]; /* Ethernet source MAC address */  
    uint16 net_ethtype; /* Ethernet type field */  
    byte net_ipvh; /* IP version and hdr length */  
    byte net_iptas; /* IP type of service */  
    uint16 net_iplen; /* IP total packet length */  
    uint16 net_ipid; /* IP datagram ID */  
    uint16 net_ipfrag; /* IP flags & fragment offset */  
    byte net_ipttl; /* IP time-to-live */  
};
```

net.h

```
byte net_ipproto; /* IP protocol (actually type) */
uint16 net_ipcksum; /* IP checksum */
uint32 net_ipsrc; /* IP source address */
uint32 net_ipdst; /* IP destination address */
union {
    struct {
        uint16 net_udpport; /* UDP source protocol port */
        uint16 net_udpdport; /* UDP destination protocol port*/
        uint16 net_udpplen; /* UDP total length */
        uint16 net_udpcksum; /* UDP checksum */
        byte net_udpdata[1500-28]; /* UDP payload (1500-above)*/
    };
    struct {
        byte net_ictype; /* ICMP message type */
        byte net_iccode; /* ICMP code field (0 for ping) */
        uint16 net_iccksum; /* ICMP message checksum */
        uint16 net_icident; /* ICMP identifier */
        uint16 net_icseq; /* ICMP sequence number */
        byte net_icdata[1500-28]; /* ICMP payload (1500-above)*/
    };
};
```

net.h

```
#define PACKLEN sizeof(struct netpacket)

extern bpid32 netbufpool; /* ID of net packet buffer pool */

struct network {
    uint32 ipucast;
    uint32 ipbcast;
    uint32 ipmask;
    uint32 ipprefix;
    uint32 iprouter;
    uint32 bootserver;
    bool8 ipvalid;
    byte ethucast[ETH_ADDR_LEN];
    byte ethbcast[ETH_ADDR_LEN];
    char bootfile[NETBOOTFILE];
};

extern struct network NetData; /* Local Network Interface */
```

net.c

```
#include <xinu.h>
#include <stdio.h>

struct network NetData;
bpid32 netbufpool;

/*
 * net_init - Initialize network data structures and processes
 */

void net_init(void)
{
    int32 nbufs; /* Total no of buffers */

    /* Initialize the network data structure */

    memset((char *)&NetData, NULLCH, sizeof(struct network));

    /* Obtain the Ethernet MAC address */

    control(ETHER0, ETH_CTRL_GET_MAC, (int32)NetData.ethucast, 0);

    memset((char *)NetData.ethbcast, 0xFF, ETH_ADDR_LEN);
}
```

net.c

```
/* Create the network buffer pool */

nbufs = UDP_SLOTS * UDP_QSIZ + ICMP_SLOTS * ICMP_QSIZ + 1;

netbufpool = mkbufpool(PACKLEN, nbufs);

/* Initialize the ARP cache */

arp_init();

/* Initialize UDP */

udp_init();

/* Initialize ICMP */

icmp_init();

/* Initialize the IP output queue */

ipoqueue.ighead = 0;
ipoqueue.igtail = 0;
ipoqueue.igsem = semcreate(0);
```


net.c

```
if((int32)ipoqueue.iqsem == SYSERR) {
    panic("Cannot create ip output queue semaphore");
    return;
}

/* Create the IP output process */
resume(create(ipout, NETSTK, NETPRIO, "ipout", 0, NULL));

/* Create a network input process */
resume(create(netin, NETSTK, NETPRIO, "netin", 0, NULL));
}
```

net.c

```
/*-----
 * netin - Repeatedly read and process the next incoming packet
 *-----
 */

process netin ()
{
    struct netpacket *pkt; /* Ptr to current packet */
    int32 retval; /* Return value from read */

    /* Do forever: read a packet from the network and process */
    while(1) {

        /* Allocate a buffer */

        pkt = (struct netpacket *)getbuf(netbufpool);

        /* Obtain next packet that arrives */

        retval = read(ETHER0, (char *)pkt, PACKLEN);
        if(retval == SYSERR) {
            panic("Cannot read from Ethernet\n");
        }
    }
}
```

net.c

```
/* Convert Ethernet Type to host order */
eth_ntoh(pkt);

/* Demultiplex on Ethernet type */
switch (pkt->net_ethtype) {

    case ETH_ARP: /* Handle ARP */
        arp_in((struct arppacket *)pkt);
        continue;

    case ETH_IP: /* Handle IP */
        ip_in(pkt);
        continue;

    case ETH_IPv6: /* Handle IPv6 */
        freebuf((char *)pkt);
        continue;

    default: /* Ignore all other incoming packets */
        freebuf((char *)pkt);
        continue;
}
}
```

net.c

```
/*-----
 * eth_hton - Convert Ethernet type field to network byte order
 *-----
 */
void eth_hton(
    struct netpacket *pktptr
)
{
    pktptr->net_ethtype = htons(pktptr->net_ethtype);
}

/*-----
 * eth_ntoh - Convert Ethernet type field to host byte order
 *-----
 */
void eth_ntoh(
    struct netpacket *pktptr
)
{
    pktptr->net_ethtype = ntohs(pktptr->net_ethtype);
}
}
```

ip.h

```
/* ip.h - Constants related to Internet Protocol version 4 (IPv4) */

#define IP_BCAST 0xffffffff /* IP local broadcast address */
#define IP_THIS 0xffffffff /* "this host" src IP address */
#define IP_ALLZEROS 0x00000000 /* The all-zeros IP address */

#define IP_ICMP 1 /* ICMP protocol type for IP */
#define IP_UDP 17 /* UDP protocol type for IP */

#define IP_ASIZE 4 /* Bytes in an IP address */
#define IP_HDR_LEN 20 /* Bytes in an IP header */
#define IP_VH 0x45 /* IP version and hdr length */

#define IP_QOSIZ 8 /* Size of IP output queue */

/* Queue of outgoing IP packets waiting for ipout process */

struct igentry {
    int32 iqhead; /* Index of next packet to send */
    int32 iqtail; /* Index of next free slot */
    sid32 iqsem; /* Semaphore that counts pkts */
    struct netpacket *iqbuf[IP_QOSIZ]; /* Circular packet queue */
};

extern struct igentry ipoqueue; /* Network output queue */
```

ip.c

```
#include <xinu.h>

struct igentry ipoqueue; /* Queue of outgoing packets */

/*
 * ip_in - Handle an IP packet that has arrived over a network
 */

void ip_in(
    struct netpacket *pktptr /* Pointer to the packet */
)
{
    int32 icmplen; /* Length of ICMP message */

    /* Verify checksum */

    if (ipcksum(pktptr) != 0) {
        kprintf("IP header checksum failed\n\n");
        freebuf((char *)pktptr);
        return;
    }

    /* Convert IP header fields to host order */

    ip_ntoh(pktptr);
```

ip.c

```
/* Verify encapsulated protocol checksums and then convert */
/* the encapsulated headers to host byte order */

switch (pktptr->net_ipproto) {

    case IP_UDP:
        /* Skipping UDP checksum for now */
        udp_ntoh(pktptr);
        break;

    case IP_ICMP:
        icmplen = pktptr->net_iphlen - IP_HDR_LEN;
        if (icmp_cksum((char *)&pktptr->net_icthdr, icmplen) != 0) {
            freebuf((char *)pktptr);
            return;
        }
        icmp_ntoh(pktptr);
        break;

    default:
        break;
}
```

ip.c

```
/* Deliver 255.255.255.255 to local stack */

if (pktptr->net_ipdst == IP_BCAST) {
    ip_local(pktptr);
    return;
}

/* If we do not yet have a valid address, accept UDP packets */
/* (to get DHCP replies) and drop others */

if (!NetData.ipvalid) {
    if (pktptr->net_ipproto == IP_UDP) {
        ip_local(pktptr);
        return;
    } else {
        freebuf((char *)pktptr);
        return;
    }
}
```

ip.c

```
/* If packet is destined for us, accept it; otherwise, drop it */
if ( (pktptr->net_ipdst == NetData.ipucast) ||
     (pktptr->net_ipdst == NetData.ipbcast) ||
     (pktptr->net_ipdst == IP_BCAST) ) {
    ip_local(pktptr);
    return;
} else {
    /* Drop the packet */
    freebuf((char *)pktptr);
    return;
}
}
```

ip.c

```
/*-----
 * ip_send - Send an outgoing IP datagram from the local stack
 *-----
 */

status ip_send(
    struct netpacket *pktptr /* Pointer to the packet */
)
{
    intmask mask; /* Saved interrupt mask */
    uint32 dest; /* Destination of the datagram */
    int32 retval; /* Return value from functions */
    uint32 nxthop; /* Next-hop address */

    mask = disable();

    /* Pick up the IP destination address from the packet */
    dest = pktptr->net_ipdst;
```

ip.c

```
/* Loop back to local stack if destination 127.0.0.0/8 */
if ((dest & 0xff000000) == 0x7f000000) {
    ip_local(pktptr);
    restore(mask);
    return OK;
}

/* Loop back if the destination matches our IP unicast address */
if (dest == NetData.ipucast) {
    ip_local(pktptr);
    restore(mask);
    return OK;
}

/* Broadcast if destination is 255.255.255.255 */
if ( (dest == IP_BCAST) ||
     (dest == NetData.ipbcast) ) {
    memcpy(pktptr->net_ethdst, NetData.ethbroadcast,
           ETH_ADDR_LEN);
    retval = ip_out(pktptr);
    restore(mask);
    return retval;
}
```

ip.c

```
/* If destination is on the local network, next hop is the
 * destination; otherwise, next hop is default router */

if ( (dest & NetData.ipmask) == NetData.ipprefix) {
    /* Next hop is the destination itself */
    nxthop = dest;
} else {
    /* Next hop is default router on the network */
    nxthop = NetData.iprouter;
}

if (nxthop == 0) { /* Dest. invalid or no default route */
    freebuf((char *)pktptr);
    return SYSERR;
}
```

ip.c

```

/* Resolve the next-hop address to get a MAC address */
retval = arp_resolve(nxthop, pktptr->net_ethdst);
if (retval != OK) {
    freebuf((char *)pktptr);
    return SYSERR;
}

/* Send the packet */
retval = ip_out(pktptr);
restore(mask);
return retval;
}

```

ip.c

```

/*-----
 * ip_local - Deliver an IP datagram to the local stack
 *-----
 */
void ip_local(
    struct netpacket *pktptr /* Pointer to the packet */
)
{
    /* Use datagram contents to determine how to process */

    switch (pktptr->net_ipproto) {

        case IP_UDP:
            udp_in(pktptr);
            return;

        case IP_ICMP:
            icmp_in(pktptr);
            return;

        default:
            freebuf((char *)pktptr);
            return;
    }
}

```

ip.c

```

/*-----
 * ip_out - Transmit an outgoing IP datagram
 *-----
 */
status ip_out(
    struct netpacket *pktptr /* Pointer to the packet */
)
{
    uint16 cksum; /* Checksum in host byte order */
    int32 len; /* Length of ICMP message */
    int32 pktlen; /* Length of entire packet */
    int32 retval; /* Value returned by write */

    /* Compute total packet length */
    pktlen = pktptr->net_iplen + ETH_HDR_LEN;

    /* Convert encapsulated protocol to network byte order */
    switch (pktptr->net_ipproto) {

        case IP_UDP:
            pktptr->net_udpcksum = 0;
            udp_hon(pktptr);

            /* ...skipping UDP checksum computation */
            break;

```

ip.c

```

        case IP_ICMP:
            icmp_hon(pktptr);

            /* Compute ICMP checksum */

            pktptr->net_iccksum = 0;
            len = pktptr->net_iplen - IP_HDR_LEN;
            cksum = icmp_cksum((char *)&pktptr->net_ictype,
                                len);
            pktptr->net_iccksum = 0xffff & htons(cksum);
            break;

        default:
            break;
    }

    /* Convert IP fields to network byte order */
    ip_hon(pktptr);

```


ip.c

```

/* Compute IP header checksum */
pktptr->net_ipcksum = 0;
cksum = ipcksum(pktptr);
pktptr->net_ipcksum = 0xffff & htons(cksum);

/* Convert Ethernet fields to network byte order */
eth_hton(pktptr);

/* Send packet over the Ethernet */

retval = write(ETHER0, (char*)pktptr, pktlen);
freebuf((char *)pktptr);

if (retval == SYSERR) {
    return SYSERR;
} else {
    return OK;
}
}

```

ip.c

```

/*-----
 * ipcksum - Compute the IP header checksum for a datagram
 *-----
 */

uint16 ipcksum(
    struct netpacket *pkt /* Pointer to the packet */
)
{
    uint16 *hptr; /* Ptr to 16-bit header values */
    int32 i; /* Counts 16-bit values in hdr */
    uint16 word; /* One 16-bit word */
    uint32 cksum; /* Computed value of checksum */

    hptr = (uint16 *) &pkt->net_ipvh;

    /* Sum 16-bit words in the packet */

    cksum = 0;
    for (i=0; i<10; i++) {
        word = *hptr++;
        cksum += (uint32) htons(word);
    }
}

```

ip.c

```

/* Add in carry, and take the ones-complement */
|
cksum += (cksum >> 16);
cksum = 0xffff & ~cksum;

/* Use all-1s for zero */

if (cksum == 0xffff) {
    cksum = 0;
}
return (uint16) (0xffff & cksum);
}

```

ip.c

```

/*-----
 * ip_ntoh - Convert IP header fields to host byte order
 *-----
 */
void ip_ntoh(
    struct netpacket *pktptr
)
{
    pktptr->net_iphlen = ntohs(pktptr->net_iphlen);
    pktptr->net_ipid = ntohs(pktptr->net_ipid);
    pktptr->net_ipfrag = ntohs(pktptr->net_ipfrag);
    pktptr->net_ipsrc = ntohl(pktptr->net_ipsrc);
    pktptr->net_ipdst = ntohl(pktptr->net_ipdst);
}

/*-----
 * ip_hton - Convert IP header fields to network byte order
 *-----
 */
void ip_hton(
    struct netpacket *pktptr
)
{
    pktptr->net_iphlen = htons(pktptr->net_iphlen);
    pktptr->net_ipid = htons(pktptr->net_ipid);
    pktptr->net_ipfrag = htons(pktptr->net_ipfrag);
    pktptr->net_ipsrc = htonl(pktptr->net_ipsrc);
    pktptr->net_ipdst = htonl(pktptr->net_ipdst);
}

```

ip.c

```
/*-----  
 * ipout - Process that transmits IP packets from the IP output queue  
 *-----  
 */  
  
process ipout(void)  
{  
    struct netpacket *pktptr; /* Pointer to next the packet */  
    struct iqentry *ipqptr; /* Pointer to IP output queue */  
    uint32 destip; /* Destination IP address */  
    uint32 nxthop; /* Next hop IP address */  
    int32 retval; /* Value returned by functions */  
  
    ipqptr = &iqueue;  
  
    while(1) {  
  
        /* Obtain next packet from the IP output queue */  
  
        wait(ipqptr->iqsem);  
        pktptr = ipqptr->iqbuf[ipqptr->iqhead++];  
        if (ipqptr->iqhead >= IP_00SIZ) {  
            ipqptr->iqhead = 0;  
        }  
    }  
}
```

ip.c

```
/* Fill in the MAC source address */  
  
memcpy(pktptr->net_ethsrc, NetData.ethucast, ETH_ADDR_LEN);  
  
/* Extract destination address from packet */  
  
destip = pktptr->net_ipdst;  
  
/* Sanity check: packets sent to ioout should *not* */  
/* contain a broadcast address. */  
  
if ((destip == IP_BCAST) || (destip == NetData.ipbcast)) {  
    kprintf("ipout: encountered a broadcast\n");  
    freebuf((char *)pktptr);  
    continue;  
}  
  
/* Check whether destination is the local computer */  
  
if (destip == NetData.ipucast) {  
    ip_local(pktptr);  
    continue;  
}
```

ip.c

```
/* Check whether destination is on the local net */  
  
if ( (destip & NetData.ipmask) == NetData.ipprefix) {  
  
    /* Next hop is the destination itself */  
  
    nxthop = destip;  
} else {  
  
    /* Next hop is default router on the network */  
  
    nxthop = NetData.iprouter;  
}  
  
if (nxthop == 0) { /* Dest. invalid or no default route */  
    freebuf((char *)pktptr);  
    continue;  
}
```

ip.c

```
/* Use ARP to resolve next-hop address */  
  
retval = arp_resolve(nxthop, pktptr->net_ethdst);  
if (retval != OK) {  
    freebuf((char *)pktptr);  
    continue;  
}  
  
/* Use ipout to Convert byte order and send */  
  
ip_out(pktptr);  
}  
}
```

ip.c

```
/*-----  
 * ip_enqueue - Deposit an outgoing IP datagram on the IP output queue  
 *-----  
 */  
status ip_enqueue(  
    struct netpacket *pktptr /* Pointer to the packet */  
)  
{  
    intmask mask; /* Saved interrupt mask */  
    struct iqentry *iptr; /* Ptr. to network output queue */  
  
    /* Ensure only one process accesses output queue at a time */  
  
    mask = disable();  
  
    /* Enqueue packet on network output queue */  
  
    iptr = &iqueue;
```

ip.c

```
if (semcount(iptr->iqsem) >= IP_OQSZ) {  
    kprintf("ipout: output queue overflow\n");  
    freebuf((char *)pktptr);  
    restore(mask);  
    return SYSERR;  
}  
iptr->iqbuf[iptr->iqtail++] = pktptr;  
if (iptr->iqtail >= IP_OQSZ) {  
    iptr->iqtail = 0;  
}  
signal(iptr->iqsem);  
restore(mask);  
return OK;  
}
```

udp.h

```
/* udp.h - Declarations pertaining to User Datagram Protocol (UDP) */  
  
#define UDP_SLOTS 6 /* Number of open UDP endpoints */  
#define UDP_QSZ 8 /* Packets enqueued per endpoint */  
  
#define UDP_DHCP_CPORT 68 /* Port number for DHCP client */  
#define UDP_DHCP_SPORT 67 /* Port number for DHCP server */  
  
/* Constants for the state of an entry */  
  
#define UDP_FREE 0 /* Entry is unused */  
#define UDP_USED 1 /* Entry is being used */  
#define UDP_RECV 2 /* Entry has a process waiting */  
  
#define UDP_ANYIF -2 /* Register an endpoint for any */  
/* interface on the machine */  
  
#define UDP_HDR_LEN 8 /* Bytes in a UDP header */
```

udp.h

```
struct udpentry { /* Entry in the UDP endpoint tbl */  
    int32 udsta; /* State of entry: free/used */  
    uint32 udremip; /* Remote IP address (zero */  
    /* means "don't care") */  
    uint16 udremport; /* Remote protocol port number */  
    uint16 udlocport; /* Local protocol port number */  
    int32 udhead; /* Index of next packet to read */  
    int32 udtail; /* Index of next slot to insert */  
    int32 udcnt; /* Count of packets enqueued */  
    pid32 uidpid; /* ID of waiting process */  
    struct netpacket *udqueue[UDP_QSZ]; /* Circular packet queue */  
};  
  
extern struct udpentry udptab[];
```

udp.c

```
/* udp.c - udp_init, udp_in, udp_register, udp_send, udp_sendto, */
/*          udp_rcv, udp_rcvaddr, udp_release, udp_ntoh, udp_hton */

#include <xinu.h>

struct udptentry udptab[UDP_SLOTS]; /* Table of UDP endpoints */

/*-----
 * udp_init - Initialize all entries in the UDP endpoint table
 *-----
 */
void udp_init(void)
{
    int32 i;      /* Index into the UDP table */

    for(i=0; i<UDP_SLOTS; i++) {
        udptab[i].udstate = UDP_FREE;
    }

    return;
}
```

udp.c

```
/*-----
 * udp_in - Handle an incoming UDP packet
 *-----
 */
void udp_in(
    struct netpacket *pktptr /* Pointer to the packet */
)
{
    intmask mask; /* Saved interrupt mask */
    int32 i; /* Index into udptab */
    struct udptentry *udptr; /* Pointer to a udptab entry */

    /* Ensure only one process can access the UDP table at a time */

    mask = disable();

    for (i=0; i<UDP_SLOTS; i++) {
        udptr = &udptab[i];
        if (udptr->udstate == UDP_FREE) {
            continue;
        }
    }
}
```

udp.c

```
if ((pktptr->net_udpport == udptr->udlocport) &&
    ((udptr->udremport == 0) ||
     (pktptr->net_udpport == udptr->udremport)) &&
    ( ((udptr->udremip==0) ||
      (pktptr->net_ipsrc == udptr->udremip)))) ) {

    /* Entry matches incoming packet */

    if (udptr->udcount < UDP_QSIZ) {
        udptr->udcount++;
        udptr->udqueue[udptr->udtail++] = pktptr;
        if (udptr->udtail >= UDP_QSIZ) {
            udptr->udtail = 0;
        }
        if (udptr->udstate == UDP_RECV) {
            udptr->udstate = UDP_USED;
            send (udptr->udpid, OK);
        }
        restore(mask);
        return;
    }
}
}
```

udp.c

```
/* No match - simply discard packet */

freebuf((char *) pktptr);
restore(mask);
return;
}
```


udp.c

```

/*-----
 * udp_register - Register a remote IP, remote port & local port to
 *               receive incoming UDP messages from the specified
 *               remote site sent to the specified local port
 *-----
 */
uid32 udp_register (
    uint32 remip, /* Remote IP address or zero */
    uint16 remport, /* Remote UDP protocol port */
    uint16 locport /* Local UDP protocol port */
)
{
    intmask mask; /* Saved interrupt mask */
    int32 slot; /* Index into udptab */
    struct udpentry *udptr; /* Pointer to udptab entry */

    /* Ensure only one process can access the UDP table at a time */
    mask = disable();

```

udp.c

```

/* See if request already registered */

for (slot=0; slot<UDP_SLOTS; slot++) {
    udptr = &udptab[slot];
    if (udptr->udstate == UDP_FREE) {
        continue;
    }

    /* Look at this entry in table */

    if ( (remport == udptr->udremport) &&
        (locport == udptr->udlocport) &&
        (remip == udptr->udremip) ) {

        /* Request is already in the table */

        restore(mask);
        return SYSERR;
    }
}

```

udp.c

```

/* Find a free slot and allocate it */

for (slot=0; slot<UDP_SLOTS; slot++) {
    udptr = &udptab[slot];
    if (udptr->udstate != UDP_FREE) {
        continue;
    }
    udptr->udlocport = locport;
    udptr->udremport = remport;
    udptr->udremip = remip;
    udptr->udcount = 0;
    udptr->udhead = udptr->udtail = 0;
    udptr->udpid = -1;
    udptr->udstate = UDP_USED;
    restore(mask);
    return slot;
}

restore(mask);
return SYSERR;
}

```

udp.c

```

/*-----
 * udp_recv - Receive a UDP packet
 *-----
 */
int32 udp_recv (
    uid32 slot, /* Slot in table to use */
    char *buff, /* Buffer to hold UDP data */
    int32 len, /* Length of buffer */
    uint32 timeout /* Read timeout in msec */
)
{
    intmask mask; /* Saved interrupt mask */
    struct udpentry *udptr; /* Pointer to udptab entry */
    umsg32 msg; /* Message from recvtime() */
    struct netpacket *pkt; /* Pointer to packet being read */
    int32 i; /* Counts bytes copied */
    int32 msglen; /* Length of UDP data in packet */
    char *udataptr; /* Pointer to UDP data */

    /* Ensure only one process can access the UDP table at a time */

    mask = disable();

```

udp.c

```

/* Verify that the slot is valid */
if ((slot < 0) || (slot >= UDP_SLOTS)) {
    restore(mask);
    return SYSERR;
}

/* Get pointer to table entry */
udptr = &udptab[slot];

/* Verify that the slot has been registered and is valid */
if (udptr->udstate != UDP_USED) {
    restore(mask);
    return SYSERR;
}

```

udp.c

```

/* Wait for a packet to arrive */
if (udptr->udcount == 0) { /* No packet is waiting */
    udptr->udstate = UDP_RECV;
    udptr->udpip = currip;
    msg = recvclr();
    msg = recvtime(timeout); /* Wait for a packet */
    udptr->udstate = UDP_USED;
    if (msg == TIMEOUT) {
        restore(mask);
        return TIMEOUT;
    } else if (msg != OK) {
        restore(mask);
        return SYSERR;
    }
}
}

```

udp.c

```

/* Packet has arrived -- dequeue it */
pkt = udptr->udqueue[udptr->udhead++];
if (udptr->udhead >= UDP_QSIZ) {
    udptr->udhead = 0;
}
udptr->udcount--;

/* Copy UDP data from packet into caller's buffer */
msglen = pkt->net_udplen - UDP_HDR_LEN;
udataptr = (char *)pkt->net_udpdata;
if (len < msglen) {
    msglen = len;
}
for (i=0; i<msglen; i++) {
    *buff++ = *udataptr++;
}
freebuf((char *)pkt);
restore(mask);
return msglen;
}

```

udp.c

```

/*-----
 * udp_recvaddr - Receive a UDP packet and record the sender's address
 *-----
 */
int32 udp_recvaddr (
    uid32 slot, /* Slot in table to use */
    uint32 *remip, /* Loc for remote IP address */
    uint16 *remport, /* Loc for remote protocol port */
    char *buff, /* Buffer to hold UDP data */
    int32 len, /* Length of buffer */
    uint32 timeout /* Read timeout in msec */
)
{
    intmask mask; /* Saved interrupt mask */
    struct udpentry *udptr; /* Pointer to udptab entry */
    umsg32 msg; /* Message from recvtime() */
    struct netpacket *pkt; /* Pointer to packet being read */
    int32 msglen; /* Length of UDP data in packet */
    int32 i; /* Counts bytes copied */
    char *udataptr; /* Pointer to UDP data */

    /* Ensure only one process can access the UDP table at a time */
    mask = disable();

```

udp.c

```

/* Verify that the slot is valid */
if ((slot < 0) || (slot >= UDP_SLOTS)) {
    restore(mask);
    return SYSERR;
}

/* Get pointer to table entry */
udp_ptr = &udptab[slot];

/* Verify that the slot has been registered and is valid */
if (udp_ptr->udstate != UDP_USED) {
    restore(mask);
    return SYSERR;
}

```

udp.c

```

/* Wait for a packet to arrive */

if (udp_ptr->udcount == 0) { /* No packet is waiting */
    udp_ptr->udstate = UDP_RECV;
    udp_ptr->udp_id = currp_id;
    msg = recvclr();
    msg = recvtime(timeout); /* Wait for a packet */
    udp_ptr->udstate = UDP_USED;
    if (msg == TIMEOUT) {
        restore(mask);
        return TIMEOUT;
    } else if (msg != OK) {
        restore(mask);
        return SYSERR;
    }
}

/* Packet has arrived -- dequeue it */

pkt = udp_ptr->udqueue[udp_ptr->udhead++];
if (udp_ptr->udhead >= UDP_QSIZ) {
    udp_ptr->udhead = 0;
}

```

udp.c

```

/* Record sender's IP address and UDP port number */

*remip = pkt->net_ipsrc;
*remport = pkt->net_udpport;

udp_ptr->udcount--;

/* Copy UDP data from packet into caller's buffer */

msglen = pkt->net_udplen - UDP_HDR_LEN;
udataptr = (char *)pkt->net_udpdata;
if (len < msglen) {
    msglen = len;
}
for (i=0; i<msglen; i++) {
    *buff++ = *udataptr++;
}
freebuf((char *)pkt);
restore(mask);
return msglen;
}

```

udp.c

```

/*-----
 * udp_send - Send a UDP packet using info in a UDP table entry
 *-----
 */
status udp_send (
    uid32 slot, /* Table slot to use */
    char *buff, /* Buffer of UDP data */
    int32 len /* Length of data in buffer */
)
{
    intmask mask; /* Saved interrupt mask */
    struct netpacket *pkt; /* Pointer to packet buffer */
    int32 pktlen; /* Total packet length */
    static uint16 ident = 1; /* Datagram IDENT field */
    char *udataptr; /* Pointer to UDP data */
    uint32 remip; /* Remote IP address to use */
    uint16 remport; /* Remote protocol port to use */
    uint16 locport; /* Local protocol port to use */
    uint32 locip; /* Local IP address taken from */
    /* the interface */
    struct udpentry *udp_ptr; /* Pointer to table entry */

    /* Ensure only one process can access the UDP table at a time */
    mask = disable();

```

udp.c

```
/* Verify that the slot is valid */
if ( (slot < 0) || (slot >= UDP_SLOTS) ) {
    restore(mask);
    return SYSERR;
}

/* Get pointer to table entry */
udptr = &udptab[slot];

/* Verify that the slot has been registered and is valid */
if (udptr->udstate == UDP_FREE) {
    restore(mask);
    return SYSERR;
}

/* Verify that the slot has a specified remote address */
remip = udptr->udremip;
if (remip == 0) {
    restore(mask);
    return SYSERR;
}
```

udp.c

```
locip = NetData.ipucast;
remport = udptr->udremport;
locport = udptr->udlocport;

/* Allocate a network buffer to hold the packet */
pkt = (struct netpacket *)getbuf(netbufpool);

if ((int32)pkt == SYSERR) {
    restore(mask);
    return SYSERR;
}

/* Compute packet length as UDP data size + fixed header size */
pktlen = ((char *)&pkt->net_udpdata - (char *)pkt) + len;
```

udp.c

```
/* Create a UDP packet in pkt */

memcpy((char *)pkt->net_ethsrc, NetData.ethucast, ETH_ADDR_LEN);
pkt->net_ethtype = 0x0800; /* Type is IP */
pkt->net_ipvh = 0x45; /* IP version and hdr length */
pkt->net_iptos = 0x00; /* Type of service */
pkt->net_iplen = pktlen - ETH_HDR_LEN; /* Total IP datagram length */
pkt->net_ipid = ident++; /* Datagram gets next IDENT */
pkt->net_ipfrag = 0x0000; /* IP flags & fragment offset */
pkt->net_ipttl = 0xff; /* IP time-to-live */
pkt->net_ipproto = IP_UDP; /* Datagram carries UDP */
pkt->net_ipcksum = 0x0000; /* initial checksum */
pkt->net_ipsrc = locip; /* IP source address */
pkt->net_ipdst = remip; /* IP destination address */

pkt->net_udpport = locport; /* Local UDP protocol port */
pkt->net_udpdport = remport; /* Remote UDP protocol port */
pkt->net_udplen = (uint16)(UDP_HDR_LEN+len); /* UDP length */
pkt->net_udpcksum = 0x0000; /* Ignore UDP checksum */
udataptr = (char *) pkt->net_udpdata;
for (; len>0; len--) {
    *udataptr++ = *buff++;
}
```

udp.c

```
/* Call ipsend to send the datagram */

ip_send(pkt);
restore(mask);
return OK;
}
```


udp.c

```

/*-----
 * udp_sendto - Send a UDP packet to a specified destination
 *-----
 */
status udp_sendto (
    uid32 slot, /* UDP table slot to use */
    uint32 remip, /* Remote IP address to use */
    uint16 remport, /* Remote protocol port to use */
    char *buff, /* Buffer of UDP data */
    int32 len /* Length of data in buffer */
)
{
    intmask mask; /* Saved interrupt mask */
    struct netpacket *pkt; /* Pointer to a packet buffer */
    int32 pktlen; /* Total packet length */
    static uint16 ident = 1; /* Datagram IDENT field */
    struct udpentry *udptr; /* Pointer to a UDP table entry */
    char *udataptr; /* Pointer to UDP data */

    /* Ensure only one process can access the UDP table at a time */
    mask = disable();

```

udp.c

```

/* Verify that the slot is valid */
if ( (slot < 0) || (slot >= UDP_SLOTS) ) {
    restore(mask);
    return SYSERR;
}

/* Get pointer to table entry */
udptr = &udptab[slot];

/* Verify that the slot has been registered and is valid */
if (udptr->udstate == UDP_FREE) {
    restore(mask);
    return SYSERR;
}

/* Allocate a network buffer to hold the packet */
pkt = (struct netpacket *)getbuf(netbufpool);

if ((int32)pkt == SYSERR) {
    restore(mask);
    return SYSERR;
}

```

udp.c

```

/* Compute packet length as UDP data size + fixed header size */
pktlen = ((char *)&pkt->net_udpdata - (char *)pkt) + len;

/* Create UDP packet in pkt */

memcpy((char *)pkt->net_ethsrc, NetData.ethucast, ETH_ADDR_LEN);
pkt->net_ethtype = 0x0800; /* Type is IP */
pkt->net_ipvh = 0x45; /* IP version and hdr length */
pkt->net_iptos = 0x00; /* Type of service */
pkt->net_iphlen = pktlen - ETH_HDR_LEN; /* total IP datagram length */
pkt->net_ipid = ident++; /* Datagram gets next IDENT */
pkt->net_ipfrag = 0x0000; /* IP flags & fragment offset */
pkt->net_ipttl = 0xff; /* IP time-to-live */
pkt->net_ipproto = IP_UDP; /* Datagram carries UDP */
pkt->net_ipcksum = 0x0000; /* Initial checksum */

```

udp.c

```

pkt->net_ipsrc = NetData.ipucast; /* IP source address */
pkt->net_ipdst = remip; /* IP destination address */
pkt->net_udpport = udptr->udlocport; /* local UDP protocol port */
pkt->net_udpport = remport; /* Remote UDP protocol port */
pkt->net_udpplen = (uint16)(UDP_HDR_LEN+len); /* UDP length */
pkt->net_udpcksum = 0x0000; /* Ignore UDP checksum */
udataptr = (char *) pkt->net_udpdata;
for (; len>0; len--) {
    *udataptr++ = *buff++;
}

/* Call ip_send to send the datagram */

ip_send(pkt);
restore(mask);
return OK;
}

```

```

/*-----
 * udp_release - Release a previously-registered UDP slot
 *-----
 */
status udp_release (
    uid32 slot /* Table slot to release */
)
{
    intmask mask; /* Saved interrupt mask */
    struct udpentry *udptr; /* Pointer to udptab entry */
    struct netpacket *pkt; /* pointer to packet being read */

    /* Ensure only one process can access the UDP table at a time */

    mask = disable();

    /* Verify that the slot is valid */

    if ( (slot < 0) || (slot >= UDP_SLOTS) ) {
        restore(mask);
        return SYSERR;
    }

    /* Get pointer to table entry */

    udptr = &udptab[slot];

```

udp.c

udp.c

```

/* Verify that the slot has been registered and is valid */

if (udptr->udstate == UDP_FREE) {
    restore(mask);
    return SYSERR;
}

/* Defer rescheduling to prevent freebuf from switching context */

resched_cntl(DEFER_START);
while (udptr->udcount > 0) {
    pkt = udptr->udqueue[udptr->udhead++];
    if (udptr->udhead >= UDP_QSIZ) {
        udptr->udhead = 0;
    }
    freebuf((char *)pkt);
    udptr->udcount--;
}
udptr->udstate = UDP_FREE;
resched_cntl(DEFER_STOP);
restore(mask);
return OK;
}

```

Summary

- Simple protocol stack
- Operation similar to device drivers with components coupled by queues
- Message passing mechanism used like a soft interrupt (with timeout)