

Operating Systems

High-Level Memory Management in Xinu

1

High-Level Memory Management

- Distinct from the lower-level `getmem()` `freemem()` interfaces
 - Provides basic memory management
 - Explicit memory allocation and release allow memory to be exhausted
- Higher-level memory management can be employed to divide resources more fairly
- A motivating example is network-centric applications that send or receive messages

2

Buffer Pools

- Buffer pools are a common design pattern
- Divide memory into a set of pools
- Each pool contains a fixed number of blocks of the same size, much like a slab allocator
 - Buffer reflects the intended use for I/O
- Each buffer pool is identified by an integer, a buffer pool ID
- Buffer allocation is synchronous in that a process requesting a buffer will be blocked until the request can be satisfied

3

Buffer Pools

```
/* bufpool.h */

#ifndef NBPOOLS
#define NBPOOLS 20 /* Maximum number of buffer pools */
#endif

#ifndef BP_MAXB
#define BP_MAXB 8192 /* Maximum buffer size in bytes */
#endif

#ifndef BP_MINB
#define BP_MINB 8 /* Minimum buffer size in bytes */
#endif
#ifndef BP_MAXN
#define BP_MAXN 2048 /* Maximum number of buffers in a pool */
#endif
```

4

```

struct bentry {          /* Description of a single buffer pool */
    struct bentry *bpnext; /* pointer to next free buffer */
    sid32 bpsem;          /* semaphore that counts buffers */
                          /* currently available in the pool */
    uint32 bpsize;        /* size of buffers in this pool */
};

extern struct bentry buftab[]; /* Buffer pool table */
extern bpid32 nbpools;        /* current number of allocated pools */

```

5

Allocating a Buffer

```

/* getbuf.c - getbuf, nbgetbuf */

#include <xinu.h>

/*-----
 * getbuf -- get a buffer from a preestablished buffer pool
 *-----
 */
char *getbuf(
    bpid32 poolid /* index of pool in buftab */
)
{
    intmask mask; /* saved interrupt mask */
    struct bentry *bpptr; /* pointer to entry in buftab */
    struct bentry *bufptr; /* pointer to a buffer */

```

6

```

    mask = disable();

    /* Check arguments */

    if ( (poolid < 0 || poolid >= nbpools) ) {
        restore(mask);
        return (char *)SYSERR;
    }
    bpptr = &buftab[poolid];

    /* Wait for pool to have > 0 buffers and allocate a buffer */

    wait(bpptr->bpsem);
    bufptr = bpptr->bpnext;

    /* Unlink buffer from pool */

    bpptr->bpnext = bufptr->bpnext;

    /* Record pool ID in first four bytes of buffer and skip */

```

7

```

    *(bpid32 *)bufptr = poolid;
    bufptr = (struct bentry *) (sizeof(bpid32) + (char *)bufptr);
    restore(mask);
    return (char *)bufptr;
}

/*-----
 * nbgetbuf - a non-blocking version of getbuf
 *-----
 */
char *nbgetbuf(
    bpid32 poolid /* index of pool in buftab */
)
{
    intmask mask;
    char *buf;

    mask = disable();

    /* Check arguments */

```

8

```

if ( (poolid < 0 || poolid >= nbpools) ) {
    restore(mask);
    return (char *)SYSERR;
}

/* If the call will block, return an error */

if (semcount(buftab[poolid].bpsem) <= 0) {
    restore(mask);
    return (char *)SYSERR;
}
buf = getbuf(poolid);
restore(mask);
return buf;
}

```

9

Returning Buffers To The Buffer Pool

```

/* freebuf.c - freebuf */

#include <xinu.h>

/*-----
 * freebuf - free a buffer that was allocated from a pool by getbuf
 *-----
 */
syscall freebuf(
    char      *bufaddr /* address of buffer to return */
)
{
    intmask mask;          /* saved interrupt mask */
    struct bentry *bpptr;   /* pointer to entry in buftab */
    bpid32 poolid;         /* ID of buffer's pool */
}

```

10

```

mask = disable();

/* Extract pool ID from integer prior to buffer address */

bufaddr -= sizeof(bpid32);
poolid = *(bpid32 *)bufaddr;
if (poolid < 0 || poolid >= nbpools) {
    restore(mask);
    return SYSERR;
}

/* Get address of correct pool entry in table */

bpptr = &buftab[poolid];

/* Insert buffer into list and signal semaphore */

((struct bentry *)bufaddr)->bpnext = bpptr->bpnext;
bpptr->bpnext = (struct bentry *)bufaddr;
signal(bpptr->bpsem);
restore(mask);
return OK;
}

```

11

Creating a Buffer Pool

```

/* mkbufpool.c - mkbufpool */

#include <xinu.h>

/*-----
 * mkbufpool - allocate memory for a buffer pool and link the
 * buffers
 *-----
 */
bpid32 mkbufpool(
    int32 bufsiz, /* size of a buffer in the pool */
    int32 numbufs /* number of buffers in the pool */
)
{
    intmask mask;          /* saved interrupt mask */
    bpid32 poolid;         /* ID of pool that is created */
    struct bentry *bpptr;   /* pointer to entry in buftab */
}

```

12

```

char *buf;          /* pointer to memory for buffer */

mask = disable();
if (bufsiz < BP_MINB || bufsiz > BP_MAXB
    || numbufs < 1 || numbufs > BP_MAXN
    || nbpools >= NBPOOLS) {
    restore(mask);
    return (bpid32)SYSERR;
}
/* Round request to a multiple of 4 bytes */

bufsiz = ( (bufsiz + 3) & (~3) );

buf = (char *)getmem( numbufs * (bufsiz+sizeof(bpid32)) );
if ((int32)buf == SYSERR) {
    restore(mask);
    return (bpid32)SYSERR;
}
poolid = nbpools++;
bpptr = &buftab[poolid];
bpptr->bpnext = (struct bentry *)buf;
bpptr->bpsize = bufsiz;

```

13

```

if ( (bpptr->bpsem = semcreate(numbufs)) == SYSERR) {
    nbpools--;
    restore(mask);
    return (bpid32)SYSERR;
}
bufsiz+=sizeof(bpid32);
for (numbufs--; numbufs>0; numbufs--) {
    bpptr = (struct bentry *)buf;
    buf += bufsiz;
    bpptr->bpnext = (struct bentry *)buf;
}
bpptr = (struct bentry *)buf;
bpptr->bpnext = (struct bentry *)NULL;
restore(mask);
return poolid;
}

```

14

Initializing the Buffer Pool Table

```

/* bufinit.c - bufinit */

#include <xinu.h>

struct bentry buftab[NBPOOLS]; /* buffer pool table */
bpid32 nbpools;

/*-----
 * bufinit -- initialize the buffer pool data structure
 *-----
 */
status bufinit(void)
{
    nbpools = 0;
    return OK;
}

```

15

end

16